

Assignment 2

Simple and Multiple Linear Regression

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-25

Course: Machine Learning 4DT905

Imports

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

Load data

```
In [2]: df = pd.read_csv("../data/Boston.csv", index_col=0)
```

Number of predictors and names

```
In [3]: df_names = df.columns.tolist()
print(f"Number of columns: {len(df_names)}")
print(f"Column names: {df_names}")
```

```
Number of columns: 14
Column names: ['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
 'ptratio', 'black', 'lstat', 'medv']
```

Statistical summary of predictors

```
In [4]: df.describe()
```

Out[4]:

	crim	zn	indus	chas	nox	rm	age
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000

Number of datapoints

In [5]: `print(f"Number of datapoints: {len(df)}")`

Number of datapoints: 506

Display data in table format

In [6]: `print(df.head(5))`

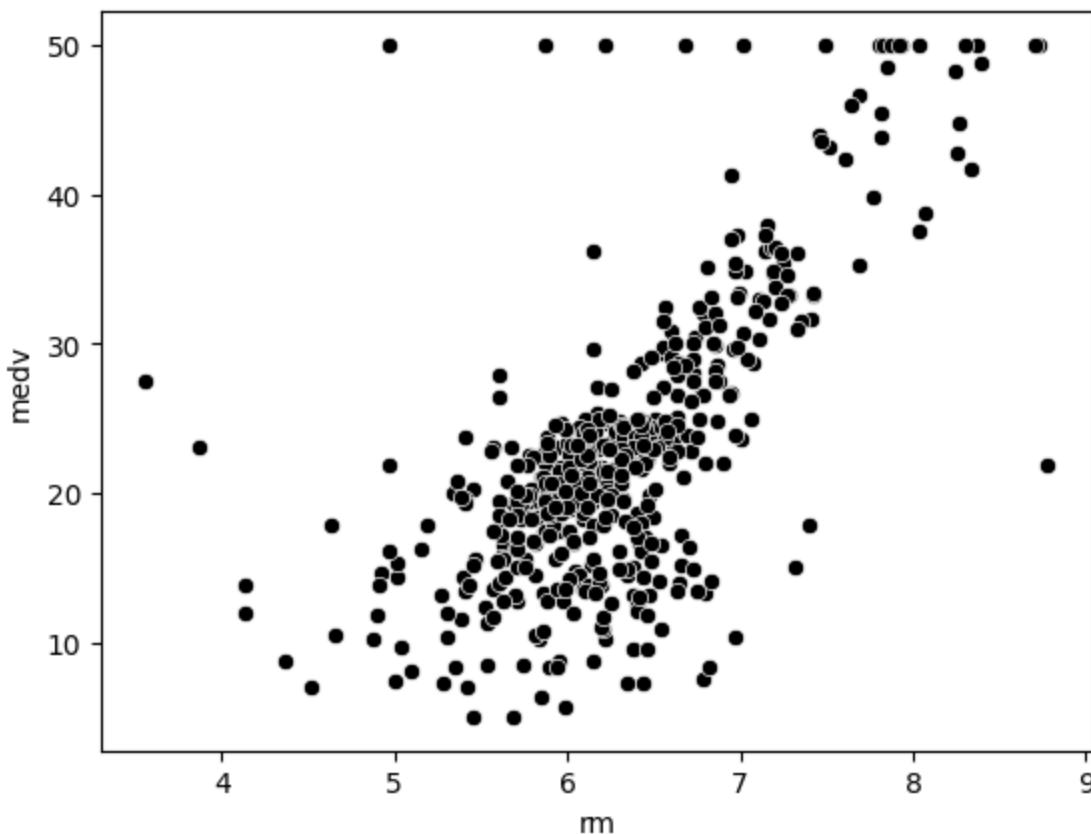
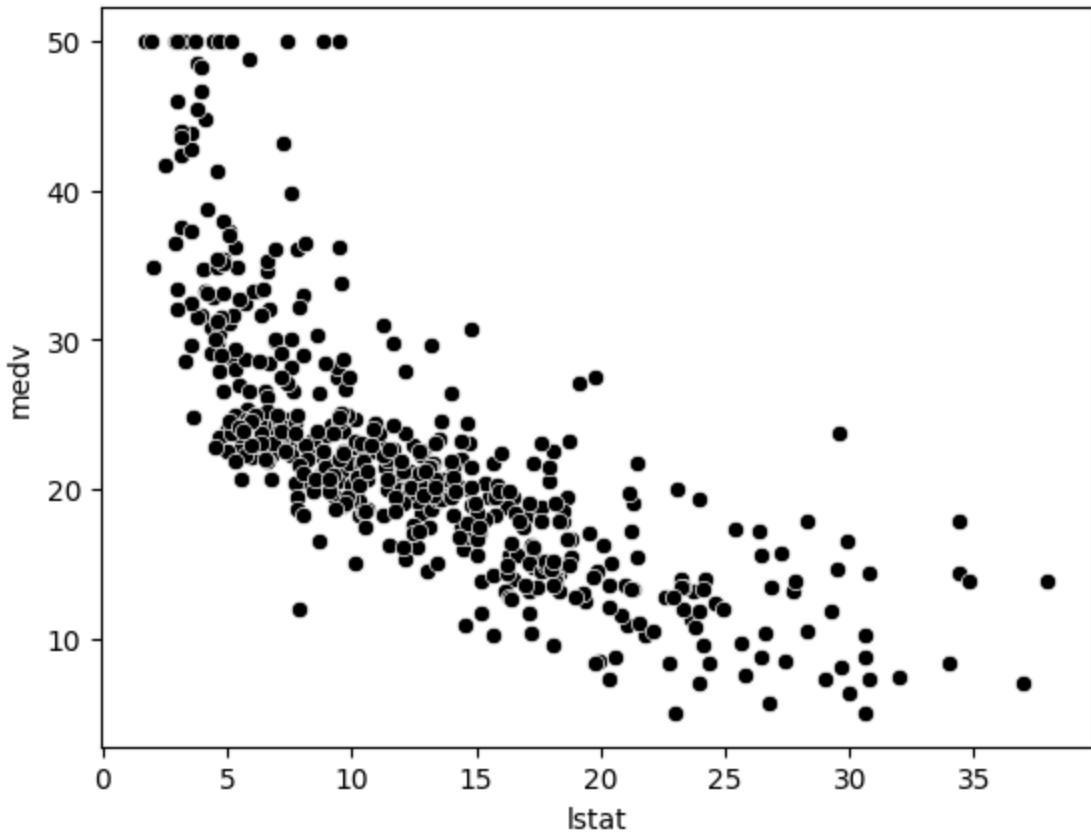
	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	
	black	lstat	medv									
1	396.90	4.98	24.0									
2	396.90	9.14	21.6									
3	392.83	4.03	34.7									
4	394.63	2.94	33.4									
5	396.90	5.33	36.2									

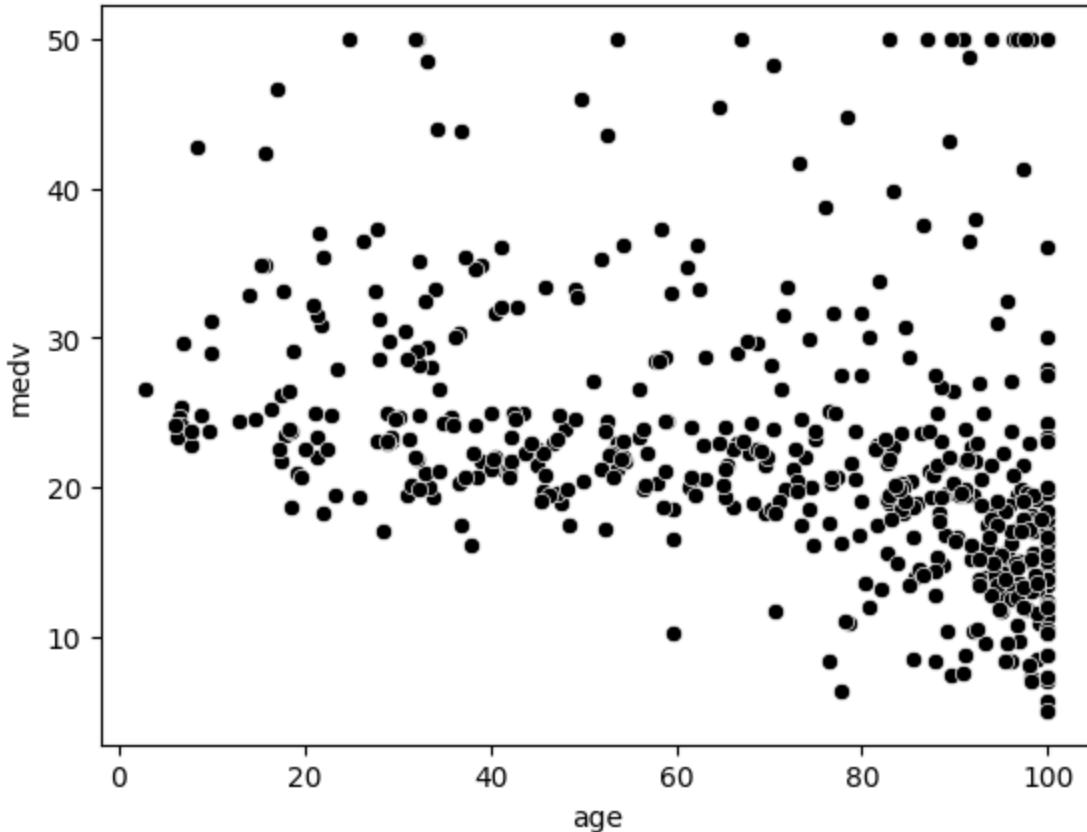
Plot lstat, rm and age against medv

In [7]: `sns.scatterplot(x="lstat", y="medv", data=df, color="black")
plt.show()

sns.scatterplot(x="rm", y="medv", data=df, color="black")
plt.show()

sns.scatterplot(x="age", y="medv", data=df, color="black")
plt.show()`





Linear regression

```
In [8]: regression1 = sm.OLS(df["medv"], sm.add_constant(df["lstat"])).fit()
print(regression1.summary())

regression2 = sm.OLS(df["medv"], sm.add_constant(df["rm"])).fit()
print(regression2.summary())

regression3 = sm.OLS(df["medv"], sm.add_constant(df["age"])).fit()
print(regression3.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.544
Model:	OLS	Adj. R-squared:	0.543
Method:	Least Squares	F-statistic:	601.6
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	5.08e-88
Time:	14:43:14	Log-Likelihood:	-1641.5
No. Observations:	506	AIC:	3287.
Df Residuals:	504	BIC:	3295.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874

Omnibus:	137.043	Durbin-Watson:	0.892
Prob(Omnibus):	0.000	Jarque-Bera (JB):	291.373
Skew:	1.453	Prob(JB):	5.36e-64
Kurtosis:	5.319	Cond. No.	29.7

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.484
Model:	OLS	Adj. R-squared:	0.483
Method:	Least Squares	F-statistic:	471.8
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	2.49e-74
Time:	14:43:14	Log-Likelihood:	-1673.1
No. Observations:	506	AIC:	3350.
Df Residuals:	504	BIC:	3359.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-34.6706	2.650	-13.084	0.000	-39.877	-29.465
rm	9.1021	0.419	21.722	0.000	8.279	9.925

Omnibus:	102.585	Durbin-Watson:	0.684
Prob(Omnibus):	0.000	Jarque-Bera (JB):	612.449
Skew:	0.726	Prob(JB):	1.02e-133
Kurtosis:	8.190	Cond. No.	58.4

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.142
Model:	OLS	Adj. R-squared:	0.140

```

Method: Least Squares F-statistic: 83.48
Date: Sun, 25 Jan 2026 Prob (F-statistic): 1.57e-18
Time: 14:43:14 Log-Likelihood: -1801.5
No. Observations: 506 AIC: 3607.
Df Residuals: 504 BIC: 3615.
Df Model: 1
Covariance Type: nonrobust
=====
            coef    std err      t      P>|t|      [0.025      0.975]
-----
const      30.9787   0.999    31.006    0.000     29.016    32.942
age        -0.1232   0.013    -9.137    0.000    -0.150    -0.097
=====
Omnibus: 170.034 Durbin-Watson: 0.613
Prob(Omnibus): 0.000 Jarque-Bera (JB): 456.983
Skew: 1.671 Prob(JB): 5.85e-100
Kurtosis: 6.240 Cond. No. 195.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation of regressions

lstat

For `lstat` gives R-squared value of (0.544), which means that approximately 54.4% of the variance can be explained by the model. The prod F-statistic (5.08e-88) indicates that the model is statistically significant due to the value < 0.05. The negative coefficient (-0.9500) indicates that as `lstat` increases, `medv` tends to decrease, suggesting an inverse relationship between these two variables.

`rm`

The same interpretation can be made for `rm`, which has an R-squared value of (0.484), prod F-statistic value of (2.49e-74) and a coefficient of (9.1021). Here the positive coefficient indicates that as `rm` increases, `medv` also tends to increase, suggesting a direct relationship between these two variables.

`age`

Same interpretation can be made for `age`, which has an R-squared value of (0.142), prod F-statistic value of (1.57e-18) and a coefficient of (-0.1232). Here the negative coefficient indicates that as `age` increases, `medv` tends to decrease, suggesting an inverse relationship between these two variables.

```
In [9]: print(regression1.conf_int())
print(regression2.conf_int())
print(regression3.conf_int())
```

```

          0      1
const  33.448457  35.659225
lstat -1.026148 -0.873951
          0      1
const -39.876641 -29.464601
rm     8.278855  9.925363
          0      1
const 29.015752  32.941604
age   -0.149647 -0.096679

```

Indicates the lower and upper bounds of the 95% confidence interval. First row (y intercept) and second row (slope of predictor). Smaller intervals in the slope indicates that the model is more precise.

The second one indicates highly positive correlation but it has a greater interval. Meanwhile the first and third one indicates negative correlation with smaller intervals.

Use model

```
In [10]: use_lstat = pd.DataFrame({"lstat": [5, 10, 15]})
use_lstat = sm.add_constant(use_lstat)
predictor1 = regression1.get_prediction(use_lstat).summary_frame(alpha=0.05)
print(predictor1[["mean", "obs_ci_lower", "obs_ci_upper"]])

use_rm = pd.DataFrame({"rm": [5, 6.5, 8]})
use_rm = sm.add_constant(use_rm)
predictor2 = regression2.get_prediction(use_rm).summary_frame(alpha=0.05)
print(predictor2[["mean", "obs_ci_lower", "obs_ci_upper"]])

use_age = pd.DataFrame({"age": [25, 50, 75]})
use_age = sm.add_constant(use_age)
predictor3 = regression3.get_prediction(use_age).summary_frame(alpha=0.05)
print(predictor3[["mean", "obs_ci_lower", "obs_ci_upper"]])
```

	mean	obs_ci_lower	obs_ci_upper
0	29.803594	17.565675	42.041513
1	25.053347	12.827626	37.279068
2	20.303101	8.077742	32.528459
	mean	obs_ci_lower	obs_ci_upper
0	10.839924	-2.214474	23.894322
1	24.493088	11.480391	37.505784
2	38.146251	25.058353	51.234149
	mean	obs_ci_lower	obs_ci_upper
0	27.899610	11.090368	44.708852
1	24.820542	8.043748	41.597335
2	21.741474	4.971031	38.511917

Interpretation of results

lstat

Inserted values for lstat where 5, 10, 15 this means that with 95% confidence the predicted medv values will respectively be between (17.56, 42.04), (12.82, 37.27) and (8.07,

32.52) approximatly.

rm & age

Same interpretation can be made for rm and age where inserted values are 5, 6.5, 8 for rm and 25, 50, 75 for age .

```
In [11]: regression = sm.OLS(df["medv"], sm.add_constant(df[["lstat", "rm", "age"]])).fit()
print(regression.summary())
```

OLS Regression Results									
Dep. Variable:	medv	R-squared:	0.639						
Model:	OLS	Adj. R-squared:	0.637						
Method:	Least Squares	F-statistic:	296.2						
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	1.20e-110						
Time:	14:43:14	Log-Likelihood:	-1582.4						
No. Observations:	506	AIC:	3173.						
Df Residuals:	502	BIC:	3190.						
Df Model:	3								
Covariance Type:	nonrobust								
	coef	std err	t	P> t	[0.025	0.975]			

const	-1.1753	3.182	-0.369	0.712	-7.427	5.076			
lstat	-0.6685	0.054	-12.298	0.000	-0.775	-0.562			
rm	5.0191	0.454	11.048	0.000	4.127	5.912			
age	0.0091	0.011	0.811	0.418	-0.013	0.031			

Omnibus:		138.819	Durbin-Watson:			0.851			
Prob(Omnibus):		0.000	Jarque-Bera (JB):			415.436			
Skew:		1.296	Prob(JB):			6.15e-91			
Kurtosis:		6.603	Cond. No.			985.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Here the R-squared value (0.639) indicates that approximately 63.9% of the variance in medv can be explained by the model. The prod F-statistic (1.20e-110) indicates that the model is statistically significant due to the value < 0.05.

```
In [12]: regression = sm.OLS(df["medv"], sm.add_constant(df.drop(columns=["medv"]))).fit()
print(regression.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.741			
Model:	OLS	Adj. R-squared:	0.734			
Method:	Least Squares	F-statistic:	108.1			
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	6.72e-135			
Time:	14:43:14	Log-Likelihood:	-1498.8			
No. Observations:	506	AIC:	3026.			
Df Residuals:	492	BIC:	3085.			
Df Model:	13					
Covariance Type:	nonrobust					
coef	std err	t	P> t			
[0.025	0.975]					
const	36.4595	5.103	7.144	0.000	26.432	46.487
crim	-0.1080	0.033	-3.287	0.001	-0.173	-0.043
zn	0.0464	0.014	3.382	0.001	0.019	0.073
indus	0.0206	0.061	0.334	0.738	-0.100	0.141
chas	2.6867	0.862	3.118	0.002	0.994	4.380
nox	-17.7666	3.820	-4.651	0.000	-25.272	-10.262
rm	3.8099	0.418	9.116	0.000	2.989	4.631
age	0.0007	0.013	0.052	0.958	-0.025	0.027
dis	-1.4756	0.199	-7.398	0.000	-1.867	-1.084
rad	0.3060	0.066	4.613	0.000	0.176	0.436
tax	-0.0123	0.004	-3.280	0.001	-0.020	-0.005
ptratio	-0.9527	0.131	-7.283	0.000	-1.210	-0.696
black	0.0093	0.003	3.467	0.001	0.004	0.015
lstat	-0.5248	0.051	-10.347	0.000	-0.624	-0.425
Omnibus:	178.041	Durbin-Watson:	1.078			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	783.126			
Skew:	1.521	Prob(JB):	8.84e-171			
Kurtosis:	8.281	Cond. No.	1.51e+04			

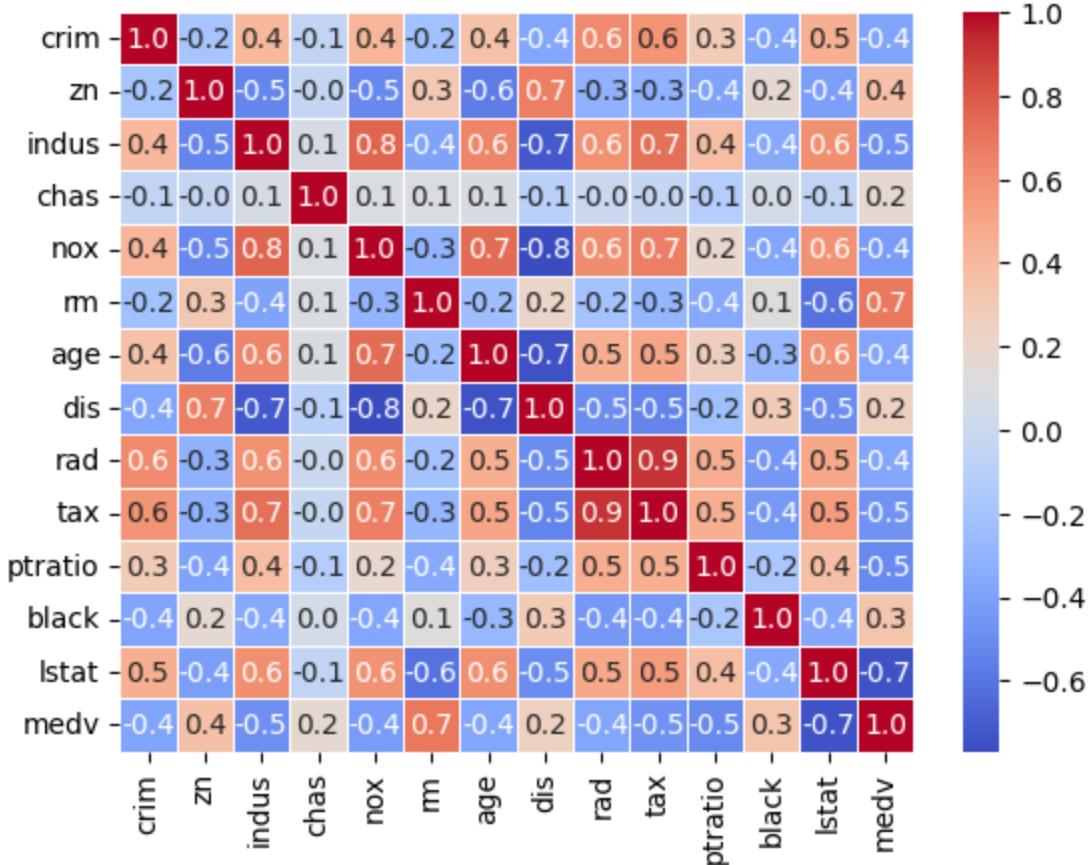
Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.51e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Here the R-squared value (0.741) indicates that approximately 74.1% of the variance in medv can be explained by the model. The prod F-statistic (6.72e-135) indicates that the model is statistically significant due to the value < 0.05.

Correlation matrix

```
In [13]: sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".1f", linewidths=0.5)
plt.show()
```



This matrix shows the correlation coefficients between each pair of variables in the dataset. A correlation coefficient close to 1 indicates a strong positive correlation, while a coefficient close to -1 indicates a strong negative correlation. A coefficient around 0 suggests no correlation between the variables.

Example interpretations:

`crim` and `zn` have a correlation coefficient of -0.2, indicating a weak negative correlation. This suggests that as the value of `crim` increases, the value of `zn` tends to decrease slightly.

Use multiple linear regression model

```
In [14]: selected_predictor_values = pd.DataFrame(
    pd.MultiIndex.from_product(
        [[5, 10, 15], [5, 6.5, 8]], names=["lstat", "rm"])
    .to_frame(index=False)
)
print(selected_predictor_values)

regression = sm.OLS(df["medv"], sm.add_constant(df[["lstat", "rm"]])).fit()
selected_predictor_values = sm.add_constant(selected_predictor_values)
predictions = regression.get_prediction(selected_predictor_values)
pred_summary = predictions.summary_frame(alpha=0.05)
```

```
print(pred_summary[["mean", "obs_ci_lower", "obs_ci_upper"]])
```

	lstat	rm
0	5	5.0
1	5	6.5
2	5	8.0
3	10	5.0
4	10	6.5
5	10	8.0
6	15	5.0
7	15	6.5
8	15	8.0

	mean	obs_ci_lower	obs_ci_upper
0	20.903875	9.889729	31.918021
1	28.546057	17.635923	39.456192
2	36.188239	25.225479	47.150999
3	17.692084	6.722152	28.662016
4	25.334266	14.437027	36.231505
5	32.976448	21.995024	43.957872
6	14.480292	3.537875	25.422709
7	22.122474	11.221204	33.023745
8	29.764656	18.747835	40.781477

Interpretation of results

Row one indicates that for value of `lstat` (5) and `rm` (5.0) will with 95% confidence result in a `medv` value between (9.88, 31.91), meanwhile row nine indicates that for value of `lstat` (15) and `rm` (8.0) will with 95% confidence result in a `medv` value between (18.74, 40.78). Same goes for all the subsequent rows with different values for `lstat` and `rm` which also provides a new boundary for all combinations of them.