



Assignment 2

Simple and Multiple Linear Regression

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-25

Course: Machine Learning 4DT905

Imports

```
In [1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

Load data

```
In [2]: df = pd.read_csv("../data/Boston.csv", index_col=0)
```

Number of predictors and names

```
In [3]: df_names = df.columns.tolist()
print(f"Number of columns: {len(df_names)}")
print(f"Column names: {df_names}")
```

```
Number of columns: 14
Column names: ['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad',
'tax', 'ptratio', 'black', 'lstat', 'medv']
```

Statistical summary of predictors

```
In [4]: df.describe()
```

Out[4]:

	crim	zn	indus	chas	nox	rm	tax	dis	rad	ptratio	medv
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	15.3	4.0900	1	296	15.3
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	17.8	4.9671	2	242	17.8
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	13.0	4.0000	1	222	13.0
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	15.0	4.0000	2	222	15.0
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	17.0	4.9671	2	242	17.0
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	18.0	5.0000	3	222	18.0
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	18.7	5.0000	3	222	18.7

Number of datapoints

```
In [5]: print(f"Number of datapoints: {len(df)}")
```

Number of datapoints: 506

Display data in table format

```
In [6]: print(df.head(5))
```

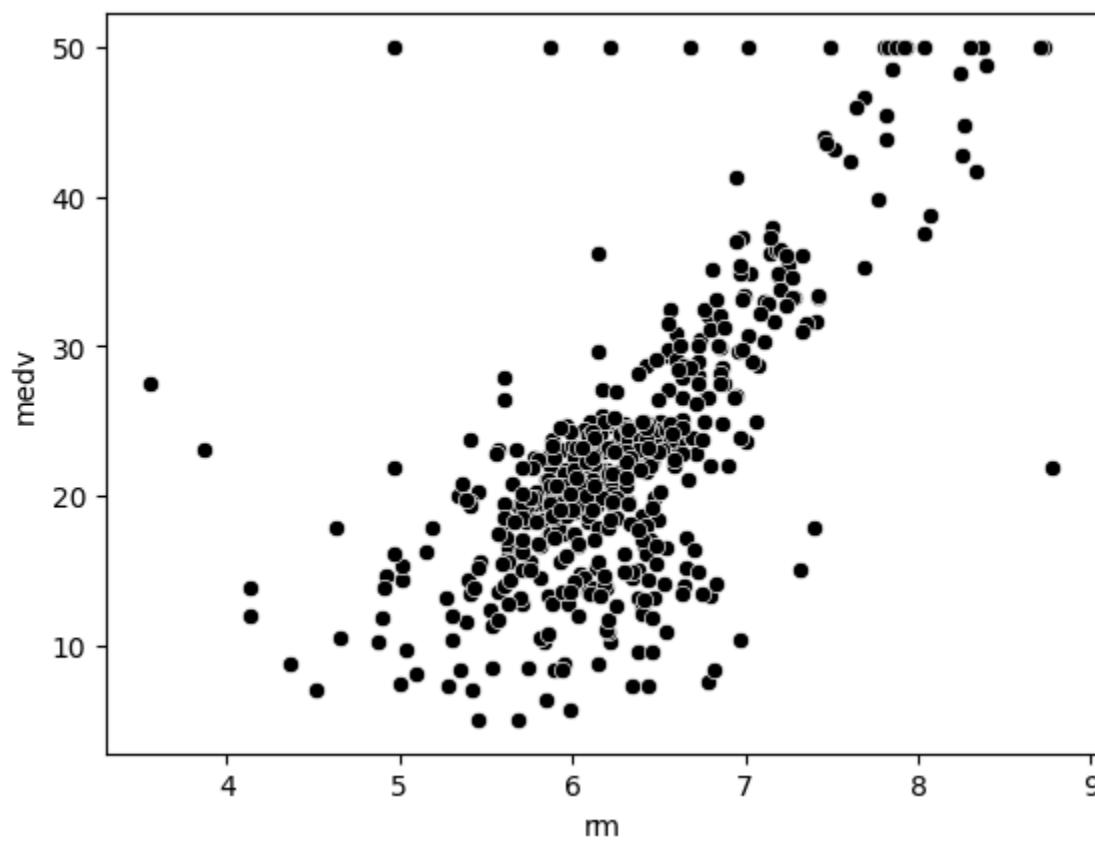
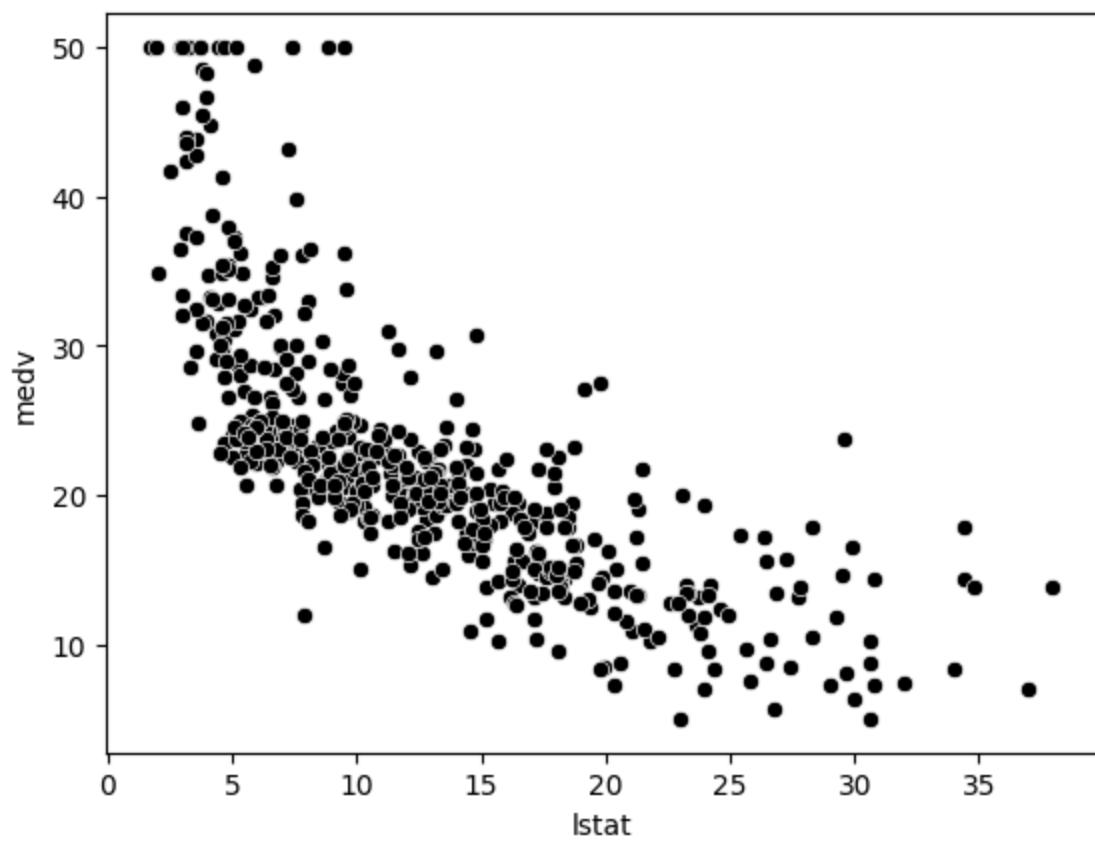
	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	medv
1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	12.6
2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	15.3
3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	15.3
4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	15.3
5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	15.3
	black	lstat	medv									
1	396.90	4.98	24.0									
2	396.90	9.14	21.6									
3	392.83	4.03	34.7									
4	394.63	2.94	33.4									
5	396.90	5.33	36.2									

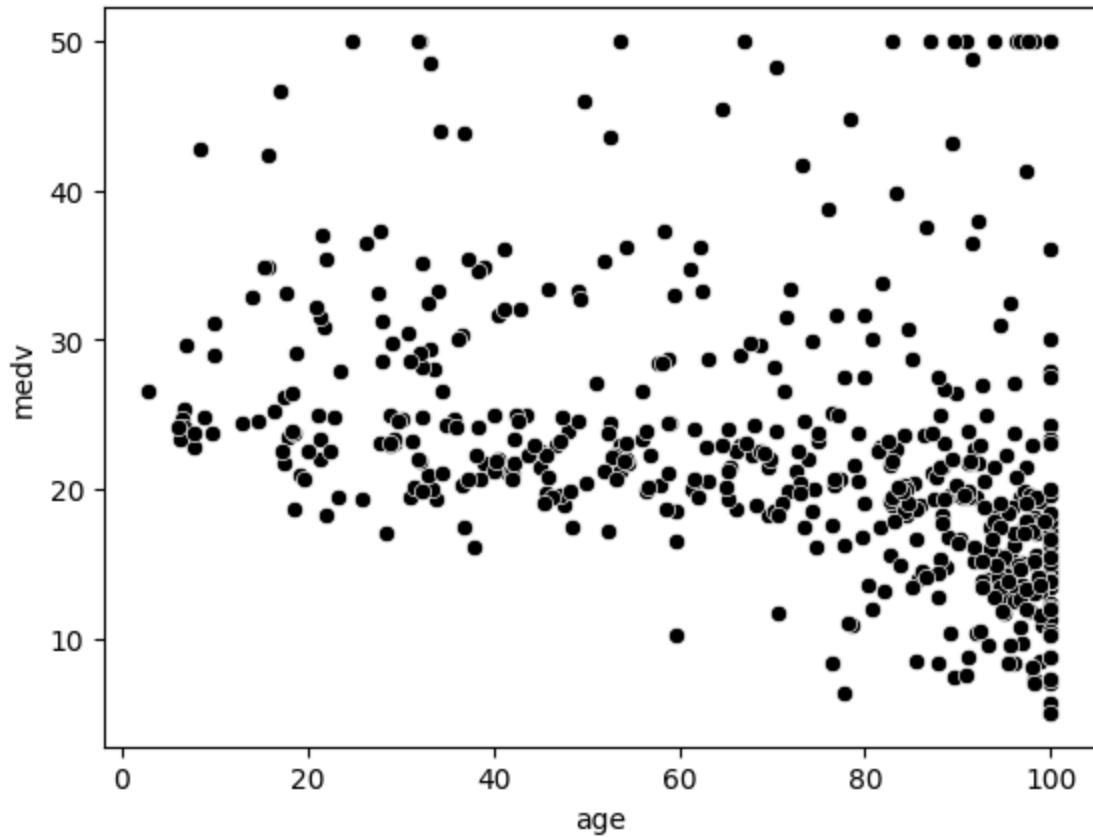
Plot lstat, rm and age against medv

```
In [7]: sns.scatterplot(x="lstat", y="medv", data=df, color="black")
plt.show()

sns.scatterplot(x="rm", y="medv", data=df, color="black")
plt.show()

sns.scatterplot(x="age", y="medv", data=df, color="black")
plt.show()
```





Linear regression

```
In [8]: regression1 = sm.OLS(df["medv"], sm.add_constant(df["lstat"])).fit()
print(regression1.summary())

regression2 = sm.OLS(df["medv"], sm.add_constant(df["rm"])).fit()
print(regression2.summary())

regression3 = sm.OLS(df["medv"], sm.add_constant(df["age"])).fit()
print(regression3.summary())
```

OLS Regression Results

```
=====
Dep. Variable:                  medv   R-squared:                 0.544
Model:                          OLS    Adj. R-squared:            0.543
Method:                         Least Squares   F-statistic:              601.6
Date:                          Sun, 25 Jan 2026   Prob (F-statistic):      5.08e-88
Time:                           14:43:14     Log-Likelihood:          -1641.5
No. Observations:                506    AIC:                      3287.
Df Residuals:                   504    BIC:                      3295.
Df Model:                        1
Covariance Type:                nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874

```
=====
Omnibus:                     137.043   Durbin-Watson:           0.892
Prob(Omnibus):                0.000    Jarque-Bera (JB):       291.373
Skew:                          1.453    Prob(JB):                 5.36e-64
Kurtosis:                      5.319   Cond. No.                  29.7
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

```
=====
Dep. Variable:                  medv   R-squared:                 0.484
Model:                          OLS    Adj. R-squared:            0.483
Method:                         Least Squares   F-statistic:              471.8
Date:                          Sun, 25 Jan 2026   Prob (F-statistic):      2.49e-74
Time:                           14:43:14     Log-Likelihood:          -1673.1
No. Observations:                506    AIC:                      3350.
Df Residuals:                   504    BIC:                      3359.
Df Model:                        1
Covariance Type:                nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-34.6706	2.650	-13.084	0.000	-39.877	-29.465
rm	9.1021	0.419	21.722	0.000	8.279	9.925

```
=====
Omnibus:                     102.585   Durbin-Watson:           0.684
Prob(Omnibus):                0.000    Jarque-Bera (JB):       612.449
Skew:                          0.726    Prob(JB):                 1.02e-133
Kurtosis:                      8.190   Cond. No.                  58.4
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.142			
Model:	OLS	Adj. R-squared:	0.140			
Method:	Least Squares	F-statistic:	83.48			
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	1.57e-18			
Time:	14:43:14	Log-Likelihood:	-1801.5			
No. Observations:	506	AIC:	3607.			
Df Residuals:	504	BIC:	3615.			
Df Model:	1					
Covariance Type:	nonrobust					
<hr/>						
	coef	std err	t	P> t	[0.025	0.975]
const	30.9787	0.999	31.006	0.000	29.016	32.942
age	-0.1232	0.013	-9.137	0.000	-0.150	-0.097
<hr/>						
Omnibus:	170.034	Durbin-Watson:	0.613			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	456.983			
Skew:	1.671	Prob(JB):	5.85e-100			
Kurtosis:	6.240	Cond. No.	195.			
<hr/>						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation of regressions

lstat

For `lstat` gives R-squared value of (0.544), which means that approximately 54.4% of the variance in `medv` can be explained by the model. The F-statistic p-value (`Prob(F-statistic) = 5.08e-88`) is less than 0.05, leading us to reject the null hypothesis that all coefficients are zero. This indicates the model is statistically significant in predicting `medv`. The negative coefficient (-0.9500) indicates that as `lstat` increases, `medv` tends to decrease, suggesting an inverse relationship between these two variables.

rm

The same interpretation can be made for `rm`, which has an R-squared value of (0.484), F-statistic p-value of (2.49e-74), and a coefficient of (9.1021). The p-value < 0.05 indicates the model is statistically significant. The positive coefficient indicates that as `rm` increases, `medv` also tends to increase, suggesting a direct relationship between these two variables.

age

Same interpretation can be made for `age`, which has an R-squared value of (0.142), F-statistic p-value of (1.57e-18), and a coefficient of (-0.1232). The p-value < 0.05 indicates the model is statistically significant. The negative coefficient

indicates that as `age` increases, `medv` tends to decrease, suggesting an inverse relationship between these two variables.

```
In [9]: print(regression1.conf_int())
print(regression2.conf_int())
print(regression3.conf_int())
```

	0	1
const	33.448457	35.659225
lstat	-1.026148	-0.873951
	0	1
const	-39.876641	-29.464601
rm	8.278855	9.925363
	0	1
const	29.015752	32.941604
age	-0.149647	-0.096679

Indicates the lower and upper bounds of the 95% confidence interval. First row (y intercept) and second row (slope of predictor). Smaller intervals in the slope indicates that the model is more precise.

The second model (`rm`) shows a positive relationship (positive coefficient) with a wider confidence interval. Meanwhile the first (`lstat`) and third (`age`) models show negative relationships (negative coefficients) with narrower confidence intervals.

Use model

```
In [10]: use_lstat = pd.DataFrame({"lstat": [5, 10, 15]})
use_lstat = sm.add_constant(use_lstat)
predictor1 = regression1.get_prediction(use_lstat).summary_frame(alpha=0.05)
print(predictor1[["mean", "obs_ci_lower", "obs_ci_upper"]])

use_rm = pd.DataFrame({"rm": [5, 6.5, 8]})
use_rm = sm.add_constant(use_rm)
predictor2 = regression2.get_prediction(use_rm).summary_frame(alpha=0.05)
print(predictor2[["mean", "obs_ci_lower", "obs_ci_upper"]])

use_age = pd.DataFrame({"age": [25, 50, 75]})
use_age = sm.add_constant(use_age)
predictor3 = regression3.get_prediction(use_age).summary_frame(alpha=0.05)
print(predictor3[["mean", "obs_ci_lower", "obs_ci_upper"]])
```

	mean	obs_ci_lower	obs_ci_upper
0	29.803594	17.565675	42.041513
1	25.053347	12.827626	37.279068
2	20.303101	8.077742	32.528459
	mean	obs_ci_lower	obs_ci_upper
0	10.839924	-2.214474	23.894322
1	24.493088	11.480391	37.505784
2	38.146251	25.058353	51.234149
	mean	obs_ci_lower	obs_ci_upper
0	27.899610	11.090368	44.708852
1	24.820542	8.043748	41.597335
2	21.741474	4.971031	38.511917

Interpretation of results

lstat

Inserted values for `lstat` where 5, 10, 15 this means that with 95% confidence the predicted `medv` values will respectively be between (17.56, 42.04), (12.82, 37.27) and (8.07, 32.52) approximately.

rm & age

Same interpretation can be made for `rm` and `age` where inserted values are 5, 6.5, 8 for `rm` and 25, 50, 75 for `age`.

```
In [11]: regression = sm.OLS(df["medv"], sm.add_constant(df[["lstat", "rm", "age"]])).f
print(regression.summary())
```

```

OLS Regression Results
=====
Dep. Variable: medv R-squared: 0.639
Model: OLS Adj. R-squared: 0.637
Method: Least Squares F-statistic: 296.2
Date: Sun, 25 Jan 2026 Prob (F-statistic): 1.20e-110
Time: 14:43:14 Log-Likelihood: -1582.4
No. Observations: 506 AIC: 3173.
Df Residuals: 502 BIC: 3190.
Df Model: 3
Covariance Type: nonrobust
=====
              coef    std err      t      P>|t|      [0.025      0.975]
-----
const      -1.1753   3.182    -0.369     0.712    -7.427     5.076
lstat      -0.6685   0.054   -12.298     0.000    -0.775    -0.562
rm         5.0191   0.454    11.048     0.000     4.127     5.912
age        0.0091   0.011     0.811     0.418    -0.013     0.031
=====
Omnibus: 138.819 Durbin-Watson: 0.851
Prob(Omnibus): 0.000 Jarque-Bera (JB): 415.436
Skew: 1.296 Prob(JB): 6.15e-91
Kurtosis: 6.603 Cond. No. 985.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Here the R-squared value (0.639) indicates that approximately 63.9% of the variance in `medv` can be explained by the model. The F-statistic p-value (`Prob(F-statistic)`) = `1.20e-110` is less than 0.05, indicating we reject the null hypothesis that all coefficients are zero. This means the model is statistically significant in predicting `medv`.

```
In [12]: regression = sm.OLS(df["medv"], sm.add_constant(df.drop(columns=["medv"]))).fit()
print(regression.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.741			
Model:	OLS	Adj. R-squared:	0.734			
Method:	Least Squares	F-statistic:	108.1			
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	6.72e-135			
Time:	14:43:14	Log-Likelihood:	-1498.8			
No. Observations:	506	AIC:	3026.			
Df Residuals:	492	BIC:	3085.			
Df Model:	13					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	36.4595	5.103	7.144	0.000	26.432	46.487
crim	-0.1080	0.033	-3.287	0.001	-0.173	-0.043
zn	0.0464	0.014	3.382	0.001	0.019	0.073
indus	0.0206	0.061	0.334	0.738	-0.100	0.141
chas	2.6867	0.862	3.118	0.002	0.994	4.380
nox	-17.7666	3.820	-4.651	0.000	-25.272	-10.262
rm	3.8099	0.418	9.116	0.000	2.989	4.631
age	0.0007	0.013	0.052	0.958	-0.025	0.027
dis	-1.4756	0.199	-7.398	0.000	-1.867	-1.084
rad	0.3060	0.066	4.613	0.000	0.176	0.436
tax	-0.0123	0.004	-3.280	0.001	-0.020	-0.005
ptratio	-0.9527	0.131	-7.283	0.000	-1.210	-0.696
black	0.0093	0.003	3.467	0.001	0.004	0.015
lstat	-0.5248	0.051	-10.347	0.000	-0.624	-0.425
Omnibus:	178.041	Durbin-Watson:	1.078			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	783.126			
Skew:	1.521	Prob(JB):	8.84e-171			
Kurtosis:	8.281	Cond. No.	1.51e+04			

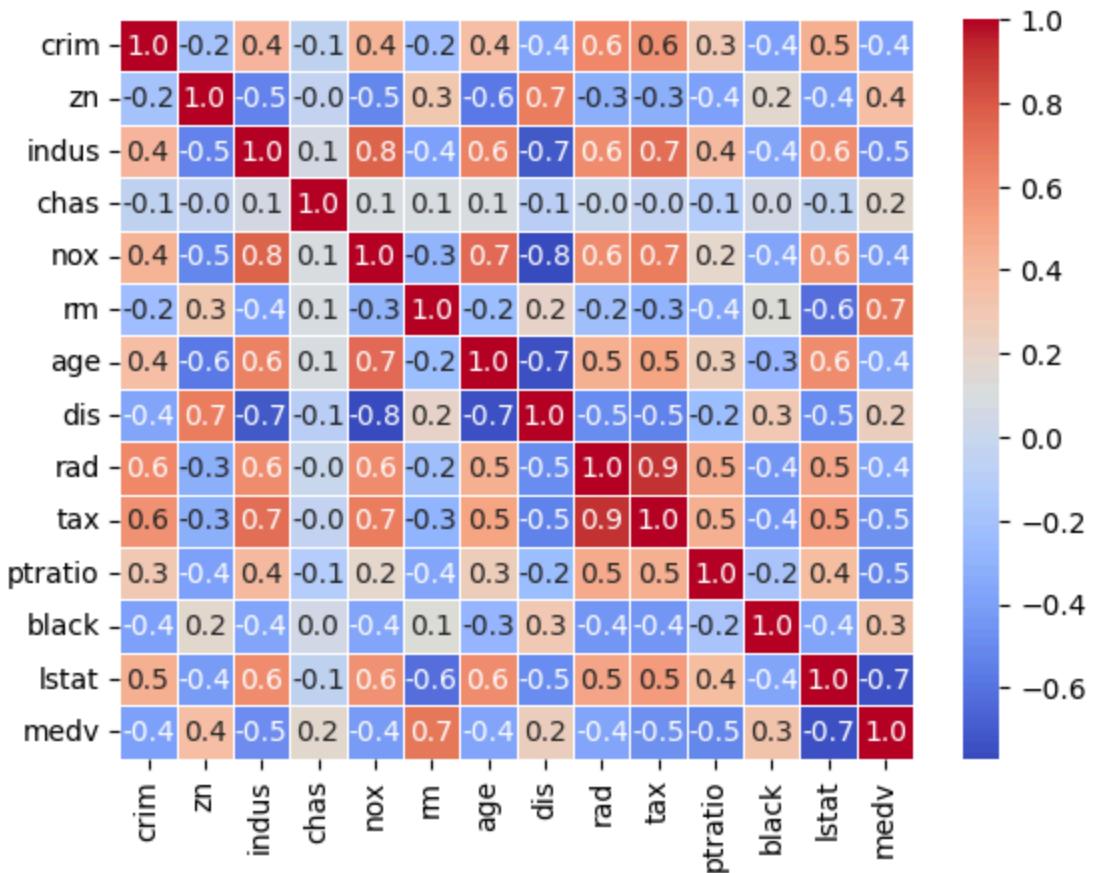
Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.51e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Here the R-squared value (0.741) indicates that approximately 74.1% of the variance in `medv` can be explained by the model. The F-statistic p-value (`Prob(F-statistic)`) = 6.72e-135) is less than 0.05, indicating we reject the null hypothesis that all coefficients are zero. This means the model is statistically significant in predicting `medv`.

Correlation matrix

```
In [13]: sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".1f", linewidths=0.5)
plt.show()
```



This matrix shows the correlation coefficients between each pair of variables in the dataset. A correlation coefficient close to 1 indicates a strong positive correlation, while a coefficient close to -1 indicates a strong negative correlation. A coefficient around 0 suggests no correlation between the variables.

Example interpretations:

`crim` and `zn` have a correlation coefficient of -0.2, indicating a weak negative correlation. This suggests that as the value of `crim` increases, the value of `zn` tends to decrease slightly.

Use multiple linear regression model

```
In [14]: selected_predictor_values = pd.DataFrame(
    pd.MultiIndex.from_product(
        [[5, 10, 15], [5, 6.5, 8]], names=["lstat", "rm"])
    .to_frame(index=False)
)
print(selected_predictor_values)

regression = sm.OLS(df["medv"], sm.add_constant(df[["lstat", "rm"]])).fit()
selected_predictor_values = sm.add_constant(selected_predictor_values)
predictions = regression.get_prediction(selected_predictor_values)
```

```

pred_summary = predictions.summary_frame(alpha=0.05)

print(pred_summary[["mean", "obs_ci_lower", "obs_ci_upper"]])

```

	lstat	rm	
0	5	5.0	
1	5	6.5	
2	5	8.0	
3	10	5.0	
4	10	6.5	
5	10	8.0	
6	15	5.0	
7	15	6.5	
8	15	8.0	
	mean	obs_ci_lower	obs_ci_upper
0	20.903875	9.889729	31.918021
1	28.546057	17.635923	39.456192
2	36.188239	25.225479	47.150999
3	17.692084	6.722152	28.662016
4	25.334266	14.437027	36.231505
5	32.976448	21.995024	43.957872
6	14.480292	3.537875	25.422709
7	22.122474	11.221204	33.023745
8	29.764656	18.747835	40.781477

Interpretation of results

Row one indicates that for value of `lstat` (5) and `rm` (5.0) will with 95% confidence result in a `medv` value between (9.88, 31.91), meanwhile row nine indicates that for value of `lstat` (15) and `rm` (8.0) will with 95% confidence result in a `medv` value between (18.74, 40.78). Same goes for all the subsequent rows with different values for `lstat` and `rm` which also provides a new boundary for all combinations of them.