

January 26, 2026

## 1 Assignment 3

### 1.1 Simple and Multiple Linear Regression pt2

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-25

Course: Machine Learning 4DT905

### 1.2 Conceptual

$$Y = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2 + \hat{\beta}_3 X_3 + \hat{\beta}_4 X_1 X_2 + \hat{\beta}_5 X_1 X_3$$

$$\hat{\beta}_0 = 50(\text{Intercept})$$

$$\hat{\beta}_1 = 20(\text{GPA})$$

$$\hat{\beta}_2 = 0.07(\text{IQ})$$

$$\hat{\beta}_3 = 35(\text{Level})$$

$$\hat{\beta}_4 = 0.01(\text{GPA} \cdot \text{IQ})$$

$$\hat{\beta}_5 = -10(\text{GPA} \cdot \text{Level})$$

$X_3 = 1$  for College, 0 for High School

1.

$$Y_c = 50 + 20X_1 + 0.07X_2 + 35 + 0.01X_1X_2 - 10X_1$$

$$Y_h = 50 + 20X_1 + 0.07X_2 + 0.01X_1X_2$$

$$Y_c - Y_h = 35 - 10X_1$$

$$35 - 10X_1 = 0 \implies X_1 = 3.5$$

True, when GPA > 3.5 High School graduates earn more than College graduates.

Answer:  *iii*

2.

$$X_1 = 4.0$$

$$X_2 = 110$$

$$X_3 = 1$$

$$Y = 50 + 20(4.0) + 0.07(110) + 35 + 0.01(4.0)(110) - 10(4.0)$$

$$Y = 137.1$$

Answer:  \$137, 100

3. False. The magnitude of a coefficient does not indicate statistical importance. To determine statistical importance we need to look at the p-values associated with that coefficient, not just its absolute value. In the presented case, the units of predictor  $X_2$  (IQ) are generally > 100. A small coefficient for the  $X_2 \cdot X_1$  term might still result in a large contribution to the model and be highly statistically significant.

Answer:  *False*

## 1.3 Practical

### 1.3.1 Imports

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

```
import numpy as np
```

### 1.3.2 Assignment 2

#### 1.3.3 Load data

```
[2]: df = pd.read_csv("../data/Boston.csv", index_col=0)
```

#### 1.3.4 Number of predictors and names

```
[3]: df_names = df.columns.tolist()
print(f"Number of columns: {len(df_names)}")
print(f"Column names: {df_names}")
```

Number of columns: 14

Column names: ['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax', 'ptratio', 'black', 'lstat', 'medv']

#### 1.3.5 Statistical summary of predictors

```
[4]: df.describe()
```

```
[4]:          crim            zn           indus          chas            nox            rm \
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean   3.613524  11.363636  11.136779  0.069170  0.554695  6.284634
std    8.601545  23.322453  6.860353  0.253994  0.115878  0.702617
min    0.006320  0.000000  0.460000  0.000000  0.385000  3.561000
25%    0.082045  0.000000  5.190000  0.000000  0.449000  5.885500
50%    0.256510  0.000000  9.690000  0.000000  0.538000  6.208500
75%    3.677083  12.500000 18.100000  0.000000  0.624000  6.623500
max    88.976200 100.000000 27.740000  1.000000  0.871000  8.780000

              age            dis           rad            tax          ptratio            black \
count  506.000000  506.000000  506.000000  506.000000  506.000000  506.000000
mean   68.574901  3.795043  9.549407  408.237154  18.455534  356.674032
std    28.148861  2.105710  8.707259 168.537116  2.164946  91.294864
min    2.900000  1.129600  1.000000 187.000000  12.600000  0.320000
25%    45.025000  2.100175  4.000000 279.000000  17.400000  375.377500
50%    77.500000  3.207450  5.000000 330.000000  19.050000  391.440000
75%    94.075000  5.188425 24.000000 666.000000  20.200000  396.225000
max   100.000000 12.126500 24.000000 711.000000  22.000000  396.900000

          lstat            medv
count  506.000000  506.000000
mean   12.653063  22.532806
std    7.141062  9.197104
min    1.730000  5.000000
25%    6.950000 17.025000
```

```
50%      11.360000  21.200000
75%      16.955000  25.000000
max      37.970000  50.000000
```

### 1.3.6 Number of datapoints

```
[5]: print(f"Number of datapoints: {len(df)}")
```

```
Number of datapoints: 506
```

### 1.3.7 Display data in table format

```
[6]: print(df.head(5))
```

```
      crim    zn  indus  chas    nox     rm    age     dis   rad   tax  ptratio \
1  0.00632  18.0   2.31      0  0.538  6.575  65.2  4.0900    1  296   15.3
2  0.02731    0.0   7.07      0  0.469  6.421  78.9  4.9671    2  242   17.8
3  0.02729    0.0   7.07      0  0.469  7.185  61.1  4.9671    2  242   17.8
4  0.03237    0.0   2.18      0  0.458  6.998  45.8  6.0622    3  222   18.7
5  0.06905    0.0   2.18      0  0.458  7.147  54.2  6.0622    3  222   18.7

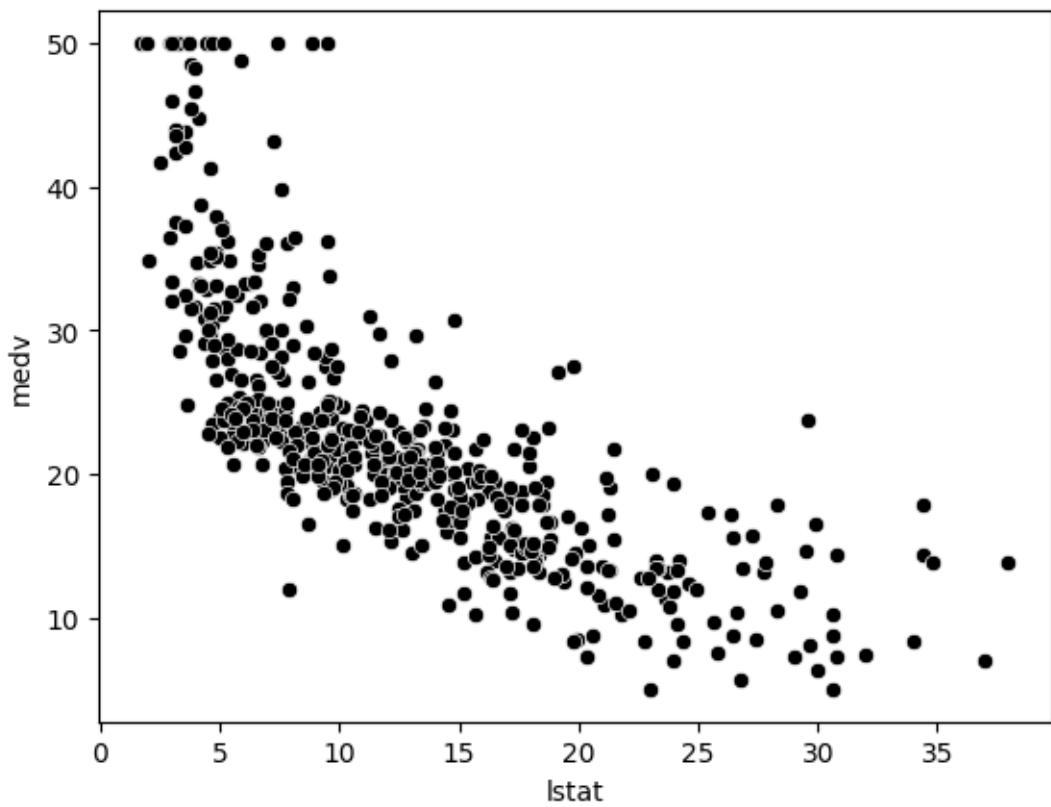
      black  lstat  medv
1  396.90   4.98  24.0
2  396.90   9.14  21.6
3  392.83   4.03  34.7
4  394.63   2.94  33.4
5  396.90   5.33  36.2
```

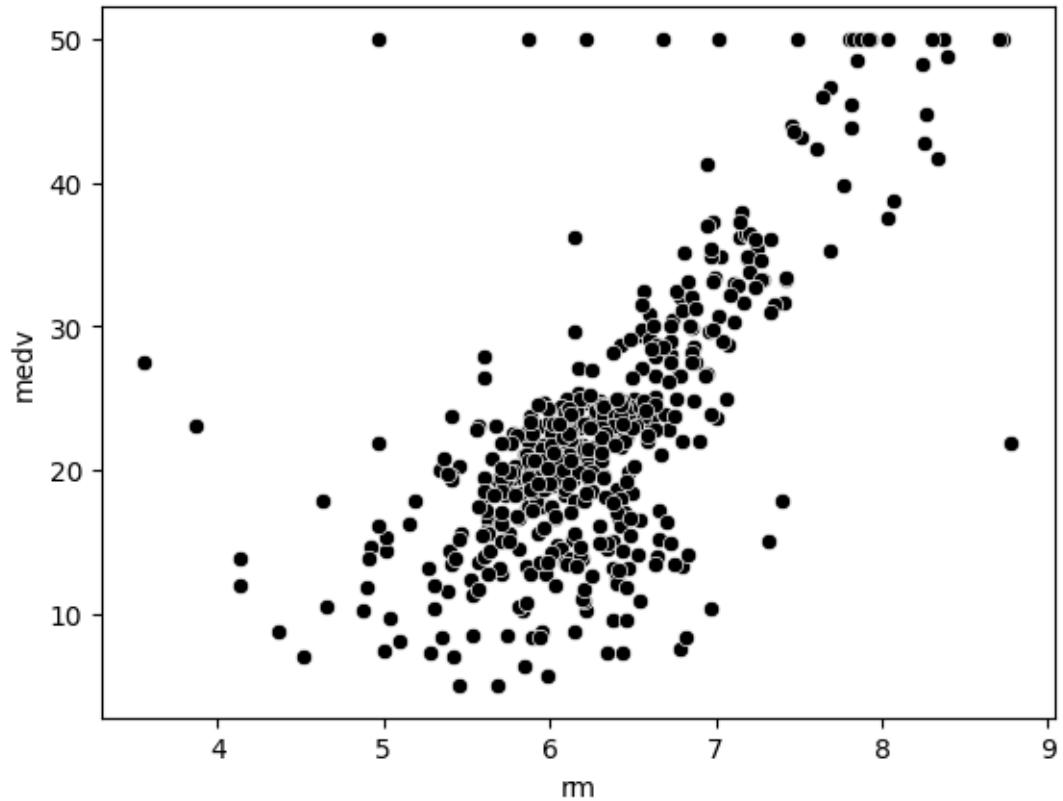
### 1.3.8 Plot lstat, rm and age against medv

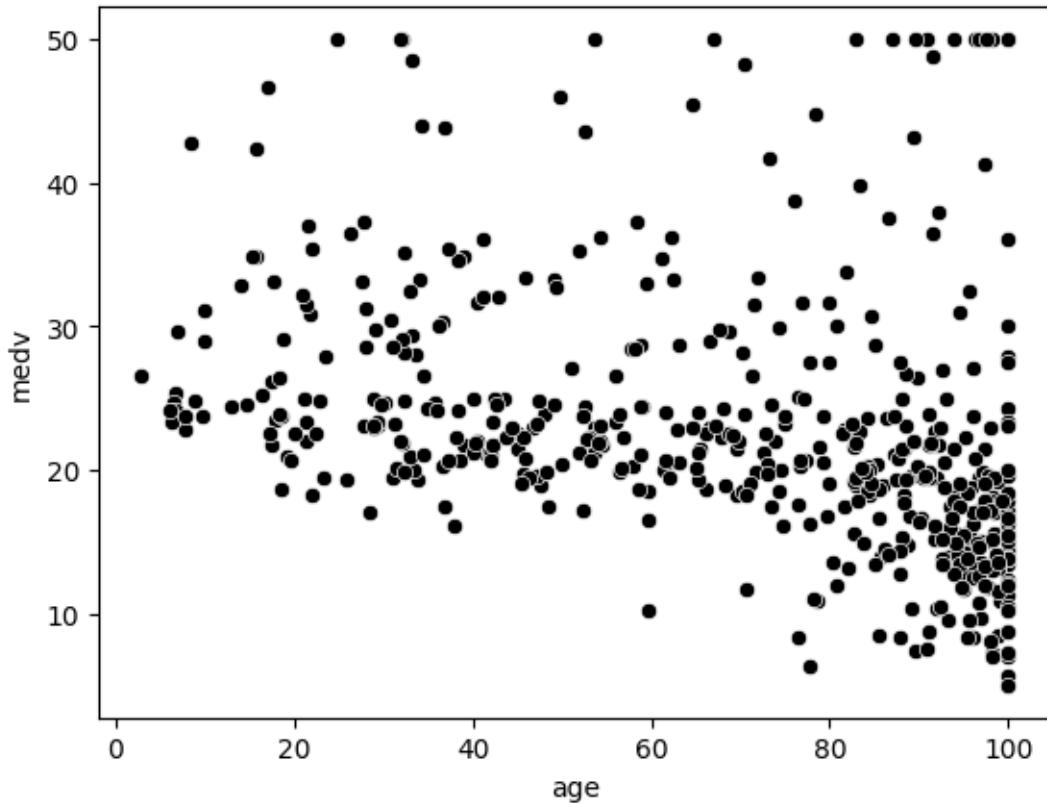
```
[7]: sns.scatterplot(x="lstat", y="medv", data=df, color="black")
plt.show()

sns.scatterplot(x="rm", y="medv", data=df, color="black")
plt.show()

sns.scatterplot(x="age", y="medv", data=df, color="black")
plt.show()
```







### 1.3.9 Linear regression

```
[8]: regression1 = sm.OLS(df["medv"], sm.add_constant(df["lstat"])).fit()
print(regression1.summary())

regression2 = sm.OLS(df["medv"], sm.add_constant(df["rm"])).fit()
print(regression2.summary())

regression3 = sm.OLS(df["medv"], sm.add_constant(df["age"])).fit()
print(regression3.summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:                  medv      R-squared:                 0.544
Model:                          OLS       Adj. R-squared:            0.543
Method:                         Least Squares   F-statistic:             601.6
Date:                          Sun, 25 Jan 2026   Prob (F-statistic):      5.08e-88
Time:                           16:43:09      Log-Likelihood:          -1641.5
No. Observations:                  506      AIC:                      3287.
Df Residuals:                      504      BIC:                      3295.
Df Model:                           1
```

Covariance Type: nonrobust						
	coef	std err	t	P> t	[0.025	0.975]
const	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874
<hr/>						
Omnibus:	137.043	Durbin-Watson:	0.892			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	291.373			
Skew:	1.453	Prob(JB):	5.36e-64			
Kurtosis:	5.319	Cond. No.	29.7			
<hr/>						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

#### OLS Regression Results

---

Dep. Variable:	medv	R-squared:	0.484			
Model:	OLS	Adj. R-squared:	0.483			
Method:	Least Squares	F-statistic:	471.8			
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	2.49e-74			
Time:	16:43:09	Log-Likelihood:	-1673.1			
No. Observations:	506	AIC:	3350.			
Df Residuals:	504	BIC:	3359.			
Df Model:	1					
Covariance Type:	nonrobust					
<hr/>						
	coef	std err	t	P> t	[0.025	0.975]
const	-34.6706	2.650	-13.084	0.000	-39.877	-29.465
rm	9.1021	0.419	21.722	0.000	8.279	9.925
<hr/>						

Omnibus:	102.585	Durbin-Watson:	0.684
Prob(Omnibus):	0.000	Jarque-Bera (JB):	612.449
Skew:	0.726	Prob(JB):	1.02e-133
Kurtosis:	8.190	Cond. No.	58.4
<hr/>			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

#### OLS Regression Results

---

Dep. Variable:	medv	R-squared:	0.142
Model:	OLS	Adj. R-squared:	0.140
Method:	Least Squares	F-statistic:	83.48
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	1.57e-18

```

Time: 16:43:09 Log-Likelihood: -1801.5
No. Observations: 506 AIC: 3607.
Df Residuals: 504 BIC: 3615.
Df Model: 1
Covariance Type: nonrobust
=====
          coef    std err      t      P>|t|      [0.025      0.975]
-----
const    30.9787    0.999    31.006    0.000    29.016    32.942
age     -0.1232    0.013   -9.137    0.000   -0.150   -0.097
=====
Omnibus: 170.034 Durbin-Watson: 0.613
Prob(Omnibus): 0.000 Jarque-Bera (JB): 456.983
Skew: 1.671 Prob(JB): 5.85e-100
Kurtosis: 6.240 Cond. No. 195.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

### 1.3.10 Interpretation of regressions

**lstat** For **lstat** gives R-squared value of (0.544), which means that approximately 54.4% of the variance can be explained by the model. The prod F-statistic (5.08e-88) indicates that the model is statistically significant due to the value < 0.05. The negative coefficient (-0.9500) indicates that as **lstat** increases, **medv** tends to decrease, suggesting an inverse relationship between these two variables.

**rm** The same interpretation can be made for **rm**, which has an R-squared value of (0.484), prod F-statistic value of (2.49e-74) and a coefficient of (9.1021). Here the positive coefficient indicates that as **rm** increases, **medv** also tends to increase, suggesting a direct relationship between these two variables.

**age** Same interpretation can be made for **age**, which has an R-squared value of (0.142), prod F-statistic value of (1.57e-18) and a coefficient of (-0.1232). Here the negative coefficient indicates that as **age** increases, **medv** tends to decrease, suggesting an inverse relationship between these two variables.

```
[9]: print(regression1.conf_int())
       print(regression2.conf_int())
       print(regression3.conf_int())
```

	0	1
const	33.448457	35.659225
lstat	-1.026148	-0.873951
	0	1
const	-39.876641	-29.464601

```

rm      8.278855  9.925363
        0           1
const   29.015752 32.941604
age     -0.149647 -0.096679

```

Indicates the lower and upper bounds of the 95% confidence interval. First row (y intercept) and second row (slope of predictor). Smaller intervals in the slope indicates that the model is more precise.

The second one indicates highly positive correlation but it has a greater interval. Meanwhile the first and third one indicates negative correlation with smaller intervals.

### 1.3.11 Use model

```
[10]: use_lstat = pd.DataFrame({"lstat": [5, 10, 15]})
use_lstat = sm.add_constant(use_lstat)
predictor1 = regression1.get_prediction(use_lstat).summary_frame(alpha=0.05)
print(predictor1[["mean", "obs_ci_lower", "obs_ci_upper"]])

use_rm = pd.DataFrame({"rm": [5, 6.5, 8]})
use_rm = sm.add_constant(use_rm)
perdictor2 = regression2.get_prediction(use_rm).summary_frame(alpha=0.05)
print(perdictor2[["mean", "obs_ci_lower", "obs_ci_upper"]])

use_age = pd.DataFrame({"age": [25, 50, 75]})
use_age = sm.add_constant(use_age)
perdictor3 = regression3.get_prediction(use_age).summary_frame(alpha=0.05)
print(perdictor3[["mean", "obs_ci_lower", "obs_ci_upper"]])
```

	mean	obs_ci_lower	obs_ci_upper
0	29.803594	17.565675	42.041513
1	25.053347	12.827626	37.279068
2	20.303101	8.077742	32.528459

	mean	obs_ci_lower	obs_ci_upper
0	10.839924	-2.214474	23.894322
1	24.493088	11.480391	37.505784
2	38.146251	25.058353	51.234149

	mean	obs_ci_lower	obs_ci_upper
0	27.899610	11.090368	44.708852
1	24.820542	8.043748	41.597335
2	21.741474	4.971031	38.511917

### Interpretation of results

**lstat** Inserted values for **lstat** where 5, 10, 15 this means that with 95% confidence the predicted medv values will respectively be between (17.56, 42.04), (12.82, 37.27) and (8.07, 32.52) approximatlly.

**rm & age** Same interpretation can be made for **rm** and **age** where inserted values are 5, 6.5, 8 for **rm** and 25, 50, 75 for **age**.

```
[11]: regression = sm.OLS(df["medv"], sm.add_constant(df[["lstat", "rm", "age"]])).  
      ↪fit()  
      print(regression.summary())
```

OLS Regression Results

---

Dep. Variable:	medv	R-squared:	0.639
Model:	OLS	Adj. R-squared:	0.637
Method:	Least Squares	F-statistic:	296.2
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	1.20e-110
Time:	16:43:09	Log-Likelihood:	-1582.4
No. Observations:	506	AIC:	3173.
Df Residuals:	502	BIC:	3190.
Df Model:	3		
Covariance Type:	nonrobust		

---

	coef	std err	t	P> t	[0.025	0.975]
const	-1.1753	3.182	-0.369	0.712	-7.427	5.076
lstat	-0.6685	0.054	-12.298	0.000	-0.775	-0.562
rm	5.0191	0.454	11.048	0.000	4.127	5.912
age	0.0091	0.011	0.811	0.418	-0.013	0.031

---

Omnibus:	138.819	Durbin-Watson:	0.851
Prob(Omnibus):	0.000	Jarque-Bera (JB):	415.436
Skew:	1.296	Prob(JB):	6.15e-91
Kurtosis:	6.603	Cond. No.	985.

---

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Here the R-squared value (0.639) indicates that approximately 63.9% of the variance in medv can be explained by the model. The prod F-statistic (1.20e-110) indicates that the model is statistically significant due to the value < 0.05.

```
[12]: regression = sm.OLS(df["medv"], sm.add_constant(df.drop(columns=["medv"]))).  
      ↪fit()  
      print(regression.summary())
```

OLS Regression Results

---

Dep. Variable:	medv	R-squared:	0.741
Model:	OLS	Adj. R-squared:	0.734
Method:	Least Squares	F-statistic:	108.1
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	6.72e-135
Time:	16:43:09	Log-Likelihood:	-1498.8
No. Observations:	506	AIC:	3026.

```

Df Residuals: 492   BIC: 3085.
Df Model: 13
Covariance Type: nonrobust
=====
      coef    std err      t      P>|t|      [0.025      0.975]
-----
const    36.4595    5.103    7.144    0.000    26.432    46.487
crim    -0.1080    0.033   -3.287    0.001   -0.173   -0.043
zn       0.0464    0.014    3.382    0.001    0.019    0.073
indus    0.0206    0.061    0.334    0.738   -0.100    0.141
chas     2.6867    0.862    3.118    0.002    0.994    4.380
nox     -17.7666   3.820   -4.651    0.000   -25.272   -10.262
rm       3.8099    0.418    9.116    0.000    2.989    4.631
age      0.0007    0.013    0.052    0.958   -0.025    0.027
dis      -1.4756    0.199   -7.398    0.000   -1.867   -1.084
rad       0.3060    0.066    4.613    0.000    0.176    0.436
tax      -0.0123    0.004   -3.280    0.001   -0.020   -0.005
ptratio   -0.9527    0.131   -7.283    0.000   -1.210   -0.696
black    0.0093    0.003    3.467    0.001    0.004    0.015
lstat    -0.5248    0.051  -10.347    0.000   -0.624   -0.425
=====
Omnibus: 178.041   Durbin-Watson: 1.078
Prob(Omnibus): 0.000   Jarque-Bera (JB): 783.126
Skew: 1.521   Prob(JB): 8.84e-171
Kurtosis: 8.281   Cond. No. 1.51e+04
=====
```

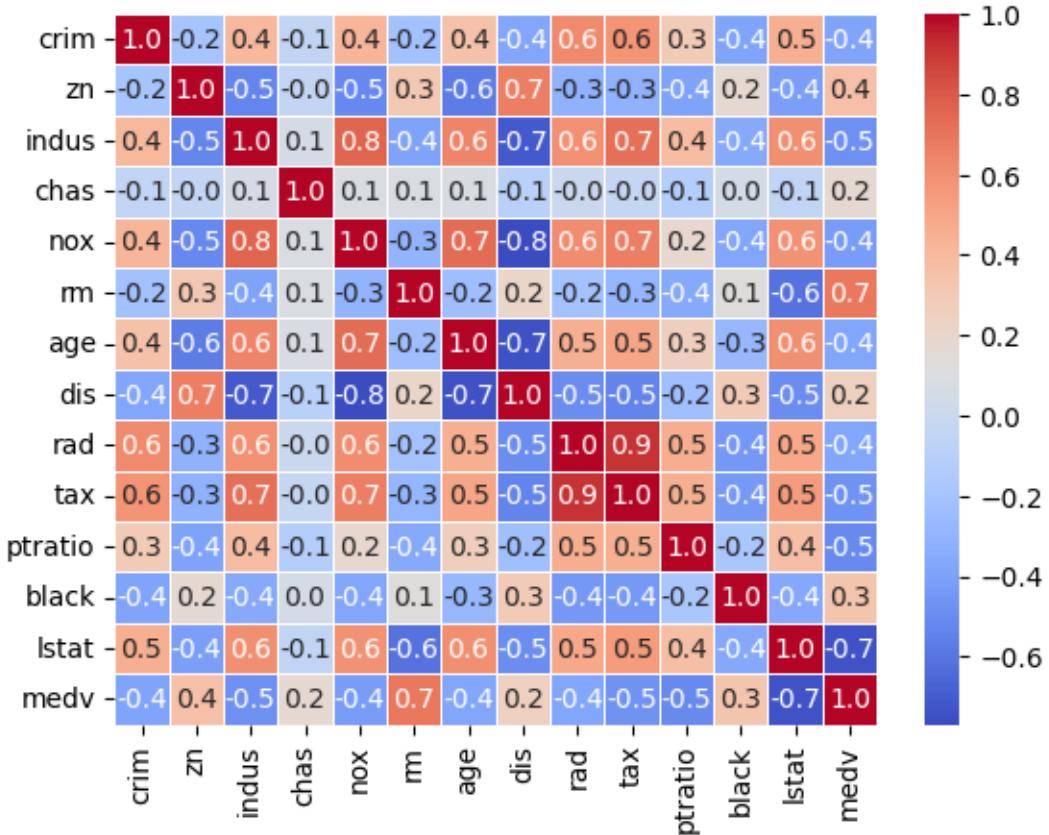
#### Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.51e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Here the R-squared value (0.741) indicates that approximately 74.1% of the variance in medv can be explained by the model. The prod F-statistic (6.72e-135) indicates that the model is statistically significant due to the value < 0.05.

#### Correlation matrix

```
[13]: sns.heatmap(df.corr(), annot=True, cmap="coolwarm", fmt=".1f", linewidths=0.5)
plt.show()
```



This matrix shows the correlation coefficients between each pair of variables in the dataset. A correlation coefficient close to 1 indicates a strong positive correlation, while a coefficient close to -1 indicates a strong negative correlation. A coefficient around 0 suggests no correlation between the variables.

Example interpretations:

`crim` and `zn` have a correlation coefficient of -0.2, indicating a weak negative correlation. This suggests that as the value of `crim` increases, the value of `zn` tends to decrease slightly.

### 1.3.12 Use multiple linear regression model

```
[14]: selected_predictor_values = pd.DataFrame(
    pd.MultiIndex.from_product(
        [[5, 10, 15], [5, 6.5, 8]], names=["lstat", "rm"]
    ).to_frame(index=False)
)
print(selected_predictor_values)

regression = sm.OLS(df["medv"], sm.add_constant(df[["lstat", "rm"]])).fit()
selected_predictor_values = sm.add_constant(selected_predictor_values)
```

```

predictions = regression.get_prediction(selected_predictor_values)
pred_summary = predictions.summary_frame(alpha=0.05)

print(pred_summary[["mean", "obs_ci_lower", "obs_ci_upper"]])

```

	lstat	rm	
0	5	5.0	
1	5	6.5	
2	5	8.0	
3	10	5.0	
4	10	6.5	
5	10	8.0	
6	15	5.0	
7	15	6.5	
8	15	8.0	
	mean	obs_ci_lower	obs_ci_upper
0	20.903875	9.889729	31.918021
1	28.546057	17.635923	39.456192
2	36.188239	25.225479	47.150999
3	17.692084	6.722152	28.662016
4	25.334266	14.437027	36.231505
5	32.976448	21.995024	43.957872
6	14.480292	3.537875	25.422709
7	22.122474	11.221204	33.023745
8	29.764656	18.747835	40.781477

**Interpretation of results** Row one indicates that for value of `lstat` (5) and `rm` (5.0) will with 95% confidence result in a `medv` value between (9.88, 31.91), meanwhile row nine indicates that for value of `lstat` (15) and `rm` (8.0) will with 95% confidence result in a `medv` value between (18.74, 40.78). Same goes for all the subsequent rows with different values for `lstat` and `rm` which also provides a new boundary for all combinations of them.

### 1.3.13 Assignment 3

```

[15]: X = df[["lstat", "rm", "nox", "dis", "ptratio"]]
Y = df["medv"]
X = sm.add_constant(X)
model1 = sm.OLS(Y, X).fit()

print(model1.summary())

```

=====			
Dep. Variable:	medv	R-squared:	0.708
Model:	OLS	Adj. R-squared:	0.705
Method:	Least Squares	F-statistic:	242.6
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	3.67e-131
Time:	16:43:09	Log-Likelihood:	-1528.7

```

No. Observations: 506 AIC: 3069.
Df Residuals: 500 BIC: 3095.
Df Model: 5
Covariance Type: nonrobust
=====
            coef    std err      t    P>|t|    [0.025    0.975]
-----
const      37.4992   4.613     8.129    0.000   28.436   46.562
lstat     -0.5811   0.048    -12.122   0.000  -0.675  -0.487
rm        4.1633   0.412     10.104   0.000   3.354   4.973
nox       -17.9966  3.261    -5.519   0.000  -24.403 -11.590
dis        -1.1847  0.168    -7.034   0.000  -1.516  -0.854
ptratio   -1.0458  0.114    -9.212   0.000  -1.269  -0.823
=====
Omnibus: 187.456 Durbin-Watson: 0.971
Prob(Omnibus): 0.000 Jarque-Bera (JB): 885.498
Skew: 1.584 Prob(JB): 5.21e-193
Kurtosis: 8.654 Cond. No. 545.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[16]: df["lstat_rm"] = df["lstat"] * df["rm"]
X = df[["lstat", "rm", "lstat_rm", "nox", "dis", "ptratio"]]
X = sm.add_constant(X)
model2 = sm.OLS(Y, X).fit()

print(model2.summary())
```

```

OLS Regression Results
=====
Dep. Variable: medv R-squared: 0.778
Model: OLS Adj. R-squared: 0.775
Method: Least Squares F-statistic: 290.8
Date: Sun, 25 Jan 2026 Prob (F-statistic): 2.48e-159
Time: 16:43:09 Log-Likelihood: -1459.9
No. Observations: 506 AIC: 2934.
Df Residuals: 499 BIC: 2963.
Df Model: 6
Covariance Type: nonrobust
=====
            coef    std err      t    P>|t|    [0.025    0.975]
-----
const      3.1518   4.880     0.646    0.519  -6.435   12.739
lstat     1.8115   0.196     9.237   0.000   1.426   2.197
rm        8.3344   0.491    16.971   0.000   7.370   9.299
=====
```

```

lstat_rm      -0.4185      0.034     -12.488      0.000      -0.484      -0.353
nox          -12.3651      2.885     -4.286      0.000     -18.033      -6.697
dis           -1.0184      0.148     -6.893      0.000     -1.309      -0.728
ptratio       -0.7152      0.103     -6.967      0.000     -0.917      -0.514
=====
Omnibus:            246.928 Durbin-Watson:           1.079
Prob(Omnibus):      0.000 Jarque-Bera (JB):        2792.613
Skew:              1.836 Prob(JB):                  0.00
Kurtosis:           13.908 Cond. No.             2.36e+03
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2.36e+03. This might indicate that there are strong multicollinearity or other numerical problems.

**Interpretation of results** By utilizing `lstat`, `rm`, `nox`, `dis` and `ptratio` columns, the model achieves an R-squared value of (0.705), but by just adding the interaction between `lstat` and `rm` the R-squared value increases to (0.775). This indicates that the interaction between `lstat` and `rm` contributes a better prediction of `medv` than just using the individual predictors alone.

### Adding non-linear term

```
[17]: df["lstat_rm_squared"] = df["lstat_rm"] ** 2
X = df[["lstat", "rm", "lstat_rm", "lstat_rm_squared", "nox", "dis", "ptratio"]]
X = sm.add_constant(X)
model3 = sm.OLS(Y, X).fit()

print(model3.summary())
```

OLS Regression Results					
Dep. Variable:	medv	R-squared:	0.781		
Model:	OLS	Adj. R-squared:	0.778		
Method:	Least Squares	F-statistic:	253.9		
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	8.05e-160		
Time:	16:43:09	Log-Likelihood:	-1455.8		
No. Observations:	506	AIC:	2928.		
Df Residuals:	498	BIC:	2961.		
Df Model:	7				
Covariance Type:	nonrobust				
====					
	coef	std err	t	P> t	[0.025
0.975]					
====					
const	10.5522	5.499	1.919	0.056	-0.253

```

21.357
lstat          1.5468     0.216      7.167      0.000      1.123
1.971
rm             7.6004     0.552     13.777      0.000      6.516
8.684
lstat_rm       -0.4468     0.035    -12.864      0.000     -0.515
-0.379
lstat_rm_squared 0.0004     0.000      2.845      0.005      0.000
0.001
nox            -12.2898    2.865     -4.290      0.000     -17.918
-6.662
dis             -1.0641    0.148     -7.209      0.000     -1.354
-0.774
ptratio         -0.7112    0.102     -6.977      0.000     -0.912
-0.511
=====
Omnibus:           217.415   Durbin-Watson:           1.059
Prob(Omnibus):    0.000    Jarque-Bera (JB):        2007.945
Skew:              1.622    Prob(JB):                  0.00
Kurtosis:          12.204   Cond. No.:        3.02e+05
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.02e+05. This might indicate that there are strong multicollinearity or other numerical problems.

### Preform ANOVA

```
[18]: ANOVA_results = sm.stats.anova_lm(model2, model3)
print(ANOVA_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	499.0	9500.381881	0.0	NaN	NaN	NaN
1	498.0	9348.435955	1.0	151.945925	8.094303	0.004623

**Conclusion from ANOVA** ANOVA test indicates that the third model (with interaction and non-linear term) is significantly better then the second model (with interaction) due to the p-value being less than 0.05 (0.004) and it having a better R-squared value.

### Add polynomial

```
[19]: for exp in range(2, 6):
    df[f"lstat_poly_{exp}"] = df["lstat"] ** exp

X = df[
    [
        "lstat",
```

```

    "rm",
    "lstat_rm",
    "lstat_poly_2",
    "lstat_poly_3",
    "lstat_poly_4",
    "lstat_poly_5",
    "nox",
    "dis",
    "ptratio",
]
]

X = sm.add_constant(X)
model4 = sm.OLS(Y, X).fit()

print(model4.summary())

```

### OLS Regression Results

Dep. Variable:	medv	R-squared:	0.792			
Model:	OLS	Adj. R-squared:	0.787			
Method:	Least Squares	F-statistic:	188.0			
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	1.80e-161			
Time:	16:43:09	Log-Likelihood:	-1443.5			
No. Observations:	506	AIC:	2909.			
Df Residuals:	495	BIC:	2956.			
Df Model:	10					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	33.9246	7.663	4.427	0.000	18.869	48.981
lstat	-5.6422	1.426	-3.957	0.000	-8.444	-2.841
rm	6.5291	0.711	9.183	0.000	5.132	7.926
lstat_rm	-0.3055	0.052	-5.878	0.000	-0.408	-0.203
lstat_poly_2	0.8633	0.187	4.622	0.000	0.496	1.230
lstat_poly_3	-0.0495	0.012	-4.153	0.000	-0.073	-0.026
lstat_poly_4	0.0013	0.000	3.795	0.000	0.001	0.002
lstat_poly_5	-1.279e-05	3.64e-06	-3.514	0.000	-1.99e-05	-5.64e-06
nox	-13.7513	2.823	-4.871	0.000	-19.298	-8.204
dis	-1.0326	0.145	-7.127	0.000	-1.317	-0.748
ptratio	-0.7407	0.101	-7.324	0.000	-0.939	-0.542
Omnibus:	232.049	Durbin-Watson:	1.116			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2275.267			
Skew:	1.742	Prob(JB):	0.00			
Kurtosis:	12.787	Cond. No.	3.38e+08			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.38e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
[20]: ANOVA_results = sm.stats.anova_lm(model2, model4)
print(ANOVA_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	499.0	9500.381881	0.0	NaN	NaN	NaN
1	495.0	8903.772453	4.0	596.609428	8.292038	0.000002

**Conclusion from ANOVA** The p-value decreased significantly (0.000002) indicating that the fourth model (with interaction and polynomial term) is statistically significantly better than the third model (with interaction and non-linear term), however with high degree polynomials there is a risk of overfitting the model to the training data, which can lead to poor generalization to new data.

```
[21]: df["log_rm"] = np.log(df["rm"])
```

```
X = df[
    [
        "lstat",
        "lstat_poly_2",
        "lstat_poly_3",
        "lstat_poly_4",
        "lstat_poly_5",
        "rm",
        "log_rm",
        "nox",
        "dis",
        "ptratio",
    ]
]
X = sm.add_constant(X)
model5 = sm.OLS(Y, X).fit()

print(model5.summary())
```

### OLS Regression Results

Dep. Variable:	medv	R-squared:	0.804
Model:	OLS	Adj. R-squared:	0.800
Method:	Least Squares	F-statistic:	202.6
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	7.10e-168
Time:	16:43:09	Log-Likelihood:	-1428.4
No. Observations:	506	AIC:	2879.
Df Residuals:	495	BIC:	2925.

```

Df Model: 10
Covariance Type: nonrobust
=====
      coef   std err      t    P>|t|    [0.025    0.975]
-----
const    172.9866   13.954   12.397   0.000   145.571   200.402
lstat   -8.5527    1.227   -6.969   0.000  -10.964  -6.141
lstat_poly_2  1.0064    0.178    5.654   0.000    0.657   1.356
lstat_poly_3 -0.0582    0.011   -5.087   0.000  -0.081  -0.036
lstat_poly_4  0.0015    0.000    4.672   0.000    0.001   0.002
lstat_poly_5 -1.521e-05  3.52e-06  -4.323   0.000  -2.21e-05 -8.3e-06
rm       25.1967    2.732    9.224   0.000   19.830   30.564
log_rm   -137.4038   16.761   -8.198   0.000  -170.336 -104.472
nox      -16.6408    2.734   -6.087   0.000  -22.012  -11.270
dis      -0.9709    0.141   -6.885   0.000  -1.248  -0.694
ptratio  -0.7843    0.097   -8.116   0.000  -0.974  -0.594
=====
Omnibus: 221.958 Durbin-Watson: 1.064
Prob(Omnibus): 0.000 Jarque-Bera (JB): 2718.500
Skew: 1.567 Prob(JB): 0.00
Kurtosis: 13.914 Cond. No. 9.63e+08
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.63e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
[22]: ANOVA_results = sm.stats.anova_lm(model2, model5)
print(ANOVA_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	499.0	9500.381881	0.0	NaN	NaN	NaN
1	495.0	8386.756361	4.0	1113.62552	16.431997	1.190966e-12

**Conclusion from ANOVA** Increase in R-squared value and the very low p-value indicates that the fifth model (with interaction, higher degree polynomial term and logarithmic term) is statistically significantly better than the fourth model (with interaction and polynomial term).

### Load data 2

```
[23]: df2 = pd.read_csv("../data/Carseats.csv", index_col=0)
print(df2.describe(), "\n")
print(df2["ShelveLoc"].value_counts(), "\n")
print(df2["Urban"].value_counts(), "\n")
print(df2["US"].value_counts())
```

```

      Sales  CompPrice     Income  Advertising  Population \
count  400.000000  400.000000  400.000000  400.000000  400.000000
mean   7.496325  124.975000  68.657500   6.635000  264.840000
std    2.824115  15.334512  27.986037   6.650364  147.376436
min    0.000000  77.000000  21.000000   0.000000  10.000000
25%    5.390000  115.000000  42.750000   0.000000  139.000000
50%    7.490000  125.000000  69.000000   5.000000  272.000000
75%    9.320000  135.000000  91.000000  12.000000  398.500000
max   16.270000  175.000000 120.000000  29.000000  509.000000

```

```

      Price        Age   Education
count  400.000000  400.000000  400.000000
mean   115.795000  53.322500  13.900000
std    23.676664  16.200297  2.620528
min    24.000000  25.000000  10.000000
25%    100.000000  39.750000  12.000000
50%    117.000000  54.500000  14.000000
75%    131.000000  66.000000  16.000000
max   191.000000  80.000000  18.000000

```

```

ShelveLoc
Medium      219
Bad         96
Good        85
Name: count, dtype: int64

```

```

Urban
Yes       282
No        118
Name: count, dtype: int64

```

```

US
Yes      258
No       142
Name: count, dtype: int64

```

```

[24]: X = pd.get_dummies(df2, columns=["ShelveLoc", "Urban", "US"])
for column in X.select_dtypes("bool"):
    X[column] = X[column].astype(int)

X = X.drop(columns=["Sales"])
X = sm.add_constant(X)
Y = df2["Sales"]
model = sm.OLS(Y, X).fit()

print(model.summary())

```

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.873		
Model:	OLS	Adj. R-squared:	0.870		
Method:	Least Squares	F-statistic:	243.4		
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	1.60e-166		
Time:	16:43:09	Log-Likelihood:	-568.99		
No. Observations:	400	AIC:	1162.		
Df Residuals:	388	BIC:	1210.		
Df Model:	11				
Covariance Type:	nonrobust				
0.975]					
const	3.3853	0.253	13.370	0.000	2.887
3.883					
CompPrice	0.0928	0.004	22.378	0.000	0.085
0.101					
Income	0.0158	0.002	8.565	0.000	0.012
0.019					
Advertising	0.1231	0.011	11.066	0.000	0.101
0.145					
Population	0.0002	0.000	0.561	0.575	-0.001
0.001					
Price	-0.0954	0.003	-35.700	0.000	-0.101
-0.090					
Age	-0.0460	0.003	-14.472	0.000	-0.052
-0.040					
Education	-0.0211	0.020	-1.070	0.285	-0.060
0.018					
ShelveLoc_Bad	-1.1405	0.118	-9.629	0.000	-1.373
-0.908					
ShelveLoc_Good	3.7096	0.121	30.652	0.000	3.472
3.948					
ShelveLoc_Medium	0.8162	0.107	7.605	0.000	0.605
1.027					
Urban_No	1.6312	0.138	11.789	0.000	1.359
1.903					
Urban_Yes	1.7541	0.139	12.629	0.000	1.481
2.027					
US_No	1.7847	0.146	12.243	0.000	1.498
2.071					
US_Yes	1.6006	0.148	10.783	0.000	1.309
1.892					
Omnibus:	0.811	Durbin-Watson:	2.013		

Prob(Omnibus):	0.667	Jarque-Bera (JB):	0.765
Skew:	0.107	Prob(JB):	0.682
Kurtosis:	2.994	Cond. No.	3.32e+18

---

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 4.43e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

**Conclusion from model summary** The model achieves an R-squared value (0.873) which indicates that  $\approx 87.3\%$  of the variance can be explained by the model. The prob F-statistic (1.60e-166) indicates that the model is statistically significant due to the value  $< 0.05$ .

```
[25]: X = df2.drop(columns=["Sales", "Population", "Education", "Age", "Urban", "US"])
X = pd.get_dummies(X, columns=["ShelveLoc"])

for column in X.select_dtypes("bool"):
    X[column] = X[column].astype(int)

X["Income:Advertising"] = df2["Income"] * df2["Advertising"]
X["Price:Age"] = df2["Price"] * df2["Age"]
Y = df2["Sales"]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()

print(model.summary())
```

### OLS Regression Results

---

Dep. Variable:	Sales	R-squared:	0.870
Model:	OLS	Adj. R-squared:	0.868
Method:	Least Squares	F-statistic:	328.2
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	2.90e-168
Time:	16:43:09	Log-Likelihood:	-573.74
No. Observations:	400	AIC:	1165.
Df Residuals:	391	BIC:	1201.
Df Model:	8		
Covariance Type:	nonrobust		

---



---

	coef	std err	t	P> t	[0.025
0.975]					
-----					
const	4.1957	0.352	11.903	0.000	3.503
4.889					

CompPrice	0.0934	0.004	22.492	0.000	0.085
0.102					
Income	0.0098	0.003	3.756	0.000	0.005
0.015					
Advertising	0.0534	0.021	2.544	0.011	0.012
0.095					
Price	-0.0759	0.003	-25.591	0.000	-0.082
-0.070					
ShelveLoc_Bad	-0.8966	0.143	-6.280	0.000	-1.177
-0.616					
ShelveLoc_Good	3.9982	0.149	26.769	0.000	3.705
4.292					
ShelveLoc_Medium	1.0942	0.133	8.221	0.000	0.833
1.356					
Income:Advertising	0.0009	0.000	3.124	0.002	0.000
0.001					
Price:Age	-0.0004	2.69e-05	-13.713	0.000	-0.000
-0.000					
<hr/>					
Omnibus:	1.537	Durbin-Watson:	1.988		
Prob(Omnibus):	0.464	Jarque-Bera (JB):	1.326		
Skew:	0.129	Prob(JB):	0.515		
Kurtosis:	3.116	Cond. No.	2.80e+19		
<hr/>					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.18e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

**Conclusion form model summary** Achieves an R-squared value (0.870) which is worse than the previous model and prob F-statistic (2.90e-168) which indicates that the model is more statistically significant due to the value < 0.05.

### Beat the teacher

```
[26]: X = df2.drop(columns=["Sales"])
X = pd.get_dummies(X, columns=["ShelveLoc", "US", "Urban"])

for column in X.select_dtypes("bool"):
    X[column] = X[column].astype(int)

X["Income:Advertising"] = df2["Income"] * df2["Advertising"]
Y = df2["Sales"]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()
```

```
print(model.summary())
```

OLS Regression Results								
Dep. Variable:	Sales	R-squared:	0.876					
Model:	OLS	Adj. R-squared:	0.872					
Method:	Least Squares	F-statistic:	227.6					
Date:	Sun, 25 Jan 2026	Prob (F-statistic):	5.48e-167					
Time:	16:43:09	Log-Likelihood:	-565.00					
No. Observations:	400	AIC:	1156.					
Df Residuals:	387	BIC:	1208.					
Df Model:	12							
Covariance Type:	nonrobust							
=====								
=====								
	coef	std err	t	P> t	[0.025			
0.975]								
-----								
const	3.5106	0.255	13.767	0.000	3.009			
4.012								
CompPrice	0.0931	0.004	22.630	0.000	0.085			
0.101								
Income	0.0107	0.003	4.125	0.000	0.006			
0.016								
Advertising	0.0684	0.022	3.043	0.003	0.024			
0.113								
Population	0.0002	0.000	0.456	0.649	-0.001			
0.001								
Price	-0.0952	0.003	-35.962	0.000	-0.100			
-0.090								
Age	-0.0454	0.003	-14.367	0.000	-0.052			
-0.039								
Education	-0.0220	0.020	-1.125	0.261	-0.060			
0.016								
ShelveLoc_Bad	-1.1051	0.118	-9.356	0.000	-1.337			
-0.873								
ShelveLoc_Good	3.7570	0.121	31.005	0.000	3.519			
3.995								
ShelveLoc_Medium	0.8586	0.107	7.988	0.000	0.647			
1.070								
US_No	1.8361	0.146	12.604	0.000	1.550			
2.123								
US_Yes	1.6744	0.150	11.199	0.000	1.380			
1.968								
Urban_No	1.6884	0.139	12.173	0.000	1.416			
1.961								
Urban_Yes	1.8222	0.140	13.031	0.000	1.547			

2.097					
Income:Advertising	0.0008	0.000	2.791	0.006	0.000
0.001					
=====					
Omnibus:	1.390	Durbin-Watson:		2.036	
Prob(Omnibus):	0.499	Jarque-Bera (JB):		1.229	
Skew:	0.131	Prob(JB):		0.541	
Kurtosis:	3.070	Cond. No.		1.15e+19	
=====					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.73e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

**Conclusion from Beat the teacher model** I was able to get slightly better R-squared value (0.876) with a prob F-statistic (5.48e-167) which is in between the two previous models significance. I did this by not dropping any predictors and adding interaction term between Income and Advertising.