

Assignment 3

Simple and Multiple Linear Regression pt2

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-02-09

Course: Machine Learning 4DT905

Conceptual

$$Y = \hat{\beta}_0 + \hat{\beta}_1 X_1 + \hat{\beta}_2 X_2 + \hat{\beta}_3 X_3 + \hat{\beta}_4 X_1 X_2 + \hat{\beta}_5 X_1 X_3$$

$$\hat{\beta}_0 = 50(\text{Intercept})$$

$$\hat{\beta}_1 = 20(\text{GPA})$$

$$\hat{\beta}_2 = 0.07(\text{IQ})$$

$$\hat{\beta}_3 = 35(\text{Level})$$

$$\hat{\beta}_4 = 0.01(\text{GPA} \cdot \text{IQ})$$

$$\hat{\beta}_5 = -10(\text{GPA} \cdot \text{Level})$$

$X_3 = 1$ for College, 0 for High School

1.

$$Y_c = 50 + 20X_1 + 0.07X_2 + 35 + 0.01X_1X_2 - 10X_1$$

$$Y_h = 50 + 20X_1 + 0.07X_2 + 0.01X_1X_2$$

$$Y_c - Y_h = 35 - 10X_1$$

$$35 - 10X_1 = 0 \implies X_1 = 3.5$$

True, when GPA > 3.5 High School graduates earn more than College graduates.

Answer: iii

2.

$$X_1 = 4.0$$

$$X_2 = 110$$

$$X_3 = 1$$

$$Y = 50 + 20(4.0) + 0.07(110) + 35 + 0.01(4.0)(110) - 10(4.0)$$

$$Y = 137.1$$

Answer: \$137,100

3. False. The magnitude of a coefficient does not indicate statistical importance. To determine statistical importance we need to look at the p-values associated with that coefficient, not just its absolute value. In the presented case, the units of predictor X_2 (IQ) are generally > 100 . A small coefficient for the $X_2 \cdot X_1$ term might still result in a large contribution to the model and be highly statistically significant.

Answer: *False*

Practical

Imports

```
In [1]: import pandas as pd  
import statsmodels.api as sm  
import numpy as np
```

Load data

```
In [2]: df = pd.read_csv("../data/Boston.csv", index_col=0)
```

```
In [3]: X = df[["lstat", "rm", "nox", "dis", "ptratio"]]  
Y = df["medv"]  
X = sm.add_constant(X)  
model1 = sm.OLS(Y, X).fit()  
  
print(model1.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.708			
Model:	OLS	Adj. R-squared:	0.705			
Method:	Least Squares	F-statistic:	242.6			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	3.67e-131			
Time:	06:38:55	Log-Likelihood:	-1528.7			
No. Observations:	506	AIC:	3069.			
Df Residuals:	500	BIC:	3095.			
Df Model:	5					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	37.4992	4.613	8.129	0.000	28.436	46.562
lstat	-0.5811	0.048	-12.122	0.000	-0.675	-0.487
rm	4.1633	0.412	10.104	0.000	3.354	4.973
nox	-17.9966	3.261	-5.519	0.000	-24.403	-11.590
dis	-1.1847	0.168	-7.034	0.000	-1.516	-0.854
ptratio	-1.0458	0.114	-9.212	0.000	-1.269	-0.823
Omnibus:	187.456	Durbin-Watson:	0.971			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	885.498			
Skew:	1.584	Prob(JB):	5.21e-193			
Kurtosis:	8.654	Cond. No.	545.			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [4]: df["lstat_rm"] = df["lstat"] * df["rm"]
X = df[["lstat", "rm", "lstat_rm", "nox", "dis", "ptratio"]]
X = sm.add_constant(X)
model2 = sm.OLS(Y, X).fit()

print(model2.summary())
```

OLS Regression Results						
Dep. Variable:	medv	R-squared:	0.778			
Model:	OLS	Adj. R-squared:	0.775			
Method:	Least Squares	F-statistic:	290.8			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	2.48e-159			
Time:	06:38:55	Log-Likelihood:	-1459.9			
No. Observations:	506	AIC:	2934.			
Df Residuals:	499	BIC:	2963.			
Df Model:	6					
Covariance Type:	nonrobust					
coef	std err	t	P> t	[0.025	0.975]	
const	3.1518	4.880	0.646	0.519	-6.435	12.739
lstat	1.8115	0.196	9.237	0.000	1.426	2.197
rm	8.3344	0.491	16.971	0.000	7.370	9.299
lstat_rm	-0.4185	0.034	-12.488	0.000	-0.484	-0.353
nox	-12.3651	2.885	-4.286	0.000	-18.033	-6.697
dis	-1.0184	0.148	-6.893	0.000	-1.309	-0.728
ptratio	-0.7152	0.103	-6.967	0.000	-0.917	-0.514
Omnibus:	246.928	Durbin-Watson:	1.079			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2792.613			
Skew:	1.836	Prob(JB):	0.00			
Kurtosis:	13.908	Cond. No.	2.36e+03			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.36e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Interpretation of results

By utilizing `lstat`, `rm`, `nox`, `dis` and `ptratio` columns, the model achieves an R-squared value of (0.705), but by just adding the interaction between `lstat` and `rm` the R-squared value increases to (0.775). This indicates that the interaction between `lstat` and `rm` contributes a better prediction of `medv` than just using the individual predictors alone.

Confidence Intervals: Examining the confidence intervals for the coefficients in model1 and model2, we should check whether they include zero (indicating significance) and their width (indicating precision of the estimate). Coefficients with confidence intervals that do not include zero are statistically significant predictors of the target variable.

Correlation and Multicollinearity: Before building the models, it would be beneficial to examine the correlation matrix to identify: (1) which predictors are most correlated with the target variable `medv`, and (2) whether there is multicollinearity between predictors (high correlation between predictor variables). High multicollinearity can affect the stability and interpretability of coefficient estimates.

Adding non-linear term

```
In [5]: df["lstat_rm_squared"] = df["lstat_rm"] ** 2
X = df[["lstat", "rm", "lstat_rm", "lstat_rm_squared", "nox", "dis", "ptratio"]]
X = sm.add_constant(X)
model3 = sm.OLS(Y, X).fit()

print(model3.summary())
```

OLS Regression Results

Dep. Variable:	medv	R-squared:	0.781			
Model:	OLS	Adj. R-squared:	0.778			
Method:	Least Squares	F-statistic:	253.9			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	8.05e-160			
Time:	06:38:55	Log-Likelihood:	-1455.8			
No. Observations:	506	AIC:	2928.			
Df Residuals:	498	BIC:	2961.			
Df Model:	7					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	10.5522	5.499	1.919	0.056	-0.253	21.357
lstat	1.5468	0.216	7.167	0.000	1.123	1.971
rm	7.6004	0.552	13.777	0.000	6.516	8.684
lstat_rm	-0.4468	0.035	-12.864	0.000	-0.515	-0.379
lstat_rm_squared	0.0004	0.000	2.845	0.005	0.000	0.001
nox	-12.2898	2.865	-4.290	0.000	-17.918	-6.662
dis	-1.0641	0.148	-7.209	0.000	-1.354	-0.774
ptratio	-0.7112	0.102	-6.977	0.000	-0.912	-0.511

Omnibus: 217.415 Durbin-Watson: 1.059
 Prob(Omnibus): 0.000 Jarque-Bera (JB): 2007.945
 Skew: 1.622 Prob(JB): 0.00
 Kurtosis: 12.204 Cond. No. 3.02e+05

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.02e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Preform ANOVA

```
In [6]: ANOVA_results = sm.stats.anova_lm(model2, model3)
print(ANOVA_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	499.0	9500.381881	0.0	NaN	NaN	NaN
1	498.0	9348.435955	1.0	151.945925	8.094303	0.004623

Conclusion from ANOVA

Hypothesis Test: The ANOVA F-test compares model2 (with interaction term) versus model3 (with interaction and non-linear term).

- Null Hypothesis (H0): The simpler model (model2) is sufficient.
- Alternative Hypothesis (H1): The more complex model (model3) provides a significantly better fit.

Result: With a p-value of 0.004 (< 0.05), we reject the null hypothesis. The ANOVA test indicates that model3 (with the squared interaction term) provides a statistically significantly better fit than model2. The F-statistic measures the improvement in model fit relative to the increase in model complexity. The significant p-value suggests that adding the non-linear term meaningfully improves the model's ability to explain variance in the target variable.

Add polynomial

```
In [7]: for exp in range(2, 6):
    df[f"lstat_poly_{exp}"] = df["lstat"] ** exp

X = df[
    [
        "lstat",
        "rm",
        "lstat_rm",
        "lstat_poly_2",
        "lstat_poly_3",
        "lstat_poly_4",
        "lstat_poly_5",
        "nox",
        "dis",
        "ptratio",
    ]
]
X = sm.add_constant(X)
model4 = sm.OLS(Y, X).fit()

print(model4.summary())
```

OLS Regression Results						
Dep. Variable:	medv	R-squared:	0.792			
Model:	OLS	Adj. R-squared:	0.787			
Method:	Least Squares	F-statistic:	188.0			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	1.80e-161			
Time:	06:38:55	Log-Likelihood:	-1443.5			
No. Observations:	506	AIC:	2909.			
Df Residuals:	495	BIC:	2956.			
Df Model:	10					
Covariance Type:	nonrobust					
coef	std err	t	P> t	[0.025	0.975]	
const	33.9246	7.663	4.427	0.000	18.869	48.981
lstat	-5.6422	1.426	-3.957	0.000	-8.444	-2.841
rm	6.5291	0.711	9.183	0.000	5.132	7.926
lstat_rm	-0.3055	0.052	-5.878	0.000	-0.408	-0.203
lstat_poly_2	0.8633	0.187	4.622	0.000	0.496	1.230
lstat_poly_3	-0.0495	0.012	-4.153	0.000	-0.073	-0.026
lstat_poly_4	0.0013	0.000	3.795	0.000	0.001	0.002
lstat_poly_5	-1.279e-05	3.64e-06	-3.514	0.000	-1.99e-05	-5.64e-06
nox	-13.7513	2.823	-4.871	0.000	-19.298	-8.204
dis	-1.0326	0.145	-7.127	0.000	-1.317	-0.748
ptratio	-0.7407	0.101	-7.324	0.000	-0.939	-0.542
Omnibus:	232.049	Durbin-Watson:	1.116			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2275.267			
Skew:	1.742	Prob(JB):	0.00			
Kurtosis:	12.787	Cond. No.	3.38e+08			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.38e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [8]: ANOVA_results = sm.stats.anova_lm(model2, model4)
print(ANOVA_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	499.0	9500.381881	0.0	NaN	NaN	NaN
1	495.0	8903.772453	4.0	596.609428	8.292038	0.000002

Conclusion from ANOVA

Hypothesis Test: The ANOVA F-test compares model2 (with interaction term) versus model4 (with interaction and polynomial terms).

- Null Hypothesis (H0): The simpler model (model2) is sufficient.
- Alternative Hypothesis (H1): The more complex model (model4) provides a significantly better fit.

Result: With a very low p-value (0.000002 << 0.05), we reject the null hypothesis. The

ANOVA test indicates that model4 (with polynomial terms) provides a statistically significantly better fit than model2. However, it's important to note that high-degree polynomials increase the risk of overfitting to the training data, which can lead to poor generalization on new data. Cross-validation should be considered to assess true predictive performance.

```
In [9]: df["log_rm"] = np.log(df["rm"])

X = df[
    [
        "lstat",
        "lstat_poly_2",
        "lstat_poly_3",
        "lstat_poly_4",
        "lstat_poly_5",
        "rm",
        "log_rm",
        "nox",
        "dis",
        "ptratio",
    ]
]
X = sm.add_constant(X)
model5 = sm.OLS(Y, X).fit()

print(model5.summary())
```

OLS Regression Results						
Dep. Variable:	medv	R-squared:	0.804			
Model:	OLS	Adj. R-squared:	0.800			
Method:	Least Squares	F-statistic:	202.6			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	7.10e-168			
Time:	06:38:55	Log-Likelihood:	-1428.4			
No. Observations:	506	AIC:	2879.			
Df Residuals:	495	BIC:	2925.			
Df Model:	10					
Covariance Type:	nonrobust					
coef	std err	t	P> t	[0.025	0.975]	
const	172.9866	13.954	12.397	0.000	145.571	200.402
lstat	-8.5527	1.227	-6.969	0.000	-10.964	-6.141
lstat_poly_2	1.0064	0.178	5.654	0.000	0.657	1.356
lstat_poly_3	-0.0582	0.011	-5.087	0.000	-0.081	-0.036
lstat_poly_4	0.0015	0.000	4.672	0.000	0.001	0.002
lstat_poly_5	-1.521e-05	3.52e-06	-4.323	0.000	-2.21e-05	-8.3e-06
rm	25.1967	2.732	9.224	0.000	19.830	30.564
log_rm	-137.4038	16.761	-8.198	0.000	-170.336	-104.472
nox	-16.6408	2.734	-6.087	0.000	-22.012	-11.270
dis	-0.9709	0.141	-6.885	0.000	-1.248	-0.694
ptratio	-0.7843	0.097	-8.116	0.000	-0.974	-0.594
Omnibus:	221.958	Durbin-Watson:	1.064			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	2718.500			
Skew:	1.567	Prob(JB):	0.00			
Kurtosis:	13.914	Cond. No.	9.63e+08			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 9.63e+08. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [10]: ANOVA_results = sm.stats.anova_lm(model2, model5)
print(ANOVA_results)
```

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	499.0	9500.381881	0.0	NaN	NaN	NaN
1	495.0	8386.756361	4.0	1113.62552	16.431997	1.190966e-12

Conclusion from ANOVA

Hypothesis Test: The ANOVA F-test compares model2 (with interaction term) versus model5 (with interaction, polynomial, and logarithmic terms).

- Null Hypothesis (H0): The simpler model (model2) is sufficient.
- Alternative Hypothesis (H1): The more complex model (model5) provides a significantly better fit.

Result: With a very low p-value (much less than 0.05), we reject the null hypothesis. The

ANOVA test indicates that model5 provides a statistically significantly better fit than model2. The increase in R-squared value demonstrates improved explanatory power. The combination of polynomial and logarithmic transformations captures both polynomial trends and logarithmic relationships in the data.

Load data 2

```
In [11]: df2 = pd.read_csv("../data/Carseats.csv", index_col=0)
print(df2.describe(), "\n")
print(df2["ShelveLoc"].value_counts(), "\n")
print(df2["Urban"].value_counts(), "\n")
print(df2["US"].value_counts())
```

	Sales	CompPrice	Income	Advertising	Population	\
count	400.000000	400.000000	400.000000	400.000000	400.000000	
mean	7.496325	124.975000	68.657500	6.635000	264.840000	
std	2.824115	15.334512	27.986037	6.650364	147.376436	
min	0.000000	77.000000	21.000000	0.000000	10.000000	
25%	5.390000	115.000000	42.750000	0.000000	139.000000	
50%	7.490000	125.000000	69.000000	5.000000	272.000000	
75%	9.320000	135.000000	91.000000	12.000000	398.500000	
max	16.270000	175.000000	120.000000	29.000000	509.000000	

	Price	Age	Education
count	400.000000	400.000000	400.000000
mean	115.795000	53.322500	13.900000
std	23.676664	16.200297	2.620528
min	24.000000	25.000000	10.000000
25%	100.000000	39.750000	12.000000
50%	117.000000	54.500000	14.000000
75%	131.000000	66.000000	16.000000
max	191.000000	80.000000	18.000000

ShelveLoc
Medium 219
Bad 96
Good 85
Name: count, dtype: int64

Urban
Yes 282
No 118
Name: count, dtype: int64

US
Yes 258
No 142
Name: count, dtype: int64

```
In [12]: X = pd.get_dummies(df2, columns=["ShelveLoc", "Urban", "US"])
for column in X.select_dtypes("bool"):
    X[column] = X[column].astype(int)

X = X.drop(columns=["Sales"])
```

```
X = sm.add_constant(X)
Y = df2["Sales"]
model = sm.OLS(Y, X).fit()

print(model.summary())
```

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.873			
Model:	OLS	Adj. R-squared:	0.870			
Method:	Least Squares	F-statistic:	243.4			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	1.60e-166			
Time:	06:38:55	Log-Likelihood:	-568.99			
No. Observations:	400	AIC:	1162.			
Df Residuals:	388	BIC:	1210.			
Df Model:	11					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	3.3853	0.253	13.370	0.000	2.887	3.883
CompPrice	0.0928	0.004	22.378	0.000	0.085	0.101
Income	0.0158	0.002	8.565	0.000	0.012	0.019
Advertising	0.1231	0.011	11.066	0.000	0.101	0.145
Population	0.0002	0.000	0.561	0.575	-0.001	0.001
Price	-0.0954	0.003	-35.700	0.000	-0.101	-0.090
Age	-0.0460	0.003	-14.472	0.000	-0.052	-0.040
Education	-0.0211	0.020	-1.070	0.285	-0.060	0.018
ShelveLoc_Bad	-1.1405	0.118	-9.629	0.000	-1.373	-0.908
ShelveLoc_Good	3.7096	0.121	30.652	0.000	3.472	3.948
ShelveLoc_Medium	0.8162	0.107	7.605	0.000	0.605	1.027
Urban_No	1.6312	0.138	11.789	0.000	1.359	1.903
Urban_Yes	1.7541	0.139	12.629	0.000	1.481	2.027
US_No	1.7847	0.146	12.243	0.000	1.498	2.071
US_Yes	1.6006	0.148	10.783	0.000	1.309	1.892
Omnibus:	0.811	Durbin-Watson:	2.013			
Prob(Omnibus):	0.667	Jarque-Bera (JB):	0.765			
Skew:	0.107	Prob(JB):	0.682			
Kurtosis:	2.994	Cond. No.	3.32e+18			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 4.43e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Conclusion from model summary

The model achieves an R-squared value of 0.873, which indicates that approximately 87.3% of the variance in sales can be explained by the model. The prob F-statistic (1.60e-166) is much less than 0.05, indicating that the model is statistically significant overall.

Coefficient Significance: When interpreting individual coefficients, we should examine their

confidence intervals. Coefficients whose 95% confidence intervals do not include zero are statistically significant predictors. The width of the confidence interval indicates the precision of our estimate - narrower intervals suggest more precise estimates.

```
In [13]: X = df2.drop(columns=["Sales", "Population", "Education", "Age", "Urban", "US"])
X = pd.get_dummies(X, columns=["ShelveLoc"])

for column in X.select_dtypes("bool"):
    X[column] = X[column].astype(int)

X["Income:Advertising"] = df2["Income"] * df2["Advertising"]
X["Price:Age"] = df2["Price"] * df2["Age"]
Y = df2["Sales"]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()

print(model.summary())
```

OLS Regression Results						
Dep. Variable:	Sales	R-squared:	0.870			
Model:	OLS	Adj. R-squared:	0.868			
Method:	Least Squares	F-statistic:	328.2			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	2.90e-168			
Time:	06:38:55	Log-Likelihood:	-573.74			
No. Observations:	400	AIC:	1165.			
Df Residuals:	391	BIC:	1201.			
Df Model:	8					
Covariance Type:	nonrobust					
<hr/>						
==						
	coef	std err	t	P> t	[0.025	0.97
5]						
<hr/>						
--						
const	4.1957	0.352	11.903	0.000	3.503	4.8
89						
CompPrice	0.0934	0.004	22.492	0.000	0.085	0.1
02						
Income	0.0098	0.003	3.756	0.000	0.005	0.0
15						
Advertising	0.0534	0.021	2.544	0.011	0.012	0.0
95						
Price	-0.0759	0.003	-25.591	0.000	-0.082	-0.0
70						
ShelveLoc_Bad	-0.8966	0.143	-6.280	0.000	-1.177	-0.6
16						
ShelveLoc_Good	3.9982	0.149	26.769	0.000	3.705	4.2
92						
ShelveLoc_Medium	1.0942	0.133	8.221	0.000	0.833	1.3
56						
Income:Advertising	0.0009	0.000	3.124	0.002	0.000	0.0
01						
Price:Age	-0.0004	2.69e-05	-13.713	0.000	-0.000	-0.0
00						
<hr/>						
Omnibus:	1.537	Durbin-Watson:	1.988			
Prob(Omnibus):	0.464	Jarque-Bera (JB):	1.326			
Skew:	0.129	Prob(JB):	0.515			
Kurtosis:	3.116	Cond. No.	2.80e+19			
<hr/>						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.18e-29. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Conclusion from model summary

This model achieves an R-squared value of 0.870, which is slightly worse than the previous model (0.873). The prob F-statistic (2.90e-168) indicates that the model is statistically significant overall ($p < 0.05$).

Coefficient Interpretation: When examining individual coefficients, we should check their confidence intervals to assess both significance (whether the interval includes zero) and precision (width of the interval). Some predictors may become insignificant when used together with other predictors due to multicollinearity - when predictor variables are highly correlated with each other. In such cases, a predictor that appears significant in isolation may become insignificant in a multiple regression model because its information is already captured by correlated predictors.

Beat the teacher

```
In [14]: X = df2.drop(columns=["Sales"])
X = pd.get_dummies(X, columns=["ShelveLoc", "US", "Urban"])

for column in X.select_dtypes("bool"):
    X[column] = X[column].astype(int)

X["Income:Advertising"] = df2["Income"] * df2["Advertising"]
Y = df2["Sales"]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()

print(model.summary())
```

OLS Regression Results

Dep. Variable:	Sales	R-squared:	0.876			
Model:	OLS	Adj. R-squared:	0.872			
Method:	Least Squares	F-statistic:	227.6			
Date:	Mon, 09 Feb 2026	Prob (F-statistic):	5.48e-167			
Time:	06:38:55	Log-Likelihood:	-565.00			
No. Observations:	400	AIC:	1156.			
Df Residuals:	387	BIC:	1208.			
Df Model:	12					
Covariance Type:	nonrobust					
<hr/>						
==						
	coef	std err	t	P> t	[0.025	0.97
5]						
<hr/>						
--						
const	3.5106	0.255	13.767	0.000	3.009	4.0
12						
CompPrice	0.0931	0.004	22.630	0.000	0.085	0.1
01						
Income	0.0107	0.003	4.125	0.000	0.006	0.0
16						
Advertising	0.0684	0.022	3.043	0.003	0.024	0.1
13						
Population	0.0002	0.000	0.456	0.649	-0.001	0.0
01						
Price	-0.0952	0.003	-35.962	0.000	-0.100	-0.0
90						
Age	-0.0454	0.003	-14.367	0.000	-0.052	-0.0
39						
Education	-0.0220	0.020	-1.125	0.261	-0.060	0.0
16						
ShelveLoc_Bad	-1.1051	0.118	-9.356	0.000	-1.337	-0.8
73						
ShelveLoc_Good	3.7570	0.121	31.005	0.000	3.519	3.9
95						
ShelveLoc_Medium	0.8586	0.107	7.988	0.000	0.647	1.0
70						
US_No	1.8361	0.146	12.604	0.000	1.550	2.1
23						
US_Yes	1.6744	0.150	11.199	0.000	1.380	1.9
68						
Urban_No	1.6884	0.139	12.173	0.000	1.416	1.9
61						
Urban_Yes	1.8222	0.140	13.031	0.000	1.547	2.0
97						
Income:Advertising	0.0008	0.000	2.791	0.006	0.000	0.0
01						
<hr/>						
Omnibus:	1.390	Durbin-Watson:	2.036			
Prob(Omnibus):	0.499	Jarque-Bera (JB):	1.229			
Skew:	0.131	Prob(JB):	0.541			
Kurtosis:	3.070	Cond. No.	1.15e+19			
<hr/>						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.73e-30. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

Conclusion from Beat the teacher model

This model achieves an R-squared value of 0.876, which is slightly better than the previous models. The prob F-statistic (5.48e-167) indicates high statistical significance ($p < 0.05$). By including all available predictors and adding an interaction term between Income and Advertising , the model captures the synergistic effect where the impact of income on sales may depend on advertising levels (or vice versa).

Model Selection Considerations: When comparing models, we should consider:

1. **Confidence Intervals:** Check that key predictors have confidence intervals that exclude zero (indicating significance) and assess the precision of estimates.
2. **Multicollinearity:** If predictors are highly correlated, some may appear insignificant even if they contain useful information.
3. **Practical Significance:** Beyond statistical significance, consider whether the improvement in R-squared justifies the added model complexity.