

# Assignment 4

## Classification

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-26

Course: Machine Learning 4DT905

### Conceptual

1. Linear Discriminant Analysis (LDA) and Quadratic Discriminant Analysis (QDA) are both generative models that assume the data follows a Gaussian (Normal) distribution. However, they differ significantly in their flexibility.

- Class Covariance:

- **LDA** assumes that all classes share the same covariance matrix. This means that the decision boundaries between classes are linear. LDA is suitable when the classes are well-separated and the assumption of equal covariance holds true.
- **QDA**, on the other hand, assumes each class to have its own covariance matrix. This results in quadratic decision boundaries, making QDA more flexible and capable of modeling more complex relationships between features. QDA is preferable when the classes have different variances or when the data is not linearly separable.

- Decision Boundaries:

- **LDA** produces linear decision boundaries, which can be limiting if the true boundary between classes is non-linear.
- **QDA** produces quadratic decision boundaries, allowing it to capture more complex patterns in the data.

- Sample Size:

- **LDA** generally requires fewer parameters to estimate (due to the shared covariance matrix), making it more stable with smaller datasets.
- **QDA** requires estimating a separate covariance matrix for each class, which can lead to overfitting if the dataset is small or if there are many features.

- Overfitting:

- **LDA** is less prone to overfitting due to its simpler model structure.
- **QDA** can overfit the training data, especially when the number of features is large relative to the number of samples.

2. K-Nearest Neighbors (KNN) is a non-parametric, instance-based learning algorithm used for classification and regression tasks.

(a) **Role of Distance Metrics:** Performance of KNN is entirely dependent on how we define "closeness". The choice of metric determines which neighbors are selected, which directly impacts the classification outcome. Common distance metrics include Euclidean, Manhattan and Minkowski distances.

- Euclidean: Measures the straight-line distance between two points in a multi-dimensional space. It works best when features are continuous and equally scaled.
- Manhattan: Measures the distance between two points by summing the absolute differences of their coordinates. It is useful when dealing with high-dimensional data or when features are discrete.
- Minkowski: A generalization of both Euclidean and Manhattan distances, controlled by a parameter  $p$ . When  $p = 2$ , it is equivalent to Euclidean distance; when  $p = 1$ , it is equivalent to Manhattan distance.

(b) **Curse of Dimensionality:** As the number of dimensions (features) increases, the volume of the feature space grows exponentially, leading to two major issues for KNN:

- Sparsity: In the high-dimensional space, data points are very spread out. Even the nearest neighbors may be far away, making it difficult to find meaningful relationships.
- Distance Convergence: As dimensions increase, the distances between points tend to converge, meaning that the difference between the nearest and farthest neighbors becomes negligible. This makes it challenging for KNN to distinguish between classes effectively and rendering the concept of "closeness" meaningless.

## Practical

### Imports

```
In [53]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
import numpy as np
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis as QDA
from sklearn.metrics import accuracy_score
from sklearn.neighbors import KNeighborsClassifier as KNN
```

### Load Data

```
In [54]: df = pd.read_csv("../data/Smarket.csv", index_col=0)
```

## Number of features and names

```
In [55]: df_names = df.columns.tolist()
print(f"Number of columns: {len(df_names)}")
print(f"Column names: {df_names}")
```

Number of columns: 9  
 Column names: ['Year', 'Lag1', 'Lag2', 'Lag3', 'Lag4', 'Lag5', 'Volume', 'Today', 'Direction']

## Statistical summary of features

```
In [56]: print(df.describe())
print(df["Direction"].value_counts())
```

	Year	Lag1	Lag2	Lag3	Lag4	\
count	1250.000000	1250.000000	1250.000000	1250.000000	1250.000000	
mean	2003.016000	0.003834	0.003919	0.001716	0.001636	
std	1.409018	1.136299	1.136280	1.138703	1.138774	
min	2001.000000	-4.922000	-4.922000	-4.922000	-4.922000	
25%	2002.000000	-0.639500	-0.639500	-0.640000	-0.640000	
50%	2003.000000	0.039000	0.039000	0.038500	0.038500	
75%	2004.000000	0.596750	0.596750	0.596750	0.596750	
max	2005.000000	5.733000	5.733000	5.733000	5.733000	
	Lag5	Volume	Today			
count	1250.000000	1250.000000	1250.000000			
mean	0.00561	1.478305	0.003138			
std	1.14755	0.360357	1.136334			
min	-4.92200	0.356070	-4.922000			
25%	-0.64000	1.257400	-0.639500			
50%	0.03850	1.422950	0.038500			
75%	0.59700	1.641675	0.596750			
max	5.73300	3.152470	5.733000			
	Direction					
Up	648					
Down	602					
Name:	count, dtype: int64					

## Number of datapoints

```
In [57]: print(f"Number of datapoints: {len(df)}")
```

Number of datapoints: 1250

## Display data in table format

```
In [58]: print(df.head(5))
```

	Year	Lag1	Lag2	Lag3	Lag4	Lag5	Volume	Today	Direction
1	2001	0.381	-0.192	-2.624	-1.055	5.010	1.1913	0.959	Up
2	2001	0.959	0.381	-0.192	-2.624	-1.055	1.2965	1.032	Up
3	2001	1.032	0.959	0.381	-0.192	-2.624	1.4112	-0.623	Down
4	2001	-0.623	1.032	0.959	0.381	-0.192	1.2760	0.614	Up
5	2001	0.614	-0.623	1.032	0.959	0.381	1.2057	0.213	Up

## Correlation matrix

In [59]:

```
sns.heatmap(
    df.drop(columns=["Direction"]).corr(), annot=True, fmt=".2f", linewidths=0.5
)
plt.show()
```



### Interpretation from matrix

The matrix indicates that only `Volume` and `Year` have some correlation, meanwhile the other features seem to be minimally correlated. The value of the correlation between `Volume` and `Year` is 0.54, indicating that increasing the year, the volume tends to increase as well.

**Multicollinearity Analysis:** The correlation between `Volume` and `Year` (0.54) shows a moderate positive relationship. As a rough guideline for initial screening, pairwise correlation thresholds of 0.7-0.8 or higher are often used to flag potential multicollinearity concerns, though this is just a heuristic and not definitive.

The lag variables (`Lag1-Lag5`) show minimal correlation with each other and with `Volume`/`Year`, suggesting that multicollinearity is not a significant concern in this dataset. This is beneficial for model stability, as highly correlated predictors can lead to unstable coefficient estimates and difficulty in interpreting individual predictor effects.

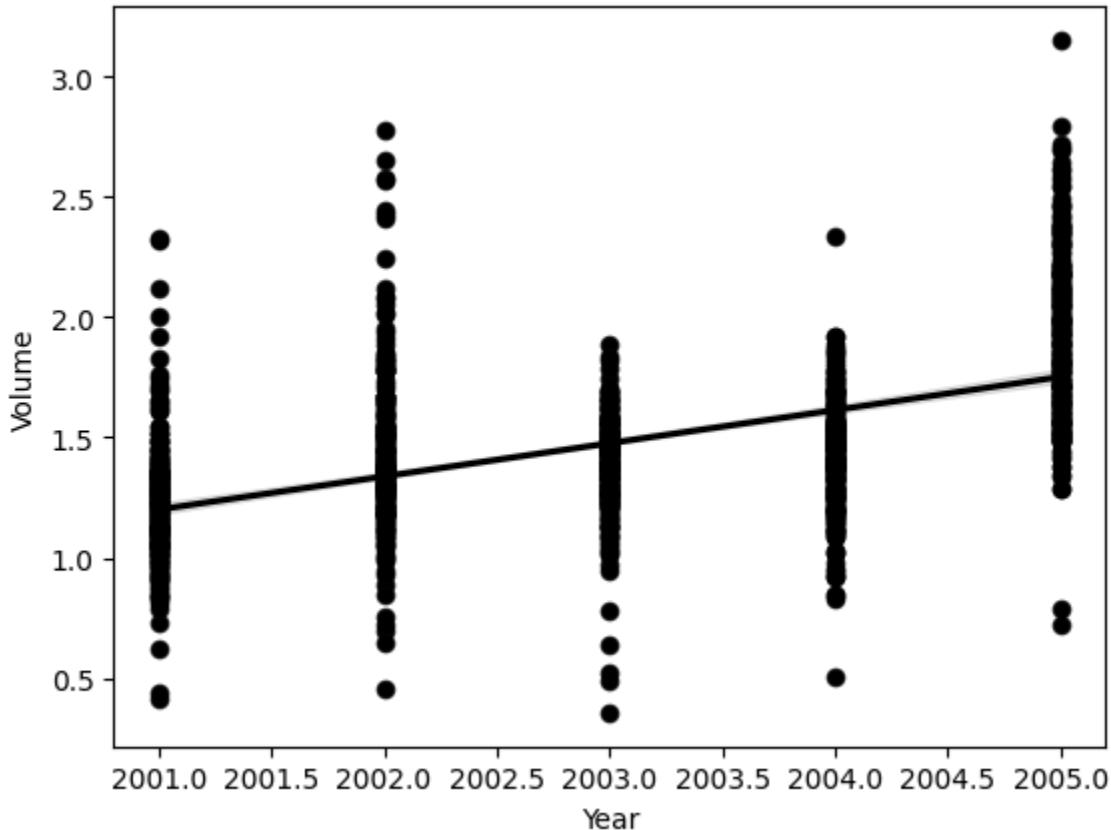
For a more rigorous and reliable assessment of multicollinearity, formal diagnostics such as Variance Inflation Factor (VIF) should be computed. VIF values above 5-10 typically indicate

problematic multicollinearity.

**Relationship with Target:** The correlation matrix only shows relationships between continuous predictors. To understand relationships with the categorical target variable `Direction`, we would need different analyses such as point-biserial correlation or visual exploration through box plots, which are not shown in this correlation matrix.

```
In [60]: sns.scatterplot(data=df, x="Year", y="Volume", color="black")
sns.regplot(data=df, x="Year", y="Volume", color="black")
```

```
Out[60]: <Axes: xlabel='Year', ylabel='Volume'>
```



### Graph interpretation

It reinforces the interpretation drawn from the correlation matrix, showing a positive trend between `Volume` and `Year`.

### Logistics Regression

```
In [61]: X = df[["Lag1", "Lag2", "Lag3", "Lag4", "Lag5", "Volume"]]
X = sm.add_constant(X)
Y = df["Direction"].map({"Up": 1, "Down": 0})
model = sm.Logit(Y, X).fit()

print(model.summary())
```

```

Optimization terminated successfully.
    Current function value: 0.691034
    Iterations 4
                    Logit Regression Results
=====
Dep. Variable:                 Direction      No. Observations:                  1250
Model:                          Logit        Df Residuals:                      1243
Method:                         MLE        Df Model:                           6
Date: Mon, 26 Jan 2026          Pseudo R-squ.:            0.002074
Time: 09:52:19                Log-Likelihood:           -863.79
converged:                      True        LL-Null:             -865.59
Covariance Type:               nonrobust   LLR p-value:            0.7319
=====
                                         coef      std err       z     P>|z|      [0.025      0.975]
-----
const      -0.1260      0.241      -0.523      0.601      -0.598      0.346
Lag1       -0.0731      0.050      -1.457      0.145      -0.171      0.025
Lag2       -0.0423      0.050      -0.845      0.398      -0.140      0.056
Lag3        0.0111      0.050      0.222      0.824      -0.087      0.109
Lag4        0.0094      0.050      0.187      0.851      -0.089      0.107
Lag5        0.0103      0.050      0.208      0.835      -0.087      0.107
Volume      0.1354      0.158      0.855      0.392      -0.175      0.446
=====
```

### Interpretation of regression

The p-values for all coefficients are  $> 0.05$ , indicating that none of the features are statistically significant at the conventional 0.05 significance level for predicting the direction.

**Confidence Intervals:** Examining the 95% confidence intervals (CI) from the regression output would provide additional insight. When a CI for a coefficient includes zero, it indicates that the predictor is not statistically significant, which aligns with the high p-values observed. The width of the confidence intervals also indicates the precision of our coefficient estimates - wider intervals suggest greater uncertainty. In this model, we would expect all confidence intervals to include zero given the insignificant p-values, confirming that none of the predictors have a reliably non-zero effect on the market direction.

**Practical Implication:** The lack of statistical significance across all predictors suggests that past lag values and volume may not be strong predictors of market direction, or that the relationship is too noisy to detect with this model and sample size.

### Use model

```
In [62]: probs = model.predict()

for indecies in range(10):
    print(f"({indecies + 1})\t {probs[indecies]}")
```

```
(1)      0.5070841334630001
(2)      0.48146787817516973
(3)      0.4811388348131379
(4)      0.5152223557927659
(5)      0.5107811625911515
(6)      0.5069564604552573
(7)      0.49265087386752915
(8)      0.5092291581449722
(9)      0.5176135261687921
(10)     0.4888377794207711
```

### Interpretation

This shows probability of the market value going up rather than down, as defined by `df["Direction"].map({"Up": 1, "Down": 0})`. A value > 0.5 indicates that the model predicts an increase in market value, while a value < 0.5 indicates a decrease.

### Confusion Matrix

```
In [63]: pred = []

for prob in probs:
    if prob > 0.5:
        pred.append("Up")
    else:
        pred.append("Down")

print(pd.crosstab(pred, df["Direction"], rownames=["Predicted"], colnames=["Actual"])
accuracy = np.mean(pred == df["Direction"])
print(f"Model accuracy: {accuracy}")
```

Actual	Down	Up
Predicted		
Down	145	141
Up	457	507

Model accuracy: 0.5216

### Interpretation

From the ten previously printed values of `probs`, we assume that the prediction is not confident in its classification, due to it hovering around 0.5. This is reinforced by the model accuracy of 52.16%, which is only slightly better than a coin flip (50%).

```
In [64]: X = df[["Lag1", "Lag2"]]
X = sm.add_constant(X)
model = sm.Logit(Y, X).fit()
probs = model.predict()

pred = []
for prob in probs:
    if prob > 0.5:
        pred.append("Up")
    else:
        pred.append("Down")
```

```
print(pd.crosstab(pred, df["Direction"], rownames=["Predicted"], colnames=["Actual"])
accuracy = np.mean(pred == df["Direction"])
print(f"Model accuracy: {accuracy}")
```

Optimization terminated successfully.  
 Current function value: 0.691361  
 Iterations 4  

Actual	Down	Up
Predicted		
Down	114	102
Up	488	546

 Model accuracy: 0.528

### Interpretation

This model's accuracy is only slightly better by using only two of the predictors (Lag1 and Lag2) instead of all of them. However, it is unclear if this is statistically significant without further testing.

**Why Simpler Model Performs Similarly:** When predictors are not statistically significant (as shown in the full model with p-values > 0.05), including them adds noise rather than signal to the model. The similar or slightly better performance of the reduced model demonstrates that Lag3, Lag4, Lag5, and Volume likely do not contribute meaningful predictive information. This aligns with the principle of parsimony - simpler models are often preferable when they achieve similar performance. Additionally, reducing the number of predictors can help prevent overfitting, especially when working with non-significant variables that may capture random noise in the training data rather than true underlying patterns.

## LDA

```
In [65]: X = df[["Lag1", "Lag2"]]
model = LDA()
model.fit(X, Y)

print(f"Prior probs. of groups: {model.priors_}")
print(f"Group means: \n {model.means_}")
print(f"Coeffs. of linear discriminants: {model.coef_}")

transform = model.transform(X)
LDA_DF = pd.DataFrame({"LDA1": transform[:, 0], "Direction": Y})

figure, axes = plt.subplots(nrows=2, ncols=1)

subset_down = LDA_DF[LDA_DF["Direction"] == 0]
axes[0].hist(subset_down["LDA1"], bins=20, density=True, color="red")
axes[0].set_xlabel("LDA Projection")
axes[0].set_ylabel("Density")
axes[0].set_title("Subset Down")

subset_up = LDA_DF[LDA_DF["Direction"] == 1]
axes[1].hist(subset_up["LDA1"], bins=20, density=True, color="green")
axes[1].set_xlabel("LDA Projection")
```

```

axes[1].set_ylabel("Density")
axes[1].set_title("Subset Up")

plt.tight_layout()
plt.show()

```

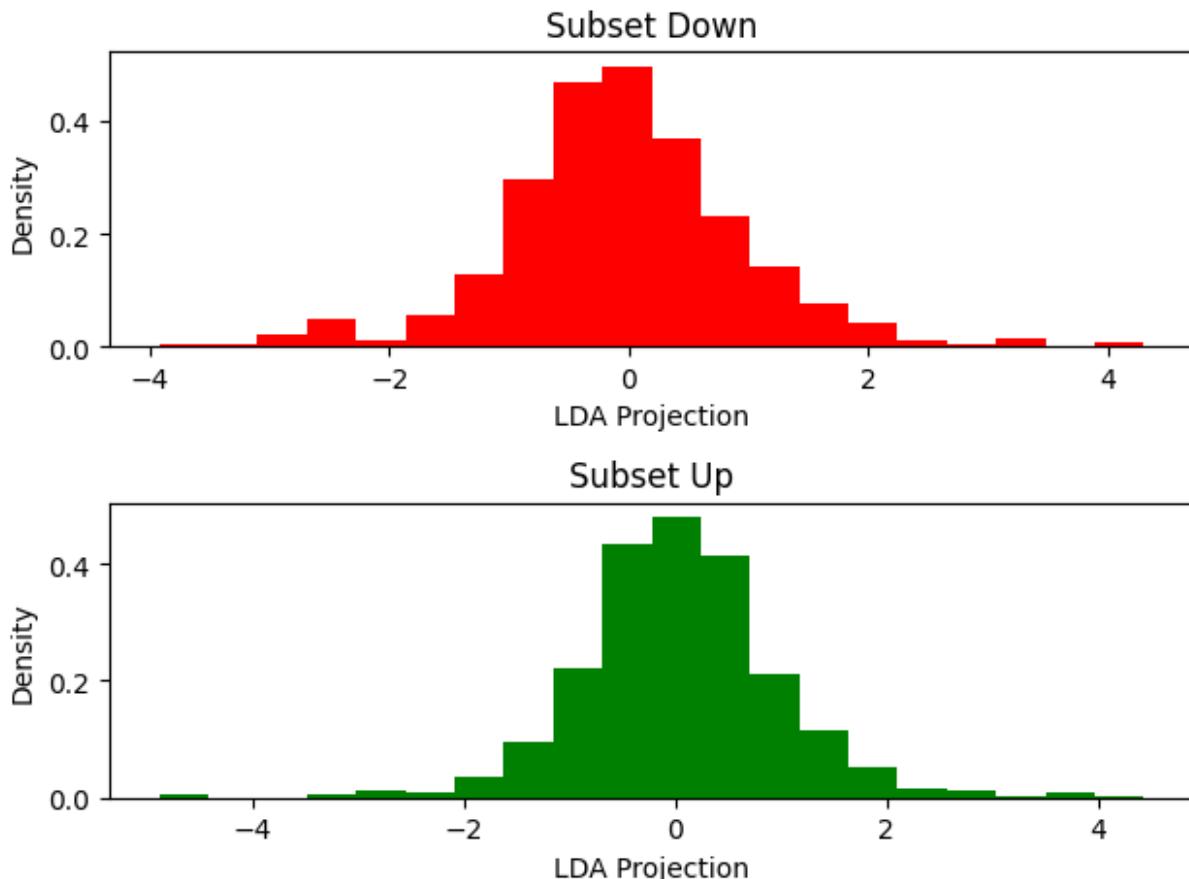
Prior probs. of groups: [0.4816 0.5184]

Group means:

[[ 0.05068605 0.03229734]

[-0.03969136 -0.02244444]]

Coeffs. of linear discriminants: [[-0.07126095 -0.04433204]]



### Interpretation

- Prior probabilities: Informs us that 51.84% of the data belongs to the "Up" class, while 48.16% belongs to the "Down" class.
- Group means: For Lag1 the mean value for the "Down" direction  $\approx 0.05$  and  $\approx 0.03$  for the "Up" direction, while Lag2 has a mean value of  $\approx -0.04$  for "Down" and  $\approx 0.02$  for "Up". This suggests that on average, Lag1 and Lag2 are slightly higher when the market goes "Up".
- Coefficients: Indicate that for an increase in Lag1 and Lag2, will result in a downward trend, as both coefficients are negative.
- Graphs: Show some overlap between the two classes for both Lag1 and Lag2, indicating that it might not be a good predictor of direction.

In [66]: LDA\_class = model.predict(X)

```

conf_matrix = pd.crosstab(LDA_class, Y, rownames=["Predicted"], colnames=["Actual"])
accuracy = accuracy_score(Y, LDA_class)

print(f"Confusion Matrix:\n{conf_matrix}")
print(f"Accuracy: {accuracy}")

```

Confusion Matrix:

Actual	0	1
Predicted		
0	114	102
1	488	546

Accuracy: 0.528

### Interpretation

Same results as for the logistics regression above, which indicates that the LDA does not contribute any more than the logistics regression.

## QDA

```

In [67]: model = QDA()
model.fit(X, Y)

print(f"Prior probs. of groups: {model.priors_}")
print(f"Group means: \n {model.means_}")

```

Prior probs. of groups: [0.4816 0.5184]  
Group means:  
[[ 0.05068605 0.03229734]  
[-0.03969136 -0.02244444]]

### Interpretation

The prior probabilities and group means are the same as for LDA because both models use the same dataset and predictors (Lag1 and Lag2). These statistics are calculated directly from the data and do not depend on the modeling assumptions.

**Key Difference:** While LDA and QDA share the same prior probabilities and group means, they differ in how they model the covariance structure. LDA assumes all classes share a common covariance matrix, while QDA allows each class to have its own covariance matrix. This means QDA can capture different variance patterns for the "Up" and "Down" classes, potentially leading to better classification if the classes have different spread or correlation patterns between predictors.

## Use model

```

In [68]: QDA_class = model.predict(X)
conf_matrix = pd.crosstab(QDA_class, Y, rownames=["Predicted"], colnames=["Actual"])
accuracy = accuracy_score(Y, QDA_class)

print(f"Confusion Matrix:\n{conf_matrix}")
print(f"Accuracy: {accuracy}")

```

```
Confusion Matrix:
Actual      0      1
Predicted
0          109    94
1          493   554
Accuracy: 0.5304
```

### Interpretation

The QDA model achieves an improved accuracy of 53.04%, which is expected due to the increased number of parameters that QDA can utilize compared to LDA.

**Why QDA Has More Parameters:** LDA estimates a single shared covariance matrix for all classes, while QDA estimates a separate covariance matrix for each class. For  $K$  classes and  $p$  predictors:

- LDA estimates:  $p(p + 1)/2$  covariance parameters (shared across classes)
- QDA estimates:  $K \times p(p + 1)/2$  covariance parameters (one per class)

In our case with 2 classes (Up/Down) and 2 predictors (Lag1, Lag2), QDA estimates twice as many covariance parameters as LDA. This increased flexibility allows QDA to model different variance-covariance structures for each class, which can lead to better classification when such differences exist in the data. However, this also means QDA is more prone to overfitting, especially with small sample sizes or many predictors, which is why the improvement should be validated on a test set.

## KNN Clustering

```
In [69]: train = df["Year"] < 2005
print(df[~train].shape)

(252, 9)
```

```
In [70]: train_X = df[["Lag1", "Lag2"]][train]
train_direction = df["Direction"][train]

test_X = df[["Lag1", "Lag2"]][~train]
test_direction = df["Direction"][~train]

KNN_model = KNN(n_neighbors=1)
KNN_model.fit(train_X, train_direction)
KNN_pred = KNN_model.predict(test_X)

conf_matrix = pd.crosstab(
    KNN_pred, test_direction, rownames=["Predicted"], colnames=["Actual"]
)
accuracy = accuracy_score(test_direction, KNN_pred)

print(f"Confusion Matrix:\n{conf_matrix}")
print(f"Accuracy: {accuracy}")
```

```
Confusion Matrix:  
Actual      Down   Up  
Predicted  
Down         43   58  
Up          68   83  
Accuracy: 0.5
```

```
In [71]: KNN_model = KNN(n_neighbors=3)  
KNN_model.fit(train_X, train_direction)  
KNN_pred = KNN_model.predict(test_X)  
  
conf_matrix = pd.crosstab(  
    KNN_pred, test_direction, rownames=["Predicted"], colnames=["Actual"]  
)  
accuracy = accuracy_score(test_direction, KNN_pred)  
  
print(f"Confusion Matrix:\n{conf_matrix}")  
print(f"Accuracy: {accuracy}")
```

```
Confusion Matrix:  
Actual      Down   Up  
Predicted  
Down         48   55  
Up          63   86  
Accuracy: 0.5317460317460317
```

### Interpretation

The value of  $k$  seems to have a decent impact on the accuracy, with  $k = 3$  yielding the highest accuracy so far of 53.17%. However, it is still unclear if this is statistically significant without further testing.