

January 26, 2026

1 Assignment 5

1.1 The Bootstrap

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-26

Course: Machine Learning 4DT905

1.1.1 Conceptual

1. Shuffle the given dataset to ensure randomness. Partition the dataset into k equal-sized groups (folds). For fold k , take the fold as *test* dataset and the remaining $k - 1$ folds as *training* dataset. Fit a model on the *training* dataset and evaluate it with the *test* dataset. Retain the resulting metrics and discard model. Compute the average of the retained metrics to estimate the models true performance.
2. *i.* **Validation set approach** is a single split of the dataset into 70% training and 30% testing. This results in a higher variance, lower data efficiency and a lower computational cost than the k-fold cross validation.
3. *ii.* **Leave-one-out cross validation (LOOCV)** is a special case of k-fold cross validation where $k = n$ where n is the number of observations. This results in a higher computational cost, lower bias and higher variance than the k-fold cross validation.

1.1.2 Practical

Imports

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import statsmodels.api as sm
import numpy as np
```

Load data

```
[2]: df = pd.read_csv("../data/Auto.csv", index_col=0)
```

Number of features and names

```
[3]: df_names = df.columns.tolist()
print(f"Number of columns: {len(df_names)}")
print(f"Column names: {df_names}")
```

```
Number of columns: 9
Column names: ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
'acceleration', 'year', 'origin', 'name']
```

Statistical summary of features

```
[5]: print(df.describe())
print(df["name"].value_counts())
```

```
          mpg    cylinders  displacement  horsepower      weight \
count  392.000000  392.000000  392.000000  392.000000  392.000000
mean   23.445918    5.471939   194.411990   104.469388  2977.584184
std    7.805007    1.705783   104.644004    38.491160   849.402560
min    9.000000    3.000000   68.000000   46.000000  1613.000000
25%   17.000000    4.000000  105.000000   75.000000  2225.250000
50%   22.750000    4.000000  151.000000   93.500000  2803.500000
75%   29.000000    8.000000  275.750000  126.000000  3614.750000
max   46.600000    8.000000  455.000000  230.000000  5140.000000

          acceleration        year       origin
count  392.000000  392.000000  392.000000
mean   15.541327   75.979592   1.576531
std    2.758864   3.683737   0.805518
min    8.000000   70.000000   1.000000
25%   13.775000   73.000000   1.000000
50%   15.500000   76.000000   1.000000
75%   17.025000   79.000000   2.000000
max   24.800000   82.000000   3.000000

name
amc matador      5
ford pinto       5
toyota corolla   5
chevrolet impala 4
amc hornet       4
..
ford mustang gl  1
vw pickup        1
dodge rampage     1
ford ranger       1
chevy s-10        1
Name: count, Length: 301, dtype: int64
```

Number of datapoints

```
[6]: print(f"Number of datapoints: {len(df)}")
```

Number of datapoints: 392

Display data in table format

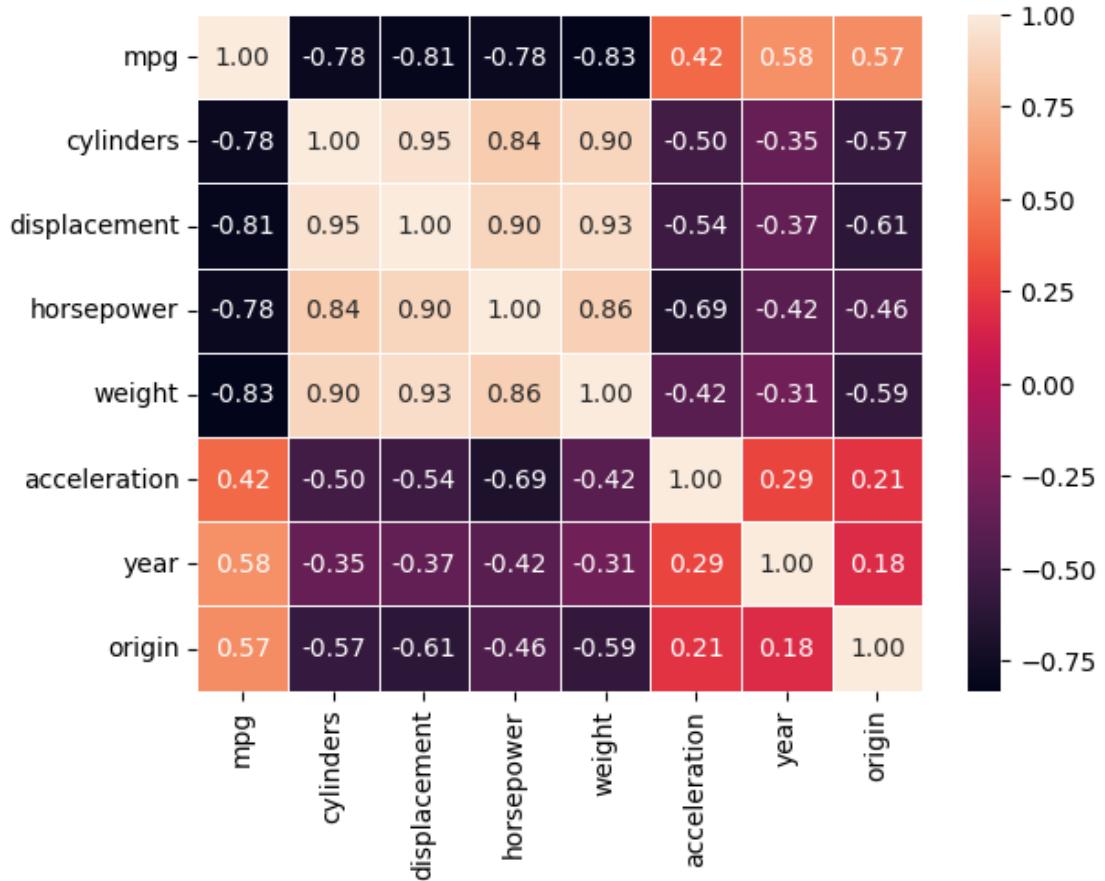
[7]: `print(df.head(5))`

```
mpg cylinders displacement horsepower weight acceleration year \
1 18.0          8         307.0        130    3504       12.0      70
2 15.0          8         350.0        165    3693       11.5      70
3 18.0          8         318.0        150    3436       11.0      70
4 16.0          8         304.0        150    3433       12.0      70
5 17.0          8         302.0        140    3449       10.5      70

origin           name
1      1  chevrolet chevelle malibu
2      1          buick skylark 320
3      1  plymouth satellite
4      1          amc rebel sst
5      1          ford torino
```

Correlation matrix

[9]: `sns.heatmap(df.drop(columns=["name"]).corr(), annot=True, fmt=".2f", linewidths=0.5)`
`plt.show()`



Accuracy Estimation Function

```
[ ]: def boot_fn(data, index):
    sample = data.iloc[index]
    X = sample["horsepower"]
    Y = sample["mpg"]
    X = sm.add_constant(X)
    model = sm.OLS(Y, X).fit()

    return model.params

print(boot_fn(df, range(392)))
```

```
const          39.935861
horsepower     -0.157845
dtype: float64
```

```
[14]: np.random.seed(42)
```

```
print(boot_fn(df, np.random.choice(392, 392, replace=True)))
```

```
const      40.466879  
horsepower -0.163738  
dtype: float64
```

```
[ ]: boot_results = np.zeros((1000, 2))  
for idx in range(1000):  
    indices = np.random.choice(392, 392, replace=True)  
    boot_results[idx, :] = boot_fn(df, indices)  
  
print(f"Standard errors: {boot_results.std(axis=0)}")
```

```
Standard errors: [0.86861119 0.00749939]
```

```
[17]: X = df["horsepower"]  
Y = df["mpg"]  
X = sm.add_constant(X)  
model = sm.OLS(Y, X).fit()  
print(model.params)  
print(model.summary())
```

```
const      39.935861  
horsepower -0.157845  
dtype: float64
```

OLS Regression Results

```
=====
```

Dep. Variable:	mpg	R-squared:	0.606			
Model:	OLS	Adj. R-squared:	0.605			
Method:	Least Squares	F-statistic:	599.7			
Date:	Mon, 26 Jan 2026	Prob (F-statistic):	7.03e-81			
Time:	10:17:19	Log-Likelihood:	-1178.7			
No. Observations:	392	AIC:	2361.			
Df Residuals:	390	BIC:	2369.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	39.9359	0.717	55.660	0.000	38.525	41.347
horsepower	-0.1578	0.006	-24.489	0.000	-0.171	-0.145

```
=====
```

```
Omnibus:          16.432   Durbin-Watson:        0.920  
Prob(Omnibus):   0.000    Jarque-Bera (JB):    17.305  
Skew:            0.492    Prob(JB):           0.000175  
Kurtosis:         3.299   Cond. No.          322.  
=====
```

Notes:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

Interpretation The standard error from the model summary is a smaller span compared to the bootstrap estimate. This is expected as the bootstrap method since it is a estimate over several subsets of the data. Meaning currently with the values given the regular estimate is better but most likely it is a bit overfitted.

```
[18]: def boot_fn_quadratic(data, index):
    sample = data.iloc[index]
    X = sample["horsepower"]
    Y = sample["mpg"]
    X = np.column_stack((X, X**2))
    X = sm.add_constant(X)
    model = sm.OLS(Y, X).fit()

    return model.params
```

```
[20]: boot_results = np.zeros((1000, 3))
for idx in range(1000):
    indices = np.random.choice(392, 392, replace=True)
    boot_results[idx, :] = boot_fn_quadratic(df, indices)

print(f"Standard errors: {boot_results.std(axis=0)}")
```

```
Standard errors: [2.14866982e+00 3.42412395e-02 1.23413967e-04]
```

```
[21]: df["horsepower_squared"] = df["horsepower"] ** 2
X = df[["horsepower", "horsepower_squared"]]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()

print(model.params)
print(model.summary())
```

```
const          56.900100
horsepower     -0.466190
horsepower_squared   0.001231
dtype: float64
```

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.688
Model:	OLS	Adj. R-squared:	0.686
Method:	Least Squares	F-statistic:	428.0
Date:	Mon, 26 Jan 2026	Prob (F-statistic):	5.40e-99
Time:	10:26:43	Log-Likelihood:	-1133.2
No. Observations:	392	AIC:	2272.
Df Residuals:	389	BIC:	2284.
Df Model:	2		

```

Covariance Type: nonrobust
=====
=====

            coef    std err          t      P>|t|      [0.025
0.975]
-----
const        56.9001     1.800     31.604      0.000     53.360
60.440
horsepower   -0.4662     0.031    -14.978      0.000    -0.527
-0.405
horsepower_squared  0.0012     0.000     10.080      0.000     0.001
0.001
=====
Omnibus:           16.158  Durbin-Watson:           1.078
Prob(Omnibus):    0.000  Jarque-Bera (JB):       30.662
Skew:             0.218  Prob(JB):                 2.20e-07
Kurtosis:         4.299  Cond. No.                1.29e+05
=====
```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.29e+05. This might indicate that there are strong multicollinearity or other numerical problems.

Interpretation Again the standard error from the model summary is a smaller span compared to the bootstrap estimate. This is due to the same reason as in the linear case.