

Assignment 5

The Bootstrap

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-26

Course: Machine Learning 4DT905

Conceptual

1. Shuffle the given dataset to ensure randomness. Partition the dataset into k equal-sized groups (folds). For each fold i (where $i = 1, \dots, k$), take fold i as the *test* dataset and the remaining $k - 1$ folds as the *training* dataset. Fit a model on the *training* dataset and evaluate it on the *test* dataset. Retain the resulting metric (e.g., MSE or accuracy) and discard the model. After repeating this for all k folds, compute the average of the k retained metrics to estimate the model's true performance. This ensures every observation is used for both training and testing exactly once.
2. *i. Validation set approach* uses a single split of the dataset (e.g., 70% training and 30% testing). Compared to k-fold cross-validation:
 - **Advantage:** Lower computational cost — only one model is trained.
 - **Disadvantages:** Higher variance in the test error estimate (the result depends heavily on which observations happen to land in the test set); lower data efficiency since a portion of data is never used for training; the test error may be overestimated because less data is available for training.
3. *ii. Leave-one-out cross validation (LOOCV)* is a special case of k-fold cross-validation where $k = n$ (the number of observations). Each model is trained on $n - 1$ observations and tested on the single left-out observation. Compared to k-fold cross-validation:
 - **Advantage:** Lower bias — the training set is nearly the full dataset, so the model performance estimate is less pessimistic.
 - **Disadvantages:** Much higher computational cost — n models must be trained; higher variance in the estimate because each test set contains only one observation, making individual evaluations noisy.

Practical

Imports

```
In [1]: import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
import statsmodels.api as sm  
import numpy as np
```

Load data

```
In [2]: df = pd.read_csv("../data/Auto.csv", index_col=0)
```

Number of features and names

```
In [3]: df_names = df.columns.tolist()  
print(f"Number of columns: {len(df_names)}")  
print(f"Column names: {df_names}")
```

```
Number of columns: 9  
Column names: ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin', 'name']
```

Statistical summary of features

```
In [5]: print(df.describe())  
print(df["name"].value_counts())
```

```
          mpg   cylinders displacement horsepower      weight \
count    392.000000  392.000000    392.000000  392.000000  392.000000
mean     23.445918    5.471939   194.411990  104.469388 2977.584184
std      7.805007    1.705783   104.644004   38.491160  849.402560
min      9.000000    3.000000    68.000000   46.000000 1613.000000
25%     17.000000    4.000000   105.000000   75.000000 2225.250000
50%     22.750000    4.000000   151.000000  93.500000 2803.500000
75%     29.000000    8.000000   275.750000  126.000000 3614.750000
max     46.600000    8.000000   455.000000  230.000000 5140.000000

          acceleration       year      origin
count    392.000000  392.000000  392.000000
mean     15.541327   75.979592   1.576531
std      2.758864   3.683737   0.805518
min      8.000000   70.000000   1.000000
25%     13.775000   73.000000   1.000000
50%     15.500000   76.000000   1.000000
75%     17.025000   79.000000   2.000000
max     24.800000   82.000000   3.000000

name
amc matador           5
ford pinto             5
toyota corolla         5
chevrolet impala      4
amc hornet             4
..
ford mustang gl        1
vw pickup               1
dodge rampage            1
ford ranger              1
chevy s-10                1
Name: count, Length: 301, dtype: int64
```

Number of datapoints

```
In [6]: print(f"Number of datapoints: {len(df)}")
```

```
Number of datapoints: 392
```

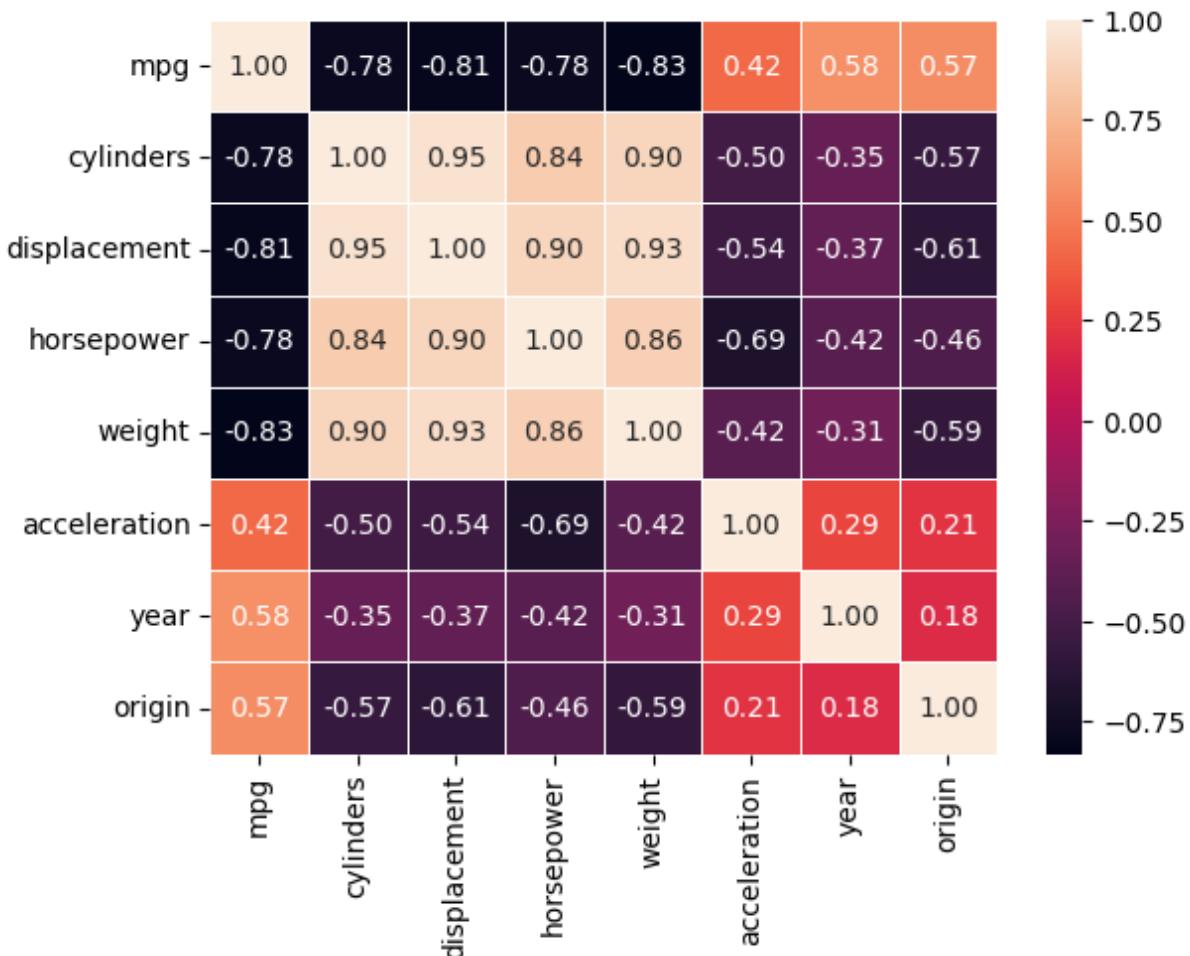
Display data in table format

In [7]: `print(df.head(5))`

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
1	18.0	8	307.0	130	3504	12.0	70	
2	15.0	8	350.0	165	3693	11.5	70	
3	18.0	8	318.0	150	3436	11.0	70	
4	16.0	8	304.0	150	3433	12.0	70	
5	17.0	8	302.0	140	3449	10.5	70	
	origin			name				
1	1	chevrolet	chevelle	malibu				
2	1	buick	skylark	320				
3	1	plymouth	satellite					
4	1	amc	rebel	sst				
5	1	ford	torino					

Correlation matrix

In [9]: `sns.heatmap(df.drop(columns=["name"]).corr(), annot=True, fmt=".2f", linewidths=0.5)`



Accuracy Estimation Function

In []: `def boot_fn(data, index):
 sample = data.iloc[index]`

```
X = sample["horsepower"]
Y = sample["mpg"]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()

return model.params

print(boot_fn(df, range(392)))
```

```
const      39.935861
horsepower -0.157845
dtype: float64
```

```
In [14]: np.random.seed(42)
```

```
print(boot_fn(df, np.random.choice(392, 392, replace=True)))

const      40.466879
horsepower -0.163738
dtype: float64
```

```
In [ ]: boot_results = np.zeros((1000, 2))
for idx in range(1000):
    indices = np.random.choice(392, 392, replace=True)
    boot_results[idx, :] = boot_fn(df, indices)

print(f"Standard errors: {boot_results.std(axis=0)}")
```

```
Standard errors: [0.86861119 0.00749939]
```

```
In [17]: X = df["horsepower"]
Y = df["mpg"]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()
print(model.params)
print(model.summary())
```

```
const      39.935861
horsepower -0.157845
dtype: float64
```

OLS Regression Results

Dep. Variable:	mpg	R-squared:	0.606
Model:	OLS	Adj. R-squared:	0.605
Method:	Least Squares	F-statistic:	599.7
Date:	Mon, 26 Jan 2026	Prob (F-statistic):	7.03e-81
Time:	10:17:19	Log-Likelihood:	-1178.7
No. Observations:	392	AIC:	2361.
Df Residuals:	390	BIC:	2369.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	39.9359	0.717	55.660	0.000	38.525	41.347
horsepower	-0.1578	0.006	-24.489	0.000	-0.171	-0.145

Omnibus:	16.432	Durbin-Watson:	0.920
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17.305
Skew:	0.492	Prob(JB):	0.000175
Kurtosis:	3.299	Cond. No.	322.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Interpretation

Comparing the standard errors (SE) for the linear model (`mpg ~ horsepower`) between the two methods:

Parameter	Bootstrap SE	OLS SE
Intercept ($\hat{\beta}_0$)	~0.84	~0.72
Horsepower ($\hat{\beta}_1$)	~0.0073	~0.0064

The OLS standard errors are smaller than the bootstrap standard errors. However, this does **not** mean the OLS estimates are more precise, rather it reflects that OLS relies on strong assumptions about the error distribution that may not hold here.

Since the true relationship between `mpg` and `horsepower` is likely nonlinear, the linear model is misspecified. When model assumptions are violated, OLS tends to **underestimate** the true standard errors, giving a false sense of precision. The bootstrap, being assumption-free and based on resampling the actual data, provides a more honest and reliable estimate of variability.

Concluding, the bootstrap SE is the more trustworthy estimate here, and the discrepancy indicates that the linear model's assumptions are not fully satisfied.

```
In [18]: def boot_fn_quadratic(data, index):
    sample = data.iloc[index]
    X = sample["horsepower"]
    Y = sample["mpg"]
    X = np.column_stack((X, X**2))
    X = sm.add_constant(X)
    model = sm.OLS(Y, X).fit()

    return model.params
```

```
In [20]: boot_results = np.zeros((1000, 3))
for idx in range(1000):
    indices = np.random.choice(392, 392, replace=True)
    boot_results[idx, :] = boot_fn_quadratic(df, indices)

print(f"Standard errors: {boot_results.std(axis=0)}")
```

Standard errors: [2.14866982e+00 3.42412395e-02 1.23413967e-04]

```
In [21]: df["horsepower_squared"] = df["horsepower"] ** 2
X = df[["horsepower", "horsepower_squared"]]
X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()

print(model.params)
print(model.summary())
```

```

const           56.900100
horsepower      -0.466190
horsepower_squared   0.001231
dtype: float64

OLS Regression Results
=====
Dep. Variable:          mpg    R-squared:       0.688
Model:                 OLS     Adj. R-squared:  0.686
Method:                Least Squares  F-statistic:    428.0
Date: Mon, 26 Jan 2026  Prob (F-statistic): 5.40e-99
Time: 10:26:43          Log-Likelihood:   -1133.2
No. Observations:      392    AIC:             2272.
Df Residuals:          389    BIC:             2284.
Df Model:                  2
Covariance Type:        nonrobust
=====

coefs      std err      t      P>|t|      [0.025      0.97
5]
-----
-- const      56.9001    1.800    31.604    0.000    53.360    60.4
40
horsepower  -0.4662    0.031   -14.978    0.000   -0.527   -0.4
05
horsepower_squared  0.0012    0.000    10.080    0.000    0.001    0.0
01
=====
Omnibus:            16.158  Durbin-Watson:      1.078
Prob(Omnibus):      0.000  Jarque-Bera (JB):  30.662
Skew:               0.218  Prob(JB):        2.20e-07
Kurtosis:            4.299  Cond. No.:      1.29e+05
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.29e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```

Interpretation

Comparing the standard errors (SE) for the quadratic model ($\text{mpg} \sim \text{horsepower} + \text{horsepower}^2$) between the two methods:

Parameter	Bootstrap SE	OLS SE
Intercept ($\hat{\beta}_0$)	~2.03	~1.80
Horsepower ($\hat{\beta}_1$)	~0.0324	~0.0311
Horsepower ² ($\hat{\beta}_2$)	~0.000117	~0.000122

In contrast to the linear model, the bootstrap and OLS standard errors are now **much closer** to each other. This is because the quadratic model provides a better fit to the true nonlinear

relationship between `mpg` and `horsepower`, meaning the model assumptions (linearity, homoscedasticity) are better satisfied.

When a model is correctly specified and its assumptions hold, the OLS standard error formulas yield estimates that agree with the assumption-free bootstrap estimates. The close correspondence here validates that the quadratic model is a more appropriate representation of the underlying data.

Concluding, the better agreement between bootstrap and OLS SE in the quadratic case, compared to the wider discrepancy in the linear case, further confirms that the quadratic model is the better fit.