

4dt905_sb224sc_A6

January 26, 2026

1 Assignment 6

1.1 Tree-based approaches

Author: Samuel Fredric Berg

Student ID: sb224sc

Date: 2026-01-26

Course: Machine Learning 4DT905

1.1.1 Conceptual

1. **Bagging** focuses on reducing the variance. Trains several models in parallel and independently on different subsets of the data. These trees are deep, complex and averaging them to “smooth out” the overfit. Meanwhile **Boosting** focuses on reducing bias. Trains models sequentially, where each model tries to correct the errors of the previous one. Also uses shallow trees and stack them to learn complex patterns with time. By creating a forest of slightly different trees cancels out the noise in the data. Using **Bootstrapping** to ensure diversity among the trees.

2.

- **Random Forests:**
 - Parallel training
 - Equal weighting
 - Every tree get one equal vote
 - Very robust to overfitting
- **AdaBoost:**
 - Sequential training
 - Dynamic weighting, higherweight prioritized next iteration
 - More accurate tree the more influence
 - Sensitive to noisy data and outliers

1.1.2 Practical

Imports

```
[3]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```

import numpy as np
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import confusion_matrix, mean_squared_error, r2_score
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor

```

Load Data

```
[4]: df = pd.read_csv("../data/Carseats.csv", index_col=0)
```

Number of features and names

```
[5]: df_names = df.columns.tolist()
print(f"Number of columns: {len(df_names)}")
print(f"Column names: {df_names}")
```

Number of columns: 11
 Column names: ['Sales', 'CompPrice', 'Income', 'Advertising', 'Population',
 'Price', 'ShelveLoc', 'Age', 'Education', 'Urban', 'US']

Statistical summary of features

```
[6]: print(df.describe())
print(df[["ShelveLoc"]].value_counts())
print(df[["Urban"]].value_counts())
print(df[["US"]].value_counts())
```

	Sales	CompPrice	Income	Advertising	Population	\
count	400.000000	400.000000	400.000000	400.000000	400.000000	
mean	7.496325	124.975000	68.657500	6.635000	264.840000	
std	2.824115	15.334512	27.986037	6.650364	147.376436	
min	0.000000	77.000000	21.000000	0.000000	10.000000	
25%	5.390000	115.000000	42.750000	0.000000	139.000000	
50%	7.490000	125.000000	69.000000	5.000000	272.000000	
75%	9.320000	135.000000	91.000000	12.000000	398.500000	
max	16.270000	175.000000	120.000000	29.000000	509.000000	

	Price	Age	Education
count	400.000000	400.000000	400.000000
mean	115.795000	53.322500	13.900000
std	23.676664	16.200297	2.620528
min	24.000000	25.000000	10.000000
25%	100.000000	39.750000	12.000000
50%	117.000000	54.500000	14.000000
75%	131.000000	66.000000	16.000000
max	191.000000	80.000000	18.000000

ShelveLoc

Medium	219
Bad	96
Good	85

```
Name: count, dtype: int64
Urban
Yes      282
No       118
Name: count, dtype: int64
US
Yes      257
No       142
Yes       1
Name: count, dtype: int64
```

Number of datapoints

```
[7]: print(f"Number of datapoints: {len(df)}")
```

```
Number of datapoints: 400
```

Display data in table format

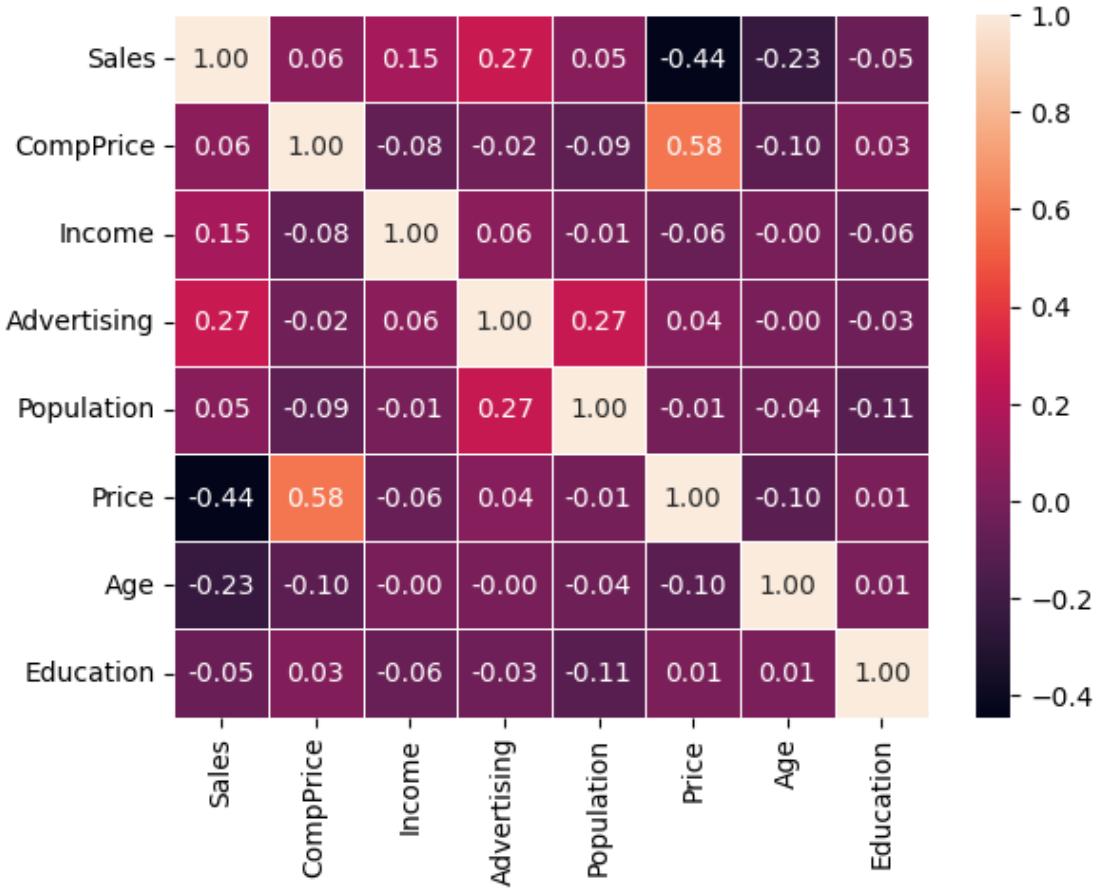
```
[8]: print(df.head(5))
```

	Sales	CompPrice	Income	Advertising	Population	Price	ShelveLoc	Age	\
1	9.50	138	73	11	276	120	Bad	42	
2	11.22	111	48	16	260	83	Good	65	
3	10.06	113	35	10	269	80	Medium	59	
4	7.40	117	100	4	466	97	Medium	55	
5	4.15	141	64	3	340	128	Bad	38	

	Education	Urban	US
1	17	Yes	Yes
2	10	Yes	Yes
3	12	Yes	Yes
4	14	Yes	Yes
5	13	Yes	No

Correlation matrix

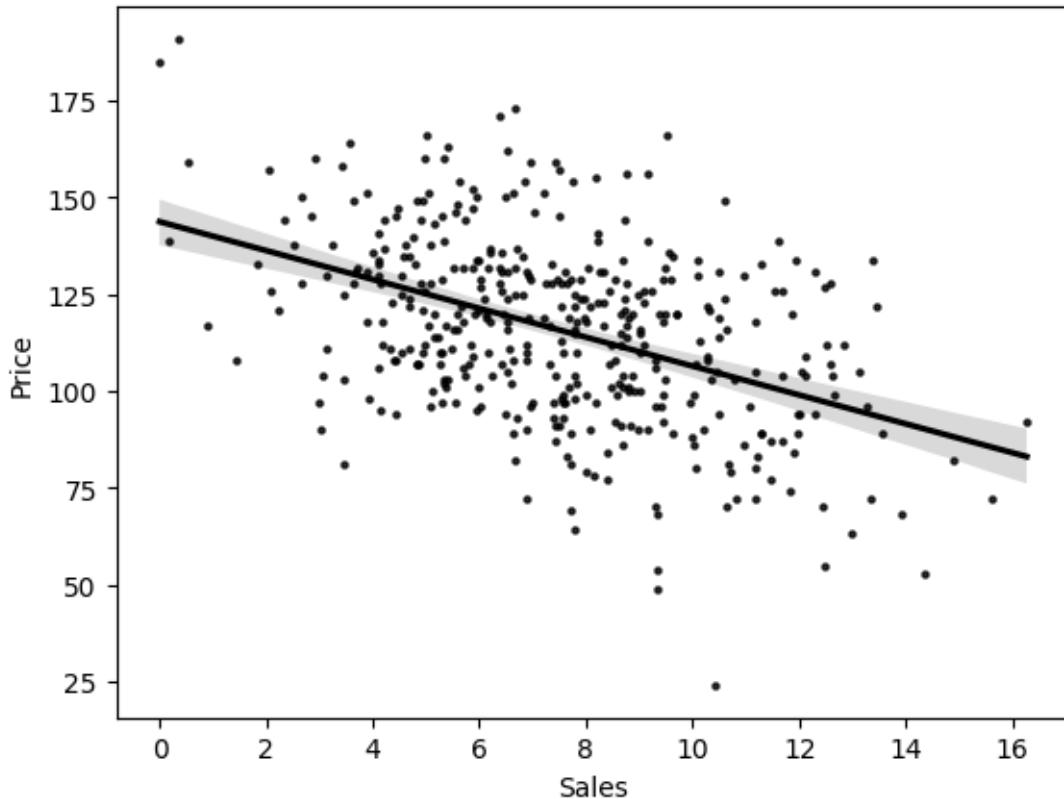
```
[9]: sns.heatmap(
    df.drop(columns=["ShelveLoc", "Urban", "US"]).corr(),
    annot=True,
    fmt=".2f",
    linewidths=0.5,
)
plt.show()
```



Scatter plot

```
[10]: sns.regplot(
    x="Sales",
    y="Price",
    data=df,
    scatter_kws={"s": 5, "color": "black"},
    line_kws={"color": "black"},  

)
plt.show()
```



Interpretation We are looking at “Sales” as our response variable and the most correlated feature is “Price”. The correlation is negative, meaning that as the price increases, the sales decrease. Also means that is the price decreases the sales increase. The second most correlated feature is “Advertising”, which has a positive correlation with sales. This means that as advertising increases, sales also increase. The least correlated features are “Population” and “Education”.

```
[11]: df_encoded = pd.get_dummies(df, columns=["ShelveLoc", "Urban", "US"])
df["High"] = ["No" if sales <= 8 else "Yes" for sales in df["Sales"]]

X = df_encoded.drop(columns=["Sales"])
Y = df["High"]

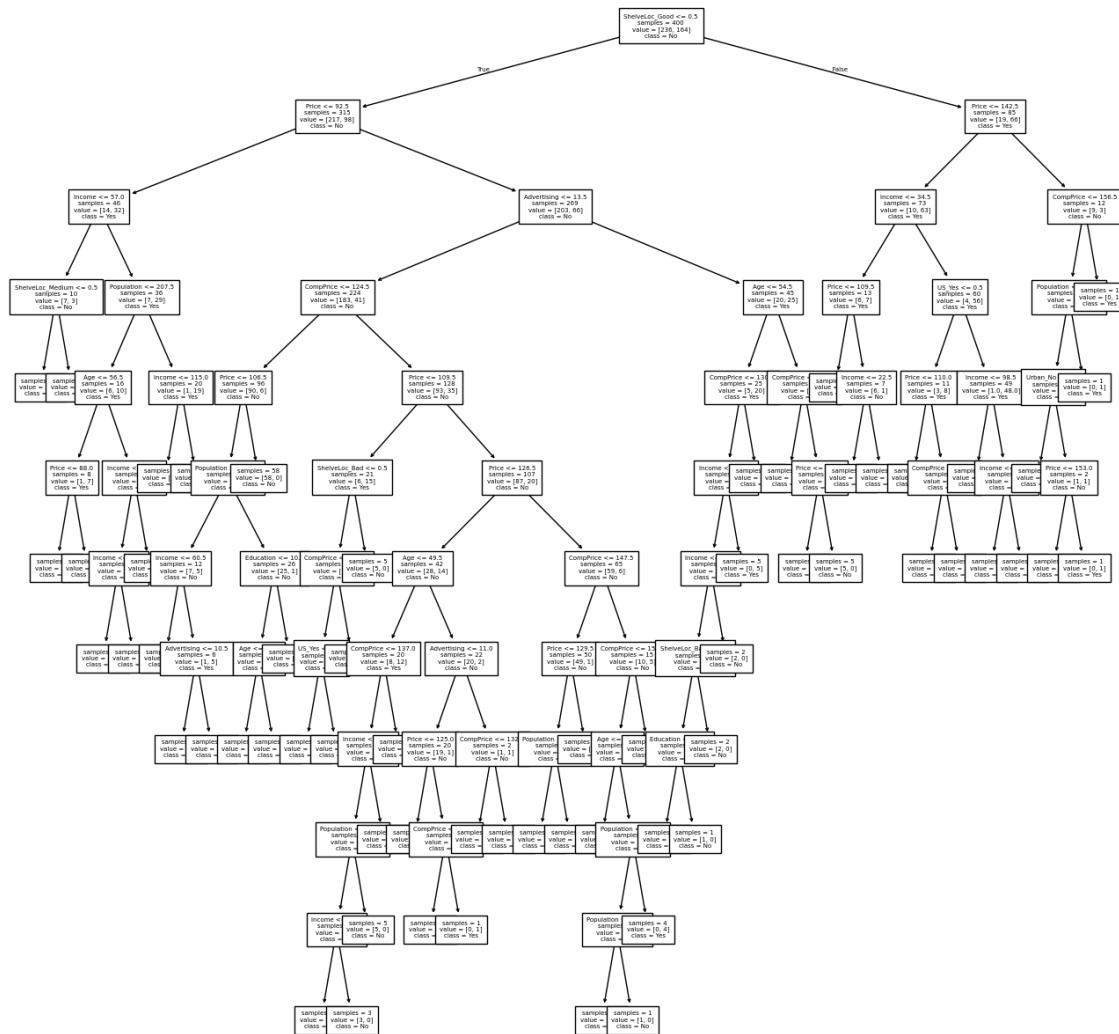
tree = DecisionTreeClassifier()
tree.fit(X, Y)

print(f"Number of leaves: {tree.get_n_leaves()}")
```

Number of leaves: 61

```
[20]: plt.figure(figsize=(16, 16))
plot_tree(
```

```
tree,
feature_names=X.columns,
class_names=["No", "Yes"],
impurity=False,
fontsize=5,
)
plt.show()
```



Interpretation Each node in the tree represents a feature and a threshold value, if the threshold is met the left branch is followed, otherwise the right branch. When it reaches the bottom of the tree it returns an answer (if “High” is Yes or No).

```
[21]: df_encoded = pd.get_dummies(df, columns=["ShelveLoc", "Urban", "US"],  
    ↪drop_first=True)  
X = df_encoded.drop(columns=["Sales", "High"])  
Y = df["High"]  
  
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.5,  
    ↪random_state=1)  
  
tree_model = DecisionTreeClassifier(random_state=1)  
tree_model.fit(X_train, Y_train)  
Y_pred = tree_model.predict(X_test)  
  
conf_matrix = confusion_matrix(Y_test, Y_pred)  
  
print(conf_matrix)
```

```
[[94 25]  
 [30 51]]
```

Interpretation

$$\begin{bmatrix} \text{True Negative (TN)} & \text{False Positive (FP)} \\ \text{False Negative (FN)} & \text{True Positive (TP)} \end{bmatrix} = \begin{bmatrix} 94 & 25 \\ 30 & 51 \end{bmatrix}$$

$$Precision = \frac{TP}{TP + FP} = \frac{51}{51 + 25} \approx 0.6710$$

$$Recall = \frac{TP}{TP + FN} = \frac{51}{51 + 30} \approx 0.6296$$

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} = \frac{51 + 94}{51 + 94 + 25 + 30} \approx 0.7250$$

$$F1 = 2 \cdot \frac{Precision \cdot Recall}{Precision + Recall} = 2 \cdot \frac{0.6710 \cdot 0.6296}{0.6710 + 0.6296} \approx 0.6496$$

This indicates a precision of 67.10%, a recall of 62.96%, an accuracy of 72.50% and this gives us an F1-score of 64.96%.

```
[26]: path = tree_model.cost_complexity_pruning_path(X_train, Y_train)  
  
cv_scores = []  
best_accuracy = 0.0  
best_alpha = 0.0  
for alpha in path ccp_alphas:  
    pruned_tree = DecisionTreeClassifier(random_state=0, ccp_alpha=alpha)  
    scores = cross_val_score(pruned_tree, X_train, Y_train, cv=5)  
    accuracy = scores.mean()  
    cv_scores.append(accuracy)
```

```

if accuracy > best_accuracy:
    best_accuracy = accuracy
    best_alpha = alpha

print(f"Best accuracy: {best_accuracy:.4f}, with alpha: {best_alpha:.4f}")

```

Best accuracy: 0.7100, with alpha: 0.0128

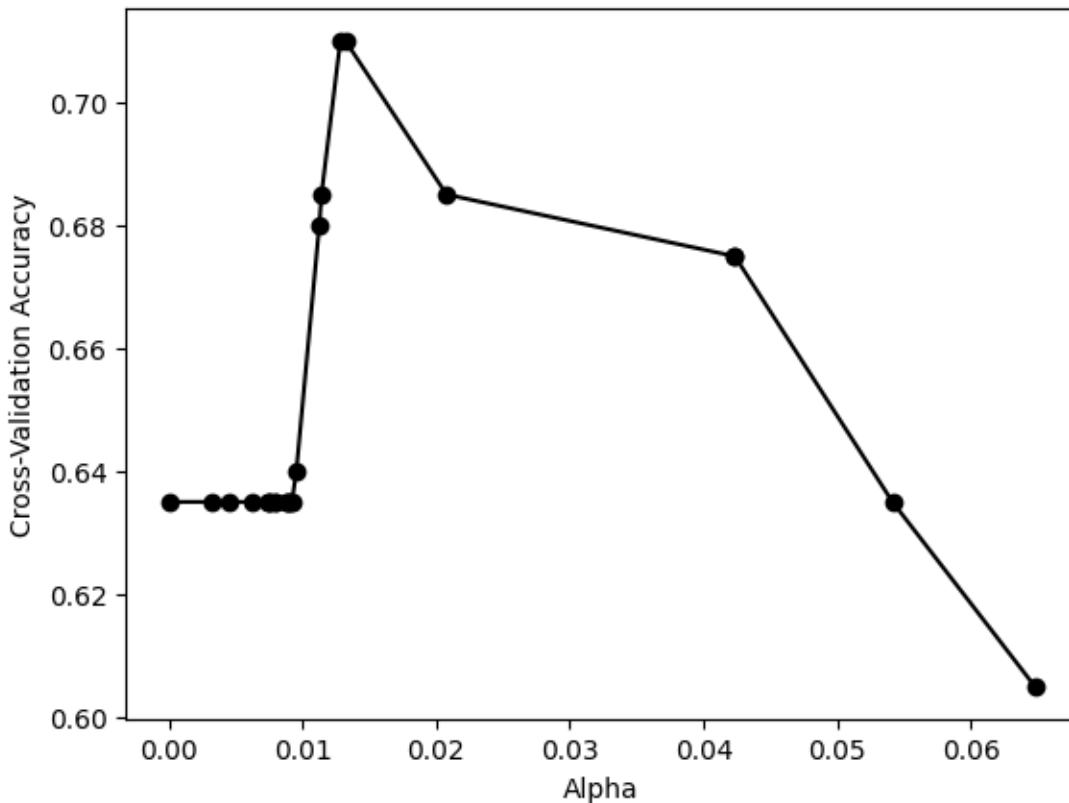
Interpretation Pruning reduces the accuracy from 72.5% to 71.00%. Meaning the pruned tree is less accurate than the unpruned tree.

[27]:

```

plt.plot(path ccp_alphas, cv_scores, marker="o", color="black")
plt.xlabel("Alpha")
plt.ylabel("Cross-Validation Accuracy")
plt.show()

```



Interpretation Previous result in graph form which indicates that at alpha 0.0128 provides us with an accuracy of 71.00%.

[28]:

```

pruned_tree = DecisionTreeClassifier(random_state=0, ccp_alpha=best_alpha)
pruned_tree.fit(X_train, Y_train)

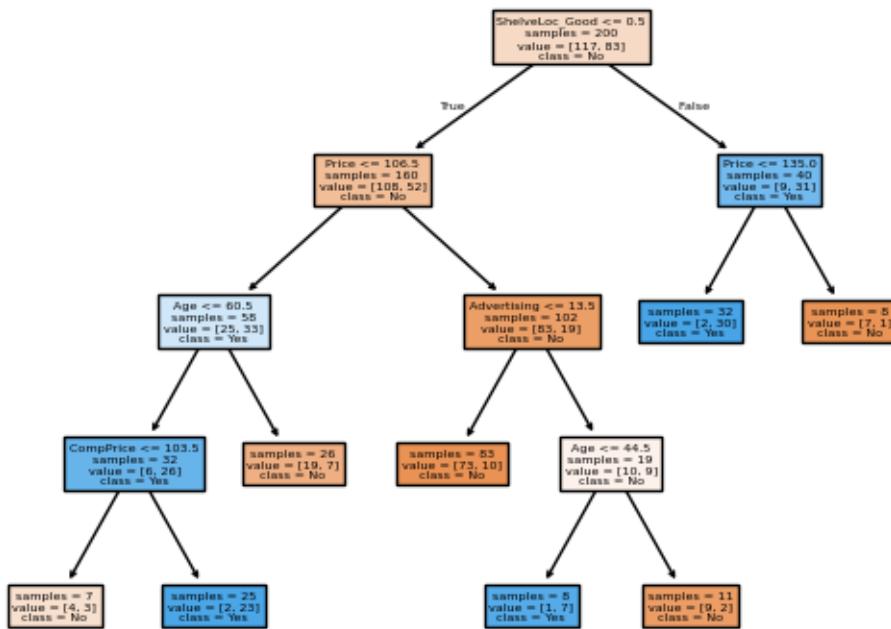
```

```

plot_tree(
    pruned_tree,
    filled=True,
    feature_names=X_train.columns,
    class_names=["No", "Yes"],
    impurity=False,
)
plt.title("Pruned Decision Tree")
plt.show()

```

Pruned Decision Tree



[29]: Y_pred = pruned_tree.predict(X_test)

```

conf_matrix = confusion_matrix(Y_test, Y_pred)
print(conf_matrix)

```

```

[[100  19]
 [ 36  45]]

```

Interpretation Results in the same accuracy as before (72.50%) with a more interpretable tree.

```
[31]: best_alpha = 0.0
for alpha in path ccp_alphas:
    pruned_tree = DecisionTreeClassifier(random_state=0, ccp_alpha=alpha)
    pruned_tree.fit(X_train, Y_train)
    if pruned_tree.get_n_leaves() <= 18:
        best_alpha = alpha
        break

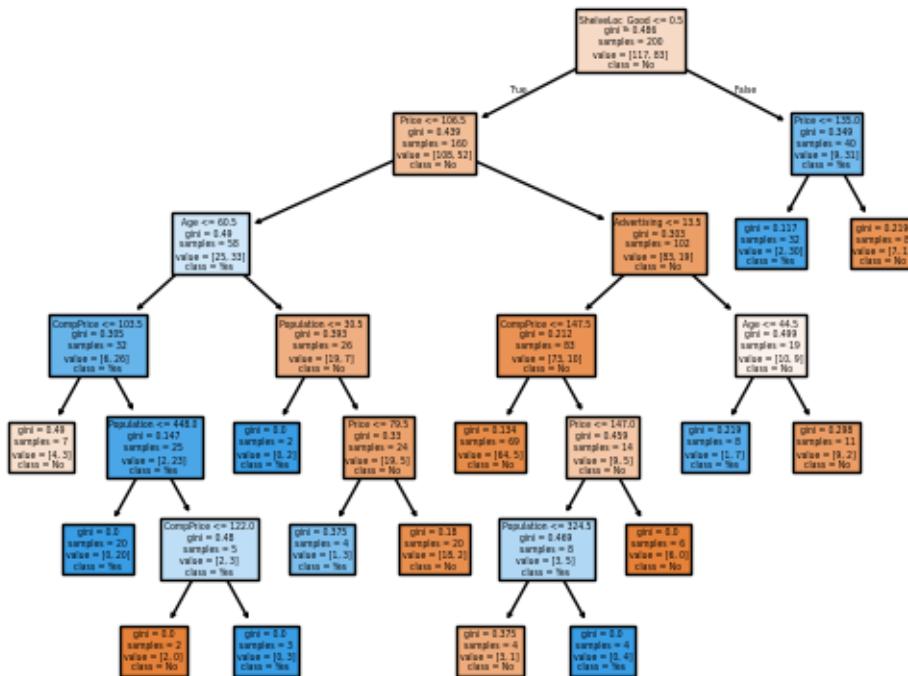
pruned_tree = DecisionTreeClassifier(random_state=0, ccp_alpha=best_alpha)
pruned_tree.fit(X_train, Y_train)
Y_pred = pruned_tree.predict(X_test)

conf_matrix = confusion_matrix(Y_test, Y_pred)

print(conf_matrix)

plt.figure()
plot_tree(
    pruned_tree, filled=True, feature_names=X_train.columns, class_names=["No", "Yes"]
)
plt.show()
```

$\begin{bmatrix} 98 & 21 \\ 33 & 48 \end{bmatrix}$



Interpretation Now we are seeing improved accuracy to 73.00%

Load Data pt2

```
[32]: df = pd.read_csv("../data/Boston.csv", index_col=0)
```

Regression tree

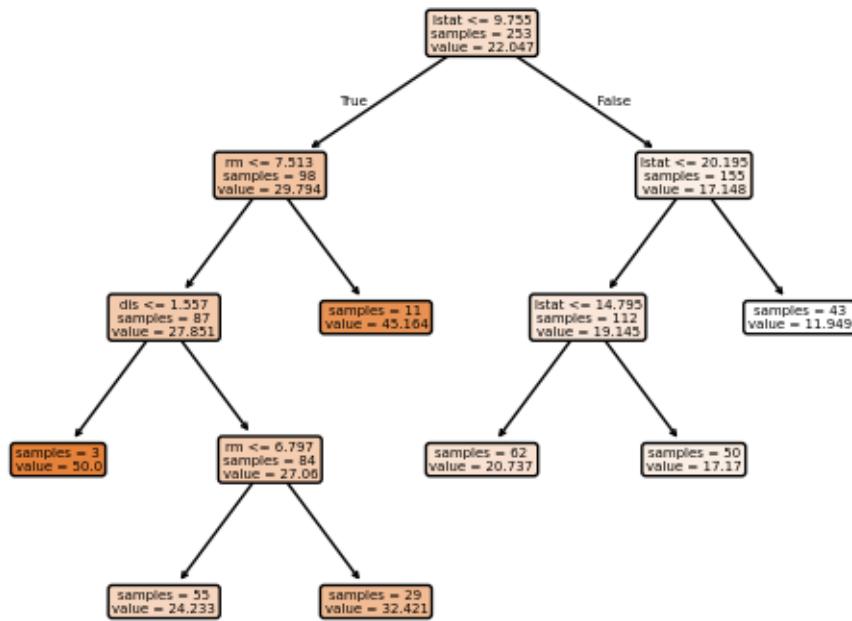
```
[33]: X = df.drop(columns=["medv"])
Y = df["medv"]
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.5, random_state=1)
tree = DecisionTreeRegressor(random_state=1, max_leaf_nodes=7)
tree.fit(X_train, Y_train)

print(f"Number of leaves: {tree.get_n_leaves()}")
print(f"Tree Depth: {tree.get_depth()}")
print(f"Residual Mean Deviance: {tree.tree_.impurity.mean():.4f}")
```

```
Number of leaves: 7
Tree Depth: 4
Residual Mean Deviance: 26.0916
```

Plot Tree

```
[35]: plot_tree(
    tree, filled=True, feature_names=X.columns, impurity=False, rounded=True, fontsize=5
)
plt.show()
```



Interpretation Residual mean deviance is a measure of how well the regression tree fits the data. The value of 26.09 indicates that on average the squared difference between the predicted and actual values is 26.09. A lower value would indicate a better fit.

```
[36]: path = tree.cost_complexity_pruning_path(X_train, Y_train)

cv_scores = []
best_accuracy = float("inf")
best_alpha = 0.0
for alpha in path ccp_alphas:
    pruned_tree = DecisionTreeRegressor(random_state=0, ccp_alpha=alpha)
    scores = -cross_val_score(
        pruned_tree, X_train, Y_train, cv=5, scoring="neg_mean_squared_error"
    )
    accuracy = scores.mean()
    cv_scores.append(accuracy)
    if accuracy < best_accuracy:
        best_accuracy = accuracy
        best_alpha = alpha

pruned_tree = DecisionTreeRegressor(random_state=0, ccp_alpha=best_alpha)
pruned_tree.fit(X_train, Y_train)
Y_pred = pruned_tree.predict(X_test)
```

```

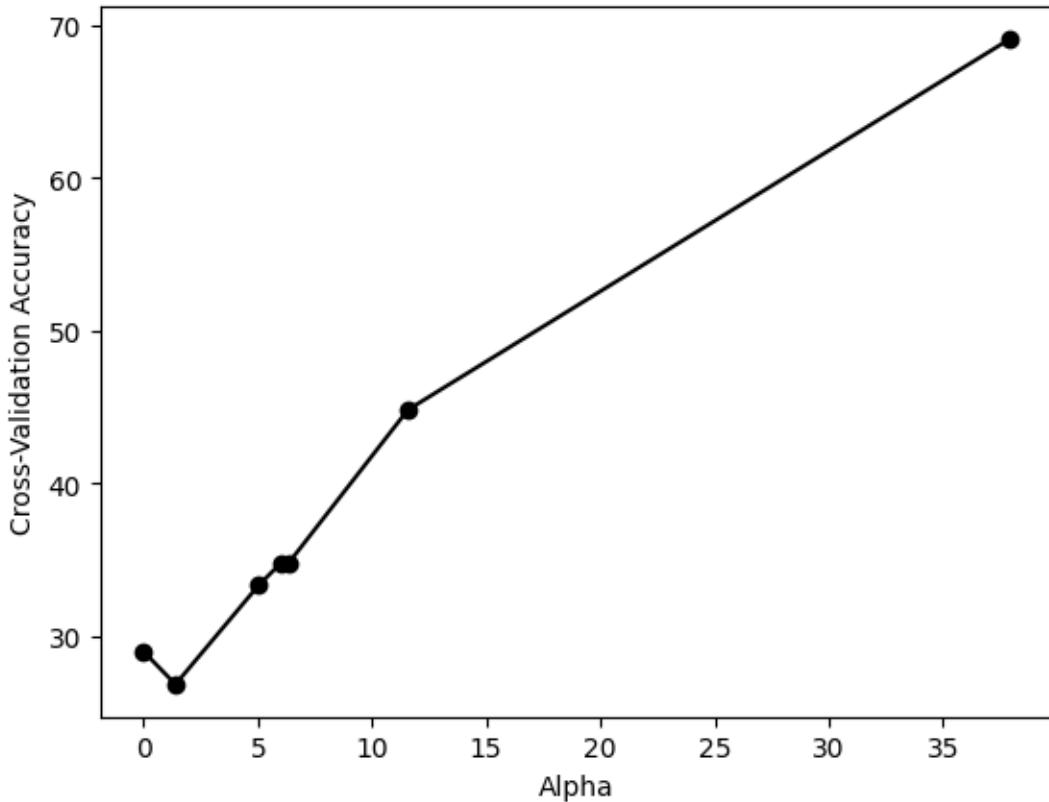
MSE = mean_squared_error(Y_test, Y_pred)

print(f"Best Error: {MSE:.4f}, with alpha: {best_alpha:.4f}")

plt.plot(path ccp_alphas, cv_scores, marker="o", color="black")
plt.xlabel("Alpha")
plt.ylabel("Cross-Validation Accuracy")
plt.show()

```

Best Error: 17.9246, with alpha: 1.3920



Interpretation An error of 17.92 (with pruning) is an improvement compared to earlier.

```

[51]: Y_hat = tree.predict(X_test)

plt.figure()
plt.scatter(Y_hat, Y_test, alpha=0.5, edgecolors="k")
plt.plot(
    [min(Y_hat), max(Y_hat)], [min(Y_pred), max(Y_pred)], color="red", 
    linestyle="--"
)

```

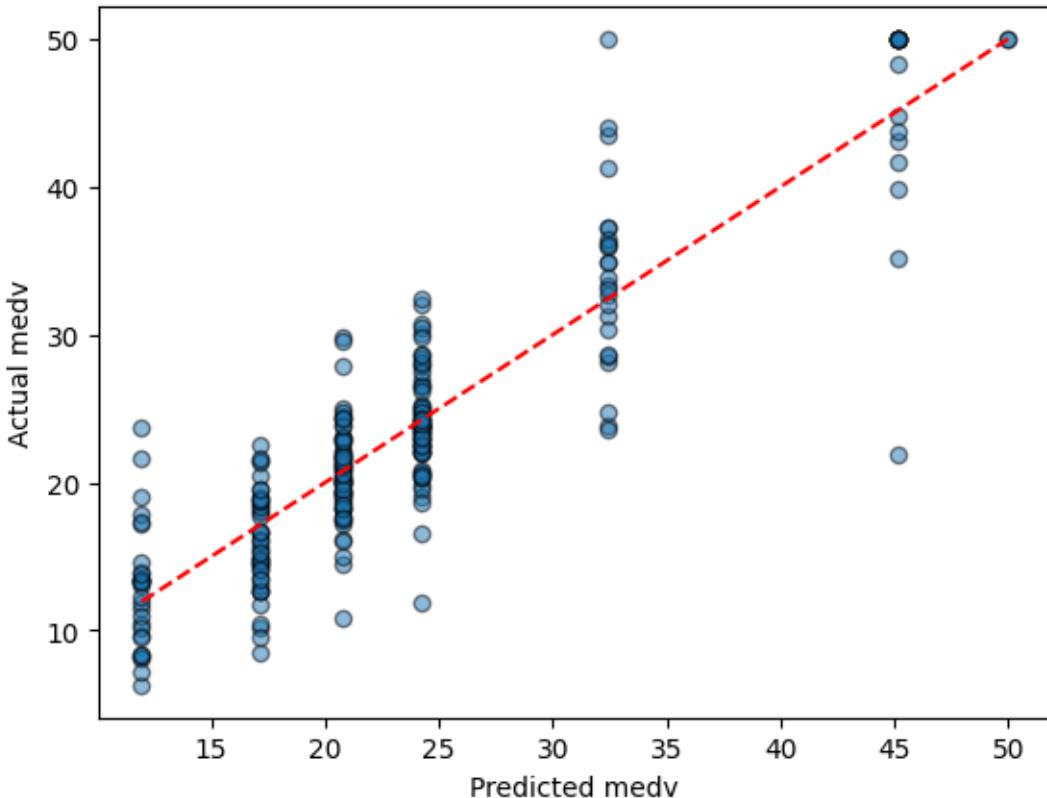
```

plt.xlabel("Predicted medv")
plt.ylabel("Actual medv")
plt.show()

MSE = np.mean((Y_hat - Y_test) ** 2)

print(f"Mean Squared Error: {MSE:.4f}")

```



Mean Squared Error: 17.9246

Interpretation It follows the the actual values decently well, with the Mean Squared Error (MSE) of 17.92 after pruning.

Reggression Bagging & Random Forests

```

[41]: random_forest = RandomForestRegressor(n_estimators=500, max_features=13,
                                          random_state=1)
random_forest.fit(X_train, Y_train)

Y_hat_bag = random_forest.predict(X_test)
MSE_bag = mean_squared_error(Y_test, Y_hat_bag)
r_squared = r2_score(Y_test, Y_hat_bag)

```

```

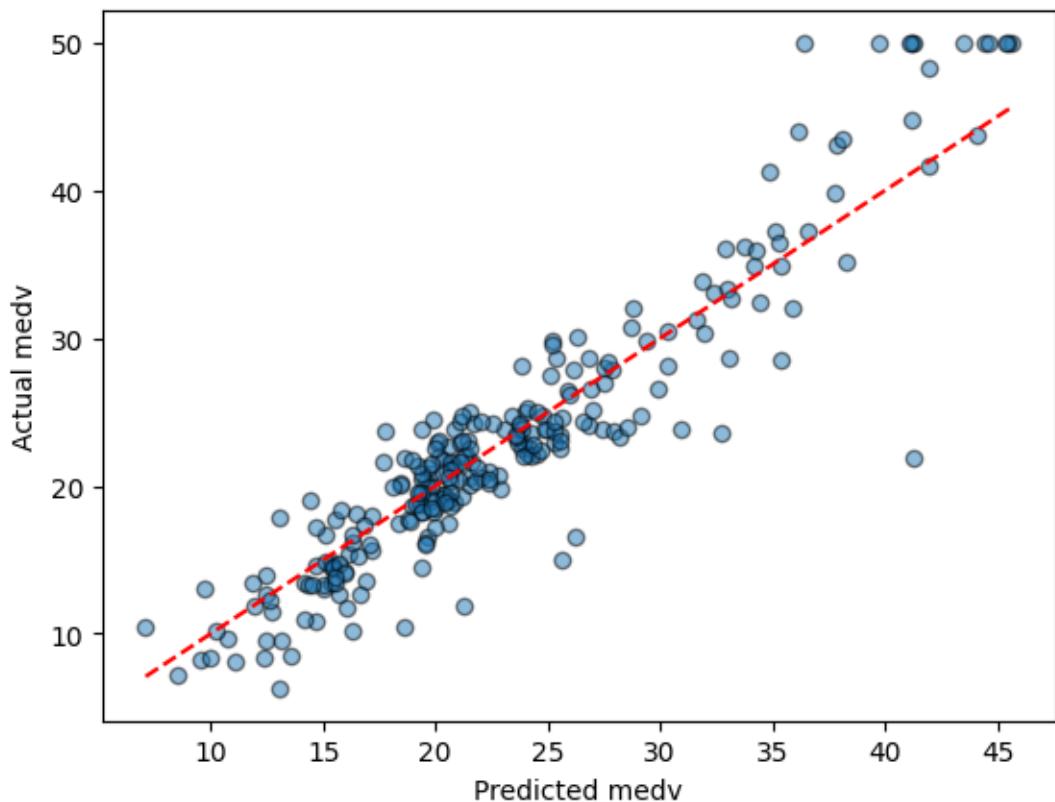
print(f"Mean Squared Error: {MSE_bag:.4f}")
print(f"R-squared: {100 * r_squared:.2f}%")

plt.figure()
plt.scatter(Y_hat_bag, Y_test, alpha=0.5, edgecolors="k")
plt.plot(
    [min(Y_hat_bag), max(Y_hat_bag)],
    [min(Y_hat_bag), max(Y_hat_bag)],
    color="red",
    linestyle="--",
)
plt.xlabel("Predicted medv")
plt.ylabel("Actual medv")
plt.show()

```

Mean Squared Error: 11.7347

R-squared: 86.63%



Interpretation Better Mean Squared Error (MSE) of 11.49 compared to the previous tree.

```
[42]: random_forest = RandomForestRegressor(n_estimators=25, max_features=13,random_state=1)
random_forest.fit(X_train, Y_train)

Y_hat_bag = random_forest.predict(X_test)
MSE_bag = mean_squared_error(Y_test, Y_hat_bag)
r_squared = r2_score(Y_test, Y_hat_bag)

print(f"Mean Squared Error: {MSE_bag:.4f}")
print(f"R-squared: {100 * r_squared:.2f}%")
```

Mean Squared Error: 13.2764

R-squared: 84.87%

Interpretation This Mean Squared Error of 13.27 is worse than the previously due to less parameters being used in the model.

```
[ ]: random_forest = RandomForestRegressor(n_estimators=500, max_features=6,random_state=1)
random_forest.fit(X_train, Y_train)

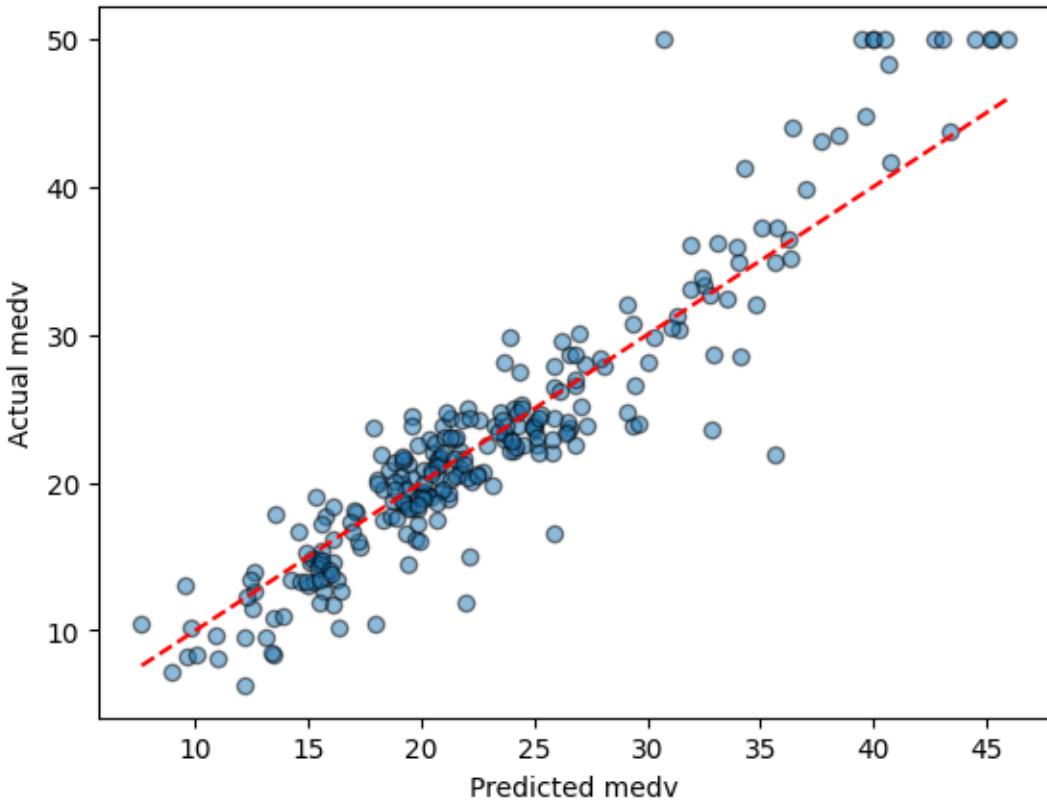
Y_hat_bag = random_forest.predict(X_test)
MSE_bag = mean_squared_error(Y_test, Y_hat_bag)
r_squared = r2_score(Y_test, Y_hat_bag)

print(f"Mean Squared Error: {MSE_bag:.4f}")
print(f"R-squared: {100 * r_squared:.2f}%")

plt.figure()
plt.scatter(Y_hat_bag, Y_test, alpha=0.5, edgecolors="k")
plt.plot(
    [min(Y_hat_bag), max(Y_hat_bag)],
    [min(Y_hat_bag), max(Y_hat_bag)],
    color="red",
    linestyle="--",
)
plt.xlabel("Predicted medv")
plt.ylabel("Actual medv")
plt.show()
```

Mean Squared Error: 11.7343

R-squared: 86.63%



Interpretation Similar MSE as before but not the same.

```
[45]: feat_importance = random_forest.feature_importances_

sorted_indices = np.argsort(feat_importance)[::-1]
feat_names = X.columns

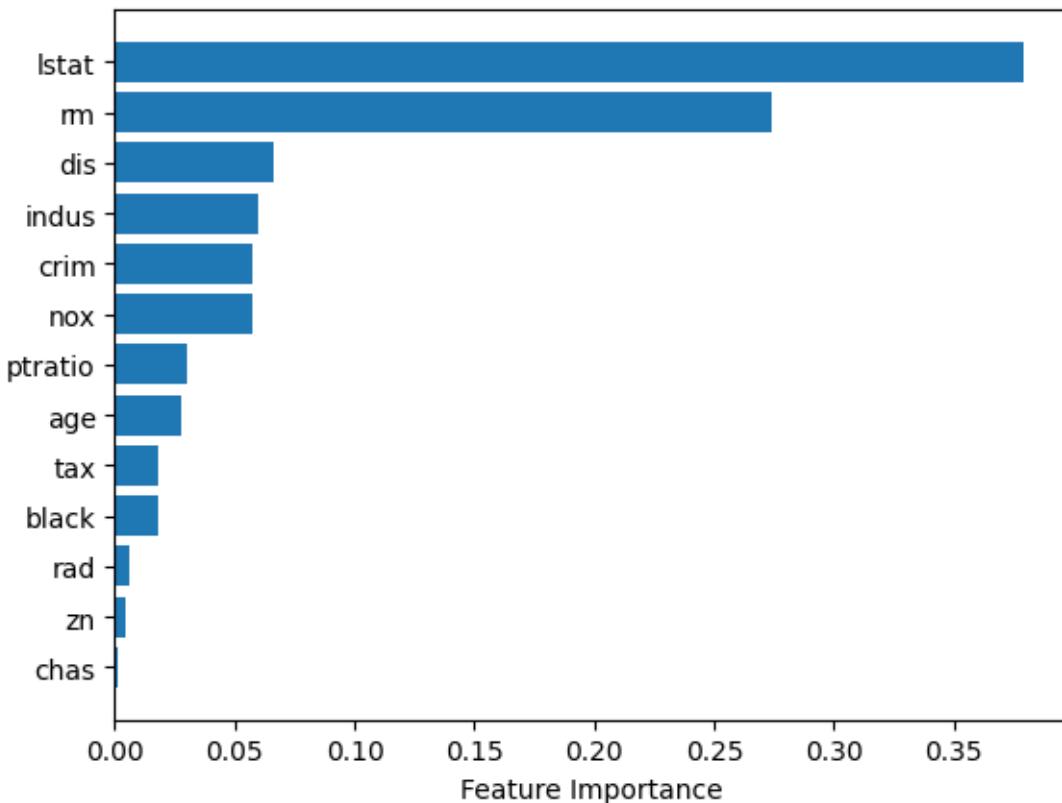
print("Feature Importances:")
for idx in sorted_indices:
    print(f"{feat_names[idx]}: {feat_importance[idx]:.4f}")

plt.figure()
plt.barh(range(len(feat_importance)), feat_importance[sorted_indices], align="center")
plt.yticks(range(len(feat_importance)), np.array(feat_names)[sorted_indices])
plt.xlabel("Feature Importance")
plt.gca().invert_yaxis()
plt.show()
```

Feature Importances:

lstat: 0.3789

```
rm: 0.2739  
dis: 0.0660  
indus: 0.0601  
crim: 0.0576  
nox: 0.0571  
ptratio: 0.0302  
age: 0.0277  
tax: 0.0182  
black: 0.0180  
rad: 0.0059  
zn: 0.0047  
chas: 0.0016
```



Interpretation The graph indicates that the `lstat` is the most important or influential feature when predicting, meanwhile `chas` is the least important feature.

```
[46]: grad_boost = GradientBoostingRegressor(  
      loss="squared_error", n_estimators=5000, learning_rate=0.001, random_state=1  
)  
grad_boost.fit(X_train, Y_train)  
feat_importance = grad_boost.feature_importances_
```

```

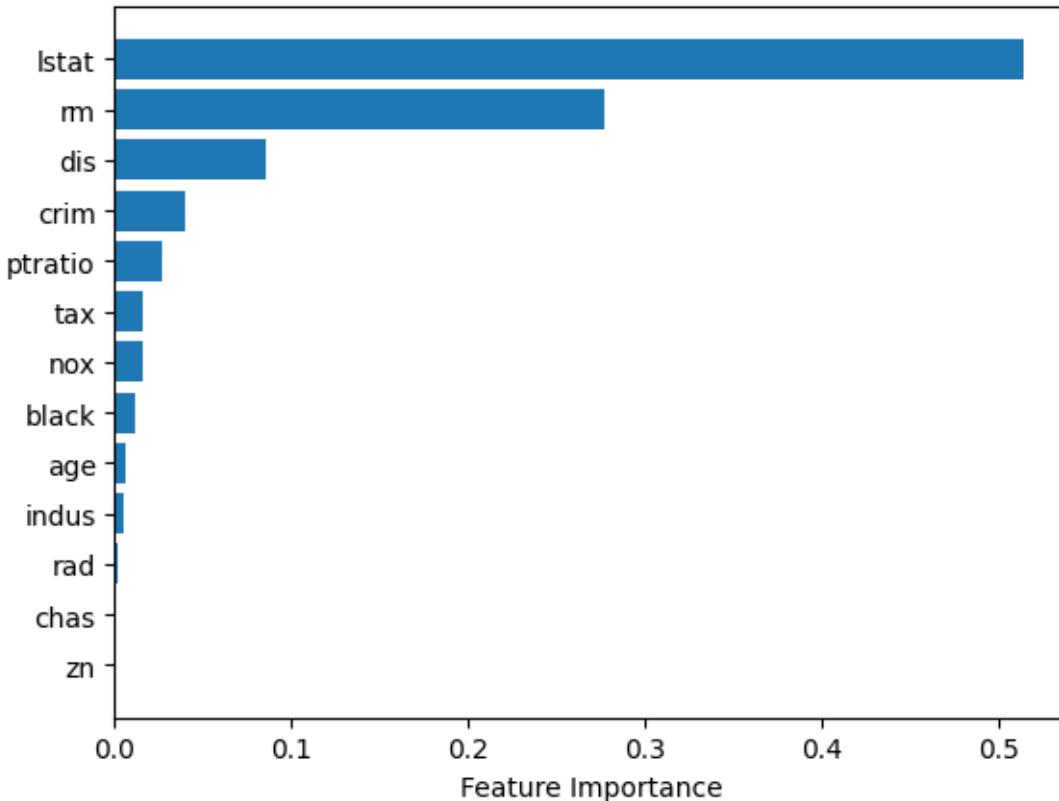
sorted_indices = np.argsort(feat_importance)[::-1]
feat_names = X.columns

for idx in sorted_indices:
    print(f"{feat_names[idx]}: {feat_importance[idx]:.4f}")

plt.figure()
plt.barh(range(len(feat_importance)), feat_importance[sorted_indices], align="center")
plt.yticks(range(len(feat_importance)), np.array(feat_names)[sorted_indices])
plt.xlabel("Feature Importance")
plt.gca().invert_yaxis()
plt.show()

```

lstat: 0.5136
rm: 0.2772
dis: 0.0851
crim: 0.0398
ptratio: 0.0266
tax: 0.0164
nox: 0.0159
black: 0.0117
age: 0.0063
indus: 0.0048
rad: 0.0018
chas: 0.0007
zn: 0.0000



Interpretation Same as before, but `zn` is now the least important feature.

```
[48]: Y_hat = grad_boost.predict(X_test)
MSE_grad_boost = mean_squared_error(Y_test, Y_hat)

print(f"Mean Squared Error: {MSE_grad_boost:.4f}")
```

Mean Squared Error: 11.0007

Interpretation Slight improvement in the Mean Squared Error to 11.00. Probably need a significance test if it is significant.

```
[50]: grad_boost = GradientBoostingRegressor(
    loss="squared_error", n_estimators=5000, learning_rate=0.2, random_state=1
)
grad_boost.fit(X_train, Y_train)
feat_importance = grad_boost.feature_importances_
sorted_indices = np.argsort(feat_importance)[::-1]
feat_names = X.columns

for idx in sorted_indices:
```

```

print(f"{feat_names[idx]}\t{feat_importance[idx]:.4f}")

plt.figure()
plt.barh(range(len(feat_importance)), feat_importance[sorted_indices], align="center")
plt.yticks(range(len(feat_importance)), np.array(feat_names)[sorted_indices])
plt.xlabel("Feature Importance")
plt.gca().invert_yaxis()
plt.show()

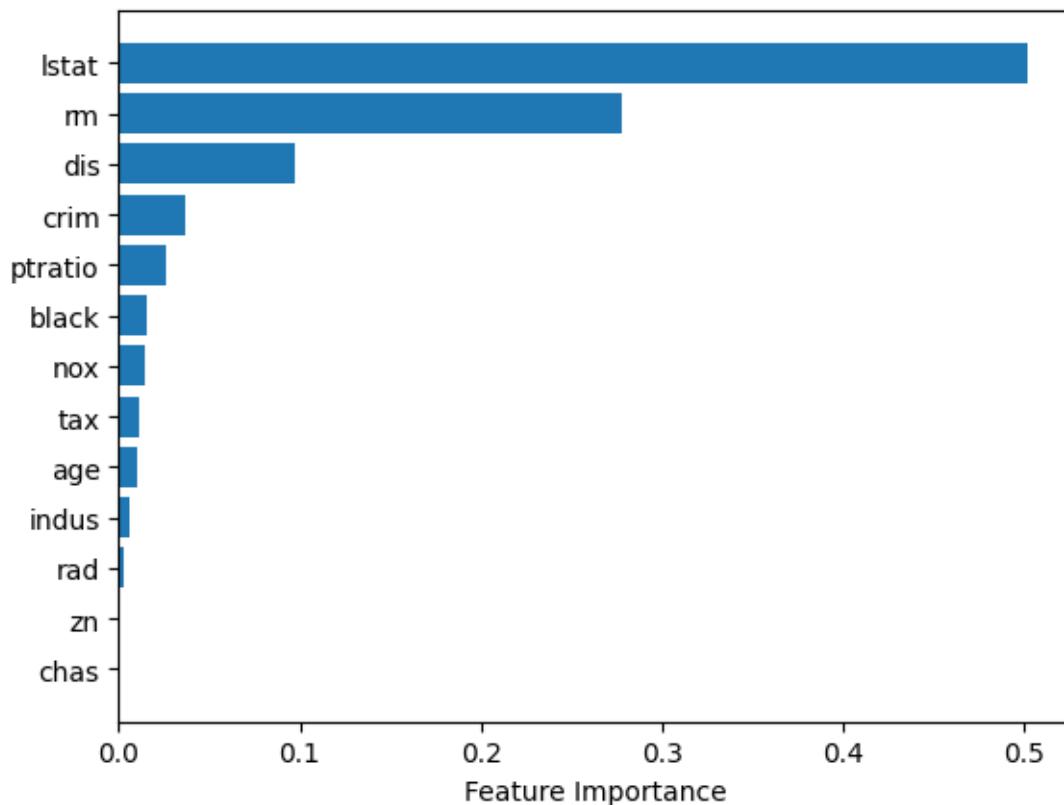
Y_hat = grad_boost.predict(X_test)
MSE_grad_boost = mean_squared_error(Y_test, Y_hat)

print(f"Mean Squared Error: {MSE_grad_boost:.4f}")

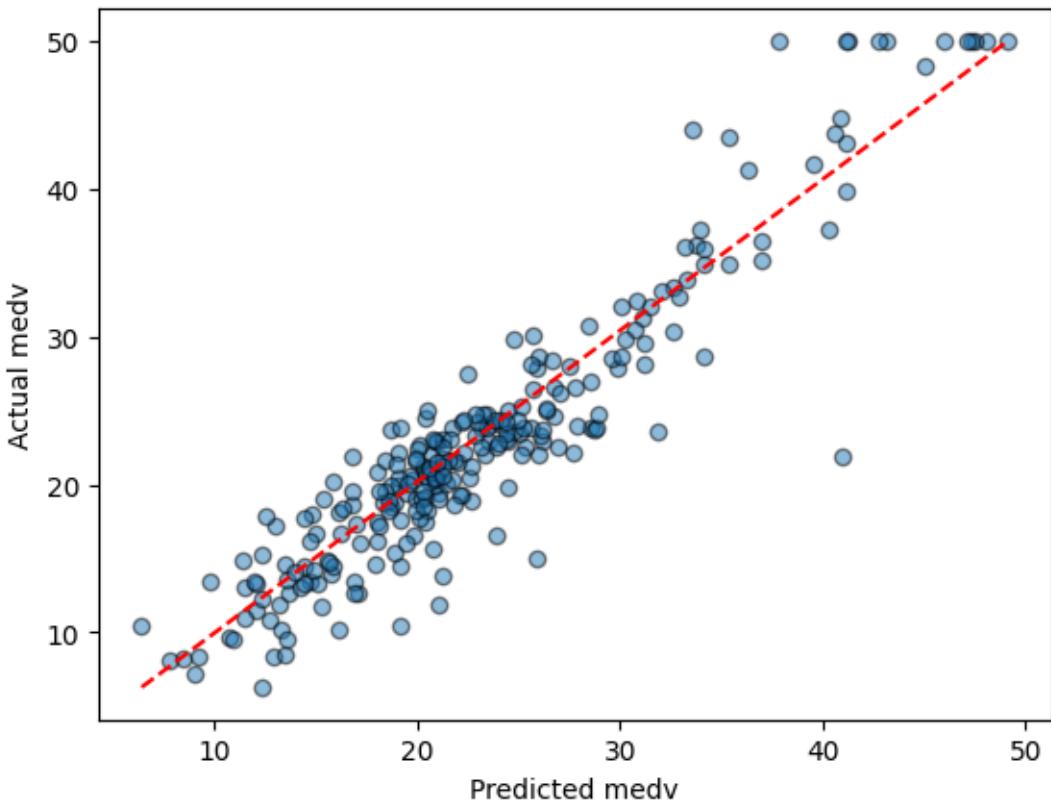
plt.figure()
plt.scatter(Y_hat, Y_test, alpha=0.5, edgecolors="k")
plt.plot(
    [min(Y_hat), max(Y_hat)], [min(Y_test), max(Y_test)], color="red", linestyle="--"
)
plt.xlabel("Predicted medv")
plt.ylabel("Actual medv")
plt.show()

```

lstat: 0.5015
rm: 0.2780
dis: 0.0971
crim: 0.0363
ptratio: 0.0259
black: 0.0159
nox: 0.0143
tax: 0.0117
age: 0.0099
indus: 0.0059
rad: 0.0027
zn: 0.0006
chas: 0.0004



Mean Squared Error: 11.0775



Interpretation The change of the learning rate to 0.02 resulted in a worse score. The above scatter plot shows the predicted vs actual values.